Foundation Models in Bioinformatics: A New Era of AI-Driven Biological Discovery

Generated at: 2026-04-03
Source URL: https://www.compbiosci.com/posts/foundation-models-in-bioinformatics-a-new-era-of-aidriven-biological-discovery

Foundation Models in Bioinformatics: A New Era of AI-Driven Biological Discovery

Abstract

Foundation Models (FMs) are instigating a paradigm shift in bioinformatics, offering powerful solutions to long-standing challenges such as limited annotated data and high data noise. This review synthesizes current advancements, focusing on the application of these large-scale, self-supervised AI models across key biological domains like genomics, proteomics, drug discovery, and single-cell analysis. We explore the core architectures of language, vision, graph, and multimodal FMs, providing a guide for researchers to select appropriate models. The article further critically examines methodological applications, persistent challenges in data and interpretability, and rigorous benchmarking studies that pit FMs against traditional methods. Finally, we discuss future trajectories, outlining how FMs are poised to fuel continued innovation in molecular biology and clinical research.

What Are Foundation Models? Redefining AI's Role in Biology

Foundation Models (FMs) represent a transformative paradigm in artificial intelligence, characterized by their training on broad data at scale and adaptability to a wide range of downstream tasks [1]. The term was coined by researchers at the Stanford Institute for Human-Centered Artificial Intelligence in 2021 to describe models that are "trained on broad data (generally using self-supervision at scale) that can be adapted to a wide range of downstream tasks" [1]. These models have fundamentally changed how data scientists approach machine learning by providing a versatile starting point for developing specialized applications more quickly and cost-effectively than training models from scratch [2].

In bioinformatics, FMs address longstanding challenges such as limited annotated data and data noise by leveraging self-supervised learning on massive, unlabeled datasets [3]. Their capacity to learn accurate representations of intricate biological data through data-intensive pretraining enables researchers to utilize these models for various downstream tasks with limited data through fine-tuning mechanisms [3]. This flexibility has positioned FMs as essential tools for tackling core biological problems including sequence analysis, structure prediction, and function annotation across diverse data modalities including DNA, RNA, proteins, and single-cell multi-omics data [3].

Core Architecture and Technical Foundation

Defining Characteristics of Foundation Models

Foundation models exhibit several distinguishing characteristics that separate them from traditional machine learning approaches:

- **Broad training data**: FMs are trained on extensive, generalized datasets that span multiple domains or modalities, enabling wide applicability [2] [1].
- **Self-supervised learning**: These models primarily use self-supervised learning objectives during pretraining, creating labels directly from the input data without human annotation [2] [4].
- **Adaptability**: A fundamental feature of FMs is their ability to be adapted through fine-tuning, prompt engineering, or other techniques to perform specialized downstream tasks [2] [1].
- **Scale**: Advanced FMs typically contain tens of billions of parameters and require substantial computational resources for training [2] [1].

The adaptability of FMs is particularly valuable in bioinformatics, where they can be fine-tuned for specific biological problems, narrowing the representation gap between general domain knowledge and specialized biological domain knowledge [3].

The Role of Self-Supervised Learning

Self-supervised learning has become the cornerstone of modern foundation models, enabling them to learn powerful representations from vast amounts of unlabeled data [5]. By designing auxiliary tasks on raw inputs, SSL removes the reliance on human-provided labels and underpins the pretraining-finetuning paradigm that has reshaped machine learning beyond traditional empirical risk minimization frameworks [5].

In the context of FMs, SSL works by creating learning signals directly from the structure of the data itself. For example, in language models, this might involve predicting masked words in a sequence, while in biological sequences, it might involve predicting masked nucleotides or amino acids [3]. The model learns meaningful representations by solving these pretext tasks, capturing underlying patterns and relationships in the data without explicit human labeling [4] [5].

Table 1: Common Self-Supervised Learning Objectives in Foundation Models

SSL Objective	Mechanism	Example Models	Biological Applications
Masked Language Modeling	Randomly masks portions of input and trains model to predict masked content	BERT, DNABERT [3]	Predicting masked nucleotides or amino acids in sequences
Contrastive Learning	Learns representations by contrasting similar and dissimilar pairs	CONCH [6]	Matching histopathology images with textual descriptions
Autoregressive Modeling	Predicts next element in a sequence given previous elements	GPT series, scGPT [2] [6]	Generating biological sequences or predicting next element
Replaced Token Detection	Distinguishes between real input elements and plausible fakes	BioELECTRA [6]	Identifying functional genomic elements

Architectural Components

Foundation models leverage sophisticated neural network architectures optimized for handling large-scale, complex data:

- **Transformer architectures**: Many FMs utilize transformer networks with self-attention mechanisms that can capture long-range dependencies in sequential data [2] [3]. The bidirectional nature of transformers in models like BERT enables comprehensive context understanding [3].
- **Graph Neural Networks**: For biological data with inherent graph structure, GNNs process graph-structured information where nodes represent biological entities and edges represent relationships [3] [7]. Models like TxGNN use GNNs to embed drugs and diseases into a latent representational space that reflects the geometry of medical knowledge graphs [7].
- **Multimodal architectures**: Advanced FMs can process and integrate multiple data types through unified architectures. For example, CONCH combines visual and language-based learning for histopathology analysis [6].

The following diagram illustrates the core self-supervised learning process that underpins foundation model pretraining:

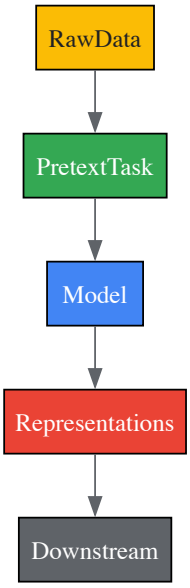

Diagram 1: Self-Supervised Learning Process

Adaptation Methodologies: From Pretraining to Specialized Applications

The Adaptation Pipeline

Foundation models transition from general pretraining to specialized applications through carefully designed adaptation methodologies. This process typically involves two main phases: pretraining and adaptation [4]. During pretraining, the model learns general representations from a massive dataset using self-supervised learning objectives. The adaptation phase then customizes the model for specific downstream tasks, which may involve fine-tuning, linear probing, or other techniques [4].

In bioinformatics, this adaptation pipeline is particularly valuable because it allows researchers to leverage knowledge learned from large-scale biological datasets and apply it to specialized problems with limited labeled data [3]. The adaptation process can be visualized as follows:

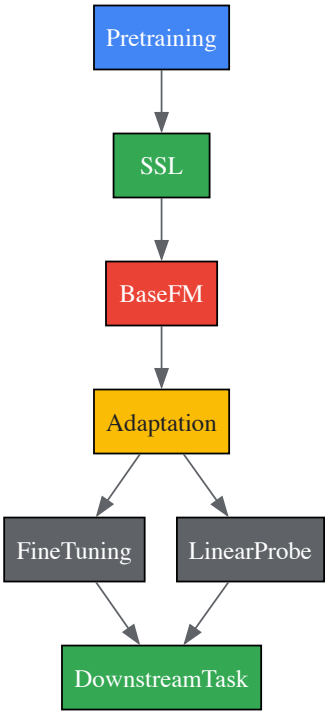

Diagram 2: Foundation Model Adaptation Pipeline

Key Adaptation Techniques

The adaptation of foundation models to specialized tasks can be accomplished through several technical approaches:

- **Fine-tuning**: This approach updates both the pretrained model parameters and the task-specific head using the downstream dataset [4]. Fine-tuning allows the model to adjust its representations to the specific nuances of the target task while retaining knowledge from pretraining.
- **Linear probing**: In this method, the pretrained model parameters remain frozen, and only a linear classification head is trained on the downstream task [4]. This approach is useful for assessing the quality of the learned representations and for scenarios with limited computational resources.
- **Prompt engineering**: Particularly prevalent in language models, this technique involves crafting input prompts to steer the model toward desired outputs without updating model parameters [2].
- **Metric learning for zero-shot prediction**: Models like TxGNN implement metric learning components that transfer knowledge from well-annotated domains to those with limited treatments by measuring similarities between entities in the learned representation space [7].

Table 2: Adaptation Techniques for Foundation Models in Bioinformatics

Adaptation Method	Parameters Updated	Data Requirements	Use Cases in Bioinformatics
Fine-tuning	All parameters	Moderate labeled data	Specializing protein language models for specific families
Linear Probing	Only final layer	Limited labeled data	Rapid prototyping, representation quality assessment
Prompt Engineering	None	No labeled data	Leveraging LLMs for biological knowledge extraction
Metric Learning	Similarity metrics	Limited to no labeled data	Zero-shot drug repurposing [7]

Foundation Models in Bioinformatics: Case Studies and Experimental Approaches

TxGNN: A Case Study in Drug Repurposing

TxGNN represents a pioneering application of foundation models in bioinformatics, specifically designed for zero-shot drug repurposing [7]. This model addresses the critical challenge of identifying therapeutic candidates for diseases with limited treatment options or no existing drugs—a scenario that affects approximately 92% of the 17,080 diseases examined in the original study [7].

Experimental Protocol for TxGNN:

- **Knowledge Graph Construction**: TxGNN is trained on a comprehensive medical knowledge graph that collates decades of biological research across 17,080 diseases, containing 9,388 indications and 30,675 contraindications [7].
- **Model Architecture**: The framework consists of two main modules:
  - **TxGNN Predictor**: A graph neural network optimized on relationships in the medical KG that produces meaningful representations for all concepts through large-scale, self-supervised pretraining [7].
  - **TxGNN Explainer**: A module that provides interpretable insights through multi-hop medical knowledge paths that form the predictive rationales [7].
- **Zero-shot Prediction Mechanism**: TxGNN implements a metric learning component that creates disease signature vectors based on neighborhood topology in the KG. When querying a specific disease, it retrieves similar diseases, generates embeddings for them, and adaptively aggregates them based on similarity to the queried disease [7].
- **Evaluation**: Under stringent zero-shot evaluation, TxGNN improved prediction accuracy for indications by 49.2% and contraindications by 35.1% compared to eight benchmark methods [7].

The following diagram illustrates the TxGNN architecture and workflow for drug repurposing:

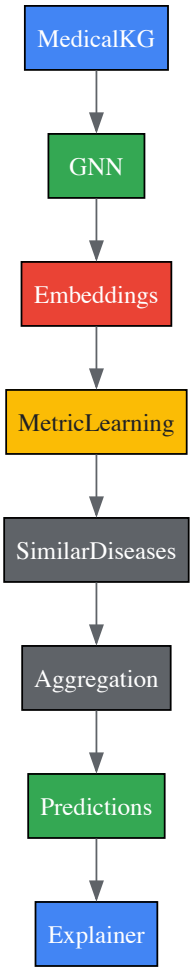

Diagram 3: TxGNN Architecture for Drug Repurposing

Experimental Protocols for Biological Foundation Models

The evaluation of foundation models in bioinformatics requires specialized experimental protocols that account for the unique characteristics of biological data and research questions:

Benchmarking Protocol for Sequence-Based FMs:

- **Data Partitioning**: Implement rigorous cross-validation strategies that account for sequence homology to avoid data leakage [3].
- **Baseline Comparison**: Compare against traditional machine learning approaches and specialized bioinformatics tools [3] [7].

- **Performance Metrics**: Utilize domain-appropriate metrics including area under the precision-recall curve (AUPRC), receiver operating characteristic (ROC) curves, and rank-based metrics [7].
- **Biological Significance Validation**: Where possible, validate computational predictions with experimental evidence from literature or new experiments [7] [6].

Interpretability Analysis Protocol:

- **Attention Visualization**: For transformer-based models, analyze attention patterns to identify biologically relevant regions in sequences [3].
- **Pathway Enrichment**: For models generating embeddings, perform enrichment analysis on genes or proteins with similar representations [6].
- **Ablation Studies**: Systematically remove components of the model or input features to assess their contribution to performance [7].

Table 3: Key Foundation Models in Bioinformatics and Their Applications

Model	Modality	Architecture	Primary Applications	Notable Achievements
AlphaFold [6]	Protein sequences	Neural networks	3D protein structure prediction	Near-experimental accuracy, 2024 Nobel Prize in Chemistry
DNABERT [3] [6]	DNA sequences	Transformer	Predicting regulatory regions	Adapts BERT to understand DNA sequences contextually
Geneformer [6]	Single-cell transcriptomics	Transformer	Predicting tissue-specific gene network dynamics	Trained on 30-95 million single-cell transcriptomes
scGPT [6]	Single-cell data	Transformer	Cell annotation, gene network inference	Distills insights from ~30 million cells
TxGNN [7]	Medical knowledge graph	Graph Neural Network	Zero-shot drug repurposing	49.2% improvement in indication prediction accuracy

Essential Research Reagents and Computational Tools

The development and application of foundation models in bioinformatics requires specialized computational "reagents" and resources. The following table details key resources that constitute the essential toolkit for researchers in this field:

Table 4: Research Reagent Solutions for Foundation Model Development in Bioinformatics

Resource Category	Specific Tools/Platforms	Function	Application Examples
Model Architectures	Transformer, GNN, VAE	Core neural network architectures for processing different data types	Transformers for sequences [3], GNNs for knowledge graphs [7]
Pretraining Frameworks	PyTorch, TensorFlow, JAX	Enable efficient large-scale model training	Implementing self-supervised learning objectives [5]
Biological Data Resources	NCBI, UniProt, Protein Data Bank	Provide structured biological data for pretraining	Protein sequences, 3D structures, genomic data [6]
Specialized Model Hubs	Hugging Face, Model Zoo	Repository of pretrained models for adaptation	Access to nearly 200,000 models [2]
Knowledge Graphs	Medical KG [7], Biological networks	Structured knowledge representation	TxGNN training [7], relationship mining
Interpretability Tools	GraphMask [7], Attention visualization	Explain model predictions and rationales	TxGNN Explainer module [7]

Challenges and Future Directions

Despite their remarkable capabilities, foundation models in bioinformatics face several significant challenges that represent opportunities for future research:

Technical and Methodological Challenges

- **Data quality and bias**: Foundation models can pick up inappropriate patterns or biases from training data, requiring careful data filtering and norm encoding [2]. In biological contexts, this may manifest as overrepresentation of well-studied organisms or pathways.
- **Interpretability and reliability**: The complex, nonlinear features extracted by FMs face challenges regarding biological interpretability and model reliability [3]. Methods like TxGNN's Explainer module represent important steps toward addressing this challenge [7].
- **Computational requirements**: Building a foundation model from scratch is expensive and requires enormous resources, with training potentially taking months [2]. This creates barriers to entry for many research institutions.
- **Generalization limitations**: While FMs exhibit strong performance across many tasks, they may struggle with out-of-distribution data or rare biological phenomena not well-represented in training data [3].

Emerging Research Directions

Several promising research directions are emerging at the intersection of foundation models and bioinformatics:

- **Multimodal foundation models**: Integrating multiple data types (e.g., sequences, structures, images, text) within unified architectures to capture complementary biological information [3] [6].
- **Federated learning approaches**: Developing methods to train FMs across distributed biological datasets while maintaining data privacy and security [3].
- **Causal representation learning**: Moving beyond correlational patterns to discover causal relationships in biological systems [3].
- **Resource-efficient adaptation**: Creating methods that enable effective adaptation of large FMs with limited computational resources and labeled data [4] [3].

The rapid evolution of foundation models continues to reshape the bioinformatics landscape, offering unprecedented opportunities to decipher complex biological systems and accelerate therapeutic development. As these models become more sophisticated, accessible, and interpretable, they are poised to become indispensable tools in the researcher's toolkit, ultimately bridging the gap between data-intensive biology and actionable biological insights.

The field of bioinformatics has undergone a profound transformation, driven by the exponential accumulation of biological data from high-throughput sequencing technologies and multi-omics approaches [8]. This data deluge posed significant challenges for analysis and interpretation, creating an urgent need for more sophisticated computational approaches [8]. Concurrently, artificial intelligence (AI) has achieved groundbreaking advances, evolving from specialized, niche models handling specific biological tasks to powerful general-purpose tools capable of addressing fundamental biological questions across multiple domains [8] [9]. This evolution has been marked by the emergence of foundation models (FMs)—large-scale AI systems pretrained on vast datasets that can be adapted to a wide range of downstream tasks [9] [10]. These models represent a paradigm shift from task-specific solutions to versatile tools that leverage transfer learning and self-supervised pretraining to capture universal patterns in biological data [10]. The integration of AI in bioinformatics has now reached a pivotal moment where these systems are not merely analytical tools but collaborative partners in scientific discovery, enabling breakthroughs in genomics, proteomics, drug discovery, and single-cell analysis [8] [9] [10].

The Trajectory of AI in Bioinformatics: From Specific to General

The Era of Traditional Machine Learning

Initial applications of AI in bioinformatics relied heavily on traditional machine learning (ML) methods, which excelled in scenarios with well-defined features and controlled experimental conditions [8]. These approaches operated within a structured framework where models learned functions to minimize objective functions for classification or regression tasks [8]. Support vector machines (SVM) and random forests (RF) became workhorses for analyzing genomic sequencing fragments, physicochemical properties of proteins, and medical imaging signals [8]. While effective for specific, limited-scale datasets, these methods required substantial feature engineering and domain expertise, constraining their applicability to broader biological questions. Their performance was closely tied to data quality and manual curation, making them suitable for targeted analyses but inadequate for extracting generalizable insights from the increasingly complex and massive biological datasets being generated [8].

The Deep Learning Revolution

The limitations of traditional ML prompted a shift toward deep learning approaches, characterized by complex neural networks that perform automated feature extraction and transformation across multiple processing layers [8]. Convolutional neural networks (CNNs) and recurrent neural networks (RNNs) demonstrated remarkable capabilities in capturing spatial and sequential dependencies in biological data, enabling more sophisticated analysis of genomic sequences and protein structures [8]. A pivotal advancement came with the introduction of Transformer architectures, which leveraged self-attention mechanisms to process sequential data with unprecedented effectiveness [8]. Models such as AlphaFold for protein structure prediction and DNABERT for genomic sequence analysis achieved landmark results that surpassed previous methods [8]. These

architectures formed the foundation for the next evolutionary step: the development of foundation models that could generalize across diverse biological domains and tasks [8] [9].

The Emergence of Foundation Models

Foundation models represent the current frontier in AI for bioinformatics, characterized by their large-scale pretraining on extensive datasets and adaptability to numerous downstream applications [9] [10]. These models overcome the fundamental limitation of earlier approaches—their narrow specialization—by learning generalizable patterns from massive, diverse biological data corpora [10]. The defining features of foundation models include: (1) training on extremely large and diverse datasets to capture universal patterns, (2) effective architectures based on transformers that model complex dependencies, and (3) the ability to fine-tune or prompt the model for new tasks, transferring learned knowledge to improve performance on target applications [10]. This paradigm shift has enabled researchers to address biological questions that previously required specialized models and extensive retraining for each new task [9].

Table 1: Quantitative Performance of AI Applications in Bioinformatics

Application Domain	AI Model/Method	Performance Metrics	Significance
Protein Structure Prediction	AlphaFold	Median 0.96 Å on CASP14	Near-atomic accuracy [8]
Single-Cell Modeling	scFMs	AvgBIO 0.82	Robust cellular representation [8]
Protein Design	AI-based Design	Up to 92% success rate	High precision engineering [8]
Cancer Detection	AI Diagnostics	Area Under Curve (AUC) 0.93	Sensitive disease identification [8]
Genomic Data Curation	GPT-4	97% correct categorization	Efficient data extraction [11]
Quantitative Trait Loci Extraction	GPT-4	61% of marker-trait associations	Automated literature mining [11]

Technical Architectures and Methodological Innovations

Transformer Architectures in Biological Foundation Models

The transformer architecture serves as the fundamental backbone for most contemporary foundation models in bioinformatics [10]. Originally developed for natural language processing, transformers utilize attention mechanisms that allow the model to learn and weight relationships between any pair of input tokens [10]. In biological applications, this enables models to determine which genes, sequences, or structural elements are most informative for specific analyses. The self-attention mechanism is particularly powerful for biological data because it can capture long-range dependencies and complex interactions that simpler architectures might miss [10]. For genomic sequences, this means identifying regulatory elements that influence distant genes; for protein structures, it means recognizing how amino acid interactions dictate folding patterns [10]. The adaptability of transformer architectures has facilitated their application across diverse biological data types, including DNA sequences, protein structures, and single-cell transcriptomes [9] [10].

Single-Cell Foundation Models: A Case Study in Architectural Innovation

Single-cell foundation models (scFMs) exemplify the sophisticated architectural innovations driving AI evolution in bioinformatics [10] [12]. These models face the unique challenge of processing non-sequential data—gene expression profiles lack the inherent ordering of words in a sentence [10]. To address this, researchers have developed innovative tokenization strategies that convert single-cell data into sequences that transformers can process effectively [10]. Common approaches include ranking genes within each cell by expression levels or partitioning genes into bins based on expression values [10]. The input layers of scFMs typically comprise three components: gene embeddings (analogous to word embeddings), value embeddings representing expression levels, and positional embeddings that provide information about gene order or rank [12]. Models such as scBERT employ bidirectional encoder architectures, while scGPT uses decoder-inspired architectures with unidirectional masked self-attention [10]. These architectural variations represent different strategies for capturing the complex relationships in single-cell data, with each offering distinct advantages for specific analytical tasks [10] [12].

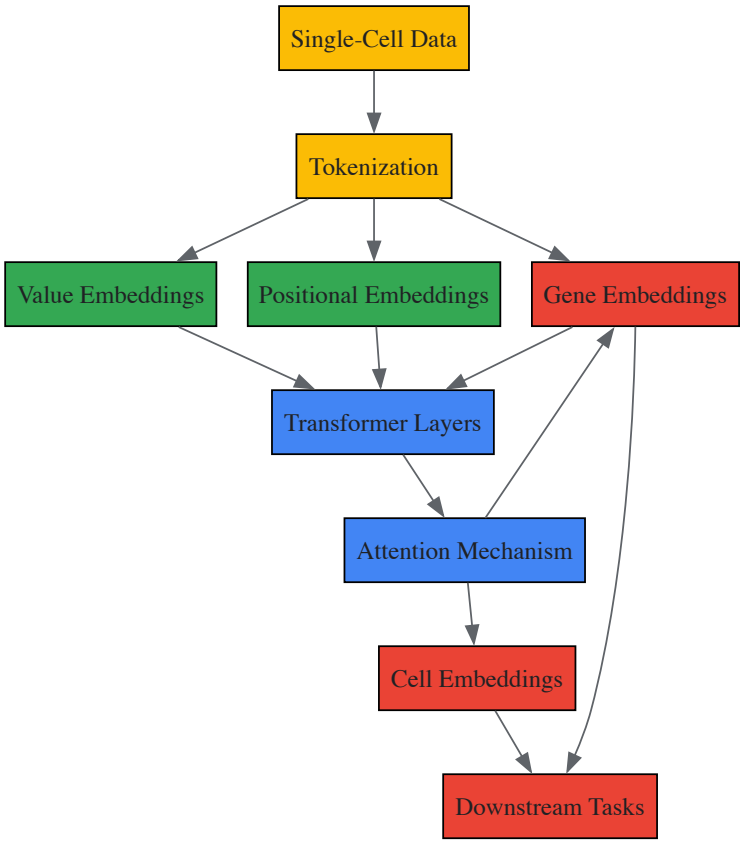

Diagram 1: Single-Cell Foundation Model Architecture

Evolutionary-Informed AI: Integrating Phylogenetic Knowledge

A recent innovation in biological AI involves explicitly incorporating evolutionary relationships into model architectures [13]. Traditional AI algorithms often struggle to analyze biological data through an evolutionary lens because they lack prior knowledge about phylogenetic trees and get confused by random patterns [13]. Researchers have addressed this limitation by developing neural networks that incorporate prior knowledge of species ancestry trees during training [13]. This approach classifies groups of four species into presumably correct ancestry trees, enabling the AI to identify patterns that have evolved throughout evolutionary history [13]. The method works not only for genetic sequence data but also for other data types, including images and structural patterns of biomolecules from various species [13]. This represents a significant advancement toward creating AI systems that reason about biological data in ways that align with established biological principles, moving beyond pattern recognition to embody deeper conceptual understanding of evolutionary processes [13].

Experimental Protocols and Benchmarking Frameworks

Benchmarking Single-Cell Foundation Models

Comprehensive benchmarking studies have emerged to evaluate the performance of scFMs under realistic biological scenarios [12]. These evaluations typically assess models across multiple task categories, including gene-level tasks (tissue specificity prediction, Gene Ontology term prediction) and cell-level tasks (dataset integration, cell type annotation, cancer cell identification, drug sensitivity prediction) [12]. Performance is measured using a combination of traditional metrics and novel biology-informed approaches such as scGraph-OntoRWR, which measures the consistency of cell type relationships captured by scFMs with prior biological knowledge [12]. The benchmarking pipeline involves extracting zero-shot gene and cell embeddings from pretrained models without additional fine-tuning, then evaluating these representations on diverse datasets that challenge models with real-world complexities like batch effects, novel cell types, and cross-tissue heterogeneity [12]. This rigorous evaluation framework provides crucial insights into model strengths and limitations, guiding researchers in selecting appropriate tools for specific biological questions.

Table 2: Single-Cell Foundation Model Performance Benchmarking

Model	Architecture Type	Batch Integration	Cell Type Annotation	Biological Relevance	Computational Efficiency
Geneformer	Transformer-based	High	Medium	High	Medium
scGPT	GPT-inspired	High	High	High	Low
scBERT	BERT-like	Medium	High	Medium	High

Model	Architecture Type	Batch Integration	Cell Type Annotation	Biological Relevance	Computational Efficiency
UCE	Custom encoder	Medium	Medium	High	Medium
scFoundation	Transformer-based	High	High	High	Low
LangCell	Language model-inspired	High	Medium	Medium	Medium

Evaluating Large Language Models for Biocuration

The application of large language models (LLMs) to biological curation tasks represents another important domain where rigorous experimental protocols have been developed [11]. These protocols typically involve comparing LLM performance against human curators for specific tasks such as categorizing manuscripts, extracting traits, and identifying marker-trait associations [11]. In one representative study, researchers used retrieval-augmented generation (RAG) to enhance GPT models with domain-specific knowledge, then evaluated their performance on curating wheat and barley genetics literature [11]. The experimental workflow involved parsing scientific PDFs, splitting text into overlapping chunks, generating vector embeddings, and querying the system with biologically relevant prompts [11]. Performance was assessed using precision metrics comparing LLM extractions against manual curation by PhD-level biologists [11]. This approach demonstrated that GPT-4 could correctly categorize manuscripts 97% of the time, extract 80% of traits, and identify 61% of marker-trait associations, highlighting both the potential and limitations of current LLMs for biological curation tasks [11].

Research Reagent Solutions: Essential Materials for scFM Research

Table 3: Essential Research Reagents and Resources for Single-Cell Foundation Model Development

Resource Category	Specific Examples	Function and Application
Data Repositories	CZ CELLxGENE, Human Cell Atlas, NCBI GEO, PanglaoDB	Provide standardized, annotated single-cell datasets for model training [10]
Processing Tools	Scipy, Scikit-learn, Scanpy	Enable data preprocessing, normalization, and quality control [12]
Model Frameworks	PyTorch, TensorFlow, Hugging Face Transformers	Provide foundational architectures for developing and training scFMs [10] [12]
Benchmarking Platforms	Custom evaluation pipelines, scGraph-OntoRWR metric	Enable standardized performance assessment and comparison [12]
Specialized Compute	GPU clusters, Cloud computing resources	Handle computational intensity of training and fine-tuning large models [10]

Applications and Impact Across Biological Domains

Genomic Sequence Analysis and Interpretation

Foundation models have revolutionized genomic sequence analysis by enabling researchers to identify evolutionarily conserved regions, mutation patterns, and critical functional domains directly from sequence information [8]. Language-inspired models treat DNA sequences as textual data, applying transformer architectures to predict regulatory elements, transcription factor binding sites, and variant effects [8] [9]. These approaches have demonstrated remarkable success in connecting sequence variations to functional consequences, providing insights into disease mechanisms and potential therapeutic targets [8]. The application of foundation models to genomics has moved beyond simple pattern recognition to enable predictive modeling of complex genotype-phenotype relationships, accelerating discovery in functional genomics and personalized medicine [8] [9].

Protein Structure Prediction and Design

The most celebrated success of AI in bioinformatics remains the breakthrough in protein structure prediction achieved by AlphaFold and related systems [8]. These models demonstrated median accuracy of 0.96 Å on the CASP14 assessment, approaching near-atomic resolution directly from amino acid sequences [8]. This achievement represented a decades-old challenge in structural biology and has fundamentally changed how researchers approach protein characterization and engineering [8]. Beyond structure prediction, foundation models have enabled innovative approaches to protein design, with AI-driven methods achieving success rates up to 92% for designing proteins with specific functions or properties [8]. These capabilities are accelerating therapeutic development, enzyme engineering, and the creation of novel biomaterials, demonstrating how foundation models can transition from analytical tools to generative platforms for biological innovation [8] [9].

Drug Discovery and Development

Foundation models are transforming drug discovery by enabling more efficient target identification, compound screening, and toxicity prediction [9]. These models integrate diverse data types—including genomic sequences, protein structures, chemical properties, and clinical outcomes—to identify promising therapeutic candidates and predict their behavior in biological systems [9]. The ability to pretrain on massive unlabeled datasets then fine-tune for specific discovery tasks has proven particularly valuable in drug development, where labeled data is often limited and expensive to acquire [9]. Foundation models can identify novel drug targets, predict drug-drug interactions, and optimize compound properties, significantly reducing the time and cost associated with traditional discovery pipelines [9]. As these models continue to improve, they promise to accelerate the development of personalized therapeutics tailored to individual genetic profiles and disease characteristics [8] [9].

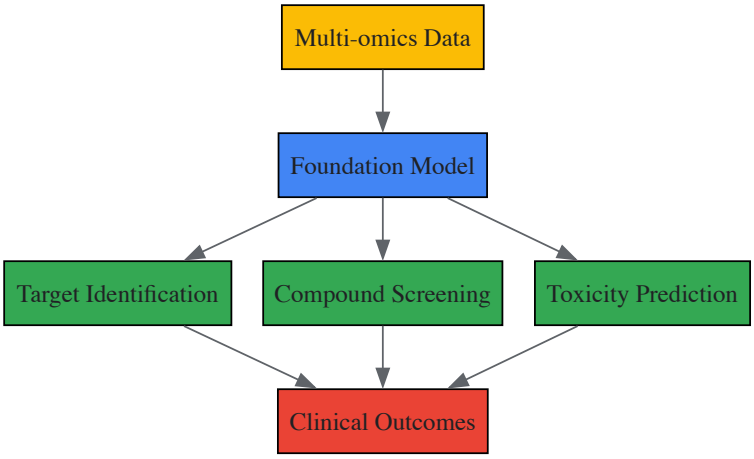

Diagram 2: Foundation Models in Drug Discovery Workflow

Challenges and Future Directions

Technical and Interpretability Limitations

Despite remarkable progress, foundation models in bioinformatics face significant technical challenges that limit their widespread adoption [8] [10] [14]. Massive and heterogeneous biological datasets frequently contain inherent noise, biases, and class imbalance issues that can compromise model performance and generalizability [8]. Biological sequences, particularly from higher organisms, often exceed the context windows of current transformer architectures, complicating effective modeling of long-range dependencies [8]. Perhaps most critically, explainability and reproducibility of AI models face heightened scrutiny from biological and medical communities [8] [14]. Black-box AI models raise concerns about decision transparency and user confidence, driving increased demand for explainable AI (XAI) approaches that can provide biological insights into model predictions [14]. Current research focuses on developing interpretation methods that reveal which features and patterns drive model decisions, enabling researchers to validate predictions against biological knowledge and generate testable hypotheses [14].

Data Quality and Integration Challenges

The effectiveness of foundation models depends critically on the quality and diversity of their training data [10]. Single-cell genomics exemplifies these challenges, as models must contend with datasets exhibiting varying sequencing depth, batch effects, technical noise, and inconsistent processing steps across different studies and platforms [10]. Assembling high-quality, non-redundant datasets for pretraining requires careful selection, filtering, and balancing of dataset compositions [10]. Furthermore, integrating multimodal data—such as combining genomic sequences, protein structures, medical imaging, and clinical text—presents additional complexities related to data alignment, representation learning, and cross-modal inference [8] [10]. Future advancements will require not only improved model architectures but also better data standardization, curation practices, and integration frameworks that preserve biological signals while mitigating technical artifacts [10].

Ethical Considerations and Future Trajectories

The expanding capabilities of foundation models in bioinformatics raise important ethical considerations, particularly regarding privacy when handling sensitive genomic and clinical patient data [8]. Establishing rigorous standards and frameworks for responsible AI deployment is essential as these models become integrated into healthcare and research settings [8]. Looking forward, researchers anticipate several promising directions for the field, including large-scale data mining across expanded biological corpora, improved cross-domain model generalization, innovations in drug design and personalized medicine, and the establishment of more open and collaborative research ecosystems [8] [15]. The integration of evolutionary perspectives represents another exciting frontier, enabling models to reason about biological data through phylogenetic relationships rather than just statistical patterns [13]. As these advancements mature, foundation models will increasingly serve not merely as analytical tools but as collaborative partners in biological discovery, helping researchers formulate hypotheses, design experiments, and interpret results across the vast complexity of living systems [8] [9] [10].

Foundation Models (FMs), large-scale machine learning models pre-trained on extensive datasets, are revolutionizing data interpretation across diverse scientific fields through self-supervised learning [10]. In bioinformatics, these models demonstrate remarkable proficiency in managing large-scale, unlabeled biological datasets, effectively addressing historical challenges related to data integration and analysis [9]. The adaptability of FMs enables their application to various downstream tasks in computational biology, consistently achieving high accuracy in representing complex biological entities and processes [9] [6].

This technical guide examines the four core architectural types of foundation models—Language, Vision, Graph, and Multimodal—within the specific context of bioinformatics research. We provide a comprehensive analysis of each architecture's fundamental principles, representative models, biological applications, and experimental methodologies. The content is structured to serve researchers, scientists, and drug development professionals seeking to understand and implement these transformative technologies in their computational workflows.

Core Foundation Model Architectures in Bioinformatics

Foundation models in bioinformatics leverage different architectural paradigms to process various types of biological data. The table below summarizes the four core types, their data handling characteristics, primary biological applications, and specific model examples.

Table 1: Core Foundation Model Architectures in Bioinformatics

Architecture Type	Primary Data Handling	Key Biological Applications	Representative Models
Language FMs [9]	Sequential data (e.g., DNA, protein sequences) [6]	Genome annotation, variant effect prediction, regulatory element identification [6]	DNABERT, Geneformer, scGPT [6] [10]
Vision FMs [9]	Image-based data (e.g., medical scans, protein structures) [16]	Medical image analysis (X-ray, MRI), protein structure prediction, histopathology [6] [16]	AlphaFold, CONCH [6] [16]
Graph FMs [9]	Graph-structured data (e.g., molecular structures, interaction networks)	Drug discovery, molecular property prediction, knowledge graph reasoning [6]	TxGNN, Novae [6]
Multimodal FMs [9]	Integrated multiple data types (e.g., text + images, multi-omics) [16]	Spatial transcriptomics, clinical diagnostics (imaging + EHR), multi-omics integration [6] [16] [10]	Nicheformer, MLLM4TS [17] [6]

Language Foundation Models

Language FMs process biological sequences—such as DNA, RNA, and proteins—as textual data, where discrete biological elements (nucleotides, amino acids) are treated as tokens or "words" in a biological language [18]. These models typically employ transformer architectures with self-attention mechanisms to capture contextual relationships within sequences [10].

Key Architectures and Pre-training Strategies:

- **Encoder-based Models (e.g., BERT-style):** Use bidirectional attention to learn from the entire sequence context simultaneously. Pre-training often involves masked language modeling, where random tokens in the input sequence are masked and the model learns to predict them based on surrounding context [10] [18].
- **Decoder-based Models (e.g., GPT-style):** Employ unidirectional masked self-attention, iteratively predicting next elements in a sequence conditioned on previous elements. This approach is particularly effective for generative tasks [10].

Bioinformatics Applications and Experimental Protocols:

- **Variant Effect Prediction:** Models like Enformer are trained to predict the effects of non-coding DNA variants on gene expression by analyzing sequence context and long-range genomic interactions (up to 100kb) [6]. Experimental validation typically involves comparing predictions with functional genomics data from assays like ChIP-seq or STARR-seq.
- **Single-Cell Analysis:** scGPT is pre-trained on approximately 30 million single-cell transcriptomes to learn gene-gene interaction patterns [6]. For cell type annotation, the protocol involves:
  - **Data Preprocessing:** Normalize raw count data and tokenize gene expression values.
  - **Embedding Generation:** Feed tokenized cells to scGPT to obtain latent embeddings.
  - **Fine-tuning:** Add a classification head and fine-tune on a smaller, labeled dataset using cross-entropy loss.
  - **Validation:** Compare predicted cell types against manual annotations or marker genes.

Table 2: Key Research Reagents for Language FM Experiments

Reagent/Resource	Function in Experimental Protocol
Reference Genome (e.g., hg38) [19]	Standardized genomic coordinate system for alignment and variant calling.

Reagent/Resource	Function in Experimental Protocol
Pre-trained Model Weights [6] [18]	Foundation for transfer learning, eliminating need for expensive pre-training.
High-Performance Computing (HPC) Cluster [19]	Provides computational capacity for training and fine-tuning large models.
Single-Cell Datasets (e.g., CZ CELLxGENE) [10]	Curated, annotated data for model fine-tuning and validation.
Containerized Software (Docker/Singularity) [19]	Ensures computational reproducibility across different environments.

Vision Foundation Models

Vision FMs process and interpret image-based biological data, including medical scans, protein structures, and cellular imagery. These models leverage convolutional neural networks and vision transformers (ViTs) to extract hierarchical visual features [16].

Key Architectures:

- **Vision Encoders:** Models like CLIP (Contrastive Language-Image Pre-training) use contrastive learning to align visual data with textual descriptions, creating a shared representation space [16].
- **Protein Structure Prediction:** AlphaFold employs novel neural architectures to predict 3D protein structures from amino acid sequences with near-experimental accuracy [6].

Bioinformatics Applications and Experimental Protocols:

- **Medical Image Analysis:** In radiology, MLLMs integrate diverse imaging modalities (CT, MRI, X-ray) with textual clinical data for tasks like automated radiology report generation and visual question answering [16]. The standard workflow involves:
  - **Image Preprocessing:** Standardize image dimensions, orientation, and intensity values.
  - **Feature Extraction:** Process images through a pre-trained vision encoder (e.g., ViT) to obtain visual embeddings.
  - **Multimodal Fusion:** Align visual embeddings with text embeddings using a multimodal connector.
  - **Task-Specific Heads:** Generate diagnostic reports or answers using the fused representations.
- **Time-Series Analysis:** The MLLM4TS framework converts multivariate time-series data into color-coded line-plot images, enabling anomaly detection and forecasting through visual pattern recognition [17].

Graph Foundation Models

Graph FMs specialize in processing graph-structured biological data, including molecular structures, protein-protein interaction networks, and knowledge graphs. These models use graph neural networks (GNNs) to capture topological relationships between biological entities [6].

Bioinformatics Applications and Experimental Protocols:

- **Drug Discovery:** TxGNN is trained on medical knowledge graphs to identify potential new purposes for existing drugs by traversing relationships between compounds, diseases, and biological pathways [6]. The experimental protocol involves:
  - **Knowledge Graph Construction:** Integrate data from sources like DrugBank, STITCH, and clinical trials.
  - **Message Passing:** Use graph attention networks to propagate information between connected nodes.
  - **Link Prediction:** Score potential drug-disease pairs using learned node embeddings.
  - **Validation:** Compare predictions against known off-label uses or conduct retrospective clinical validation.
- **Satial Biology Analysis:** Novae uses graph-based learning to correct for batch effects in spatial transcriptomics data, enabling more accurate comparison of spatial domains across different samples and experiments [6].

Multimodal Foundation Models

Multimodal FMs represent the most advanced architecture, capable of processing and integrating multiple data types simultaneously. These models align representations from different modalities into a shared semantic space, enabling complex cross-modal reasoning [16].

Key Architectural Components:

- **Modality-Specific Encoders:** Transform raw data from each modality (image, text, etc.) into meaningful representations [16].
- **Multimodal Connector:** Bridges the modality gap by mapping non-textual data to representations interpretable by Large Language Models (LLMs). Connector types include:
  - **Projection-based:** Employ Multi-Layer Perceptrons (MLPs) to transform visual data to align with language representations [16].
  - **Query-based:** Utilize trainable 'query tokens' to extract salient visual details [16].
  - **Fusion-based:** Facilitate feature-level integration through cross-attention mechanisms [16].
- **Pre-trained LLM Backbone:** Serves as the cognitive engine, providing reasoning capabilities without requiring additional fine-tuning for multimodal inputs [16].

Training Strategies: Multimodal FMs are typically developed through three sequential stages [16]:

- **Pre-training:** The multimodal connector learns to align visual and textual representations using autoregressive captioning on image-text pairs.
- **Instruction Tuning:** The model is fine-tuned with diverse natural language instructions and multimodal inputs to reliably follow complex directives.
- **Alignment Tuning:** The model's outputs are optimized to better reflect human preferences through reinforcement learning from human feedback.

Bioinformatics Applications and Experimental Protocols:

- **Spatial Transcriptomics:** Nicheformer is trained on both dissociated single-cell transcriptomics data (>57 million cells) and spatially-resolved transcriptomics data (>53 million cells) to make context-specific predictions about cellular microenvironments [6]. The protocol involves:
  - **Data Alignment:** Map dissociated cells to spatial coordinates using transfer learning.
  - **Contextual Prediction:** Infer spatial organization and cell-cell communication patterns.
  - **Biological Validation:** Validate predictions using techniques like multiplexed FISH or immunohistochemistry.
- **Clinical Diagnostics:** MLLMs in radiology integrate imaging data with clinical notes and laboratory results to generate comprehensive diagnostic reports [16]. Implementation requires robust validation against clinician interpretations and monitoring for hallucinated findings.

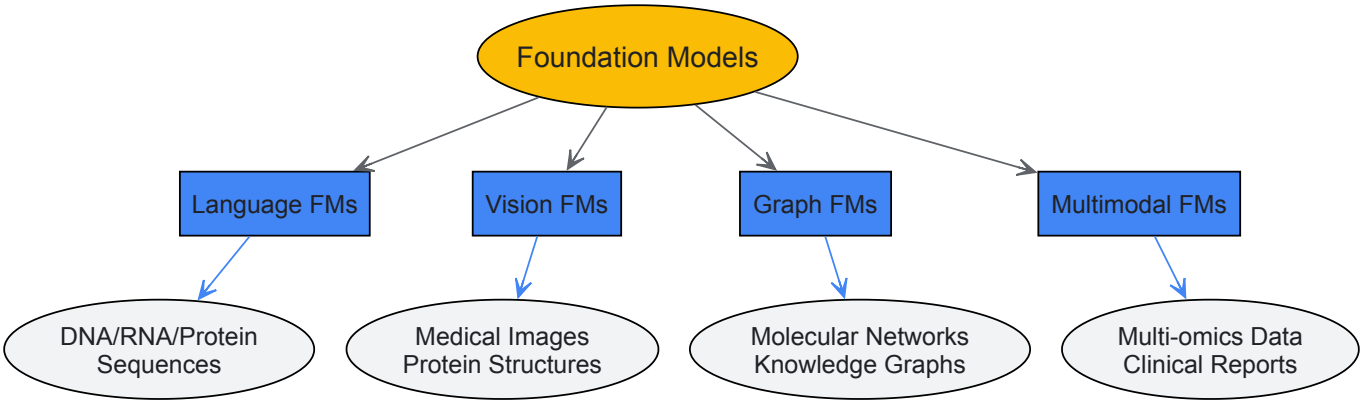

Foundation Models Classification in Bioinformatics

Implementation and Workflow Management

Successfully implementing foundation models in bioinformatics requires robust computational infrastructure and standardized workflow management. Nextflow and nf-core provide a critical framework for creating reproducible, scalable, and portable analysis pipelines [20].

Essential Implementation Rules:

- **Strategic Planning:** Align computational infrastructure with institutional priorities and research needs, identifying available HPC resources and existing knowledge gaps [20].
- **Workflow Standardization:** Transition from ad hoc scripts to containerized Nextflow pipelines, ensuring reproducibility through version control and standardized file formats [19] [20].
- **Containerized Environments:** Encapsulate software in Docker or Singularity containers to maintain consistency across different computational environments [19].
- **Quality Assurance:** Implement comprehensive testing strategies at unit, integration, system, and end-to-end levels, validating pipelines against standard truth sets like GIAB for germline variant calling [19].
- **Community Engagement:** Facilitate continuous learning through seminars, hackathons, and collaborative platforms to build sustainable bioinformatics capacity [20].

Bioinformatics Analysis Workflow with Foundation Models

Future Perspectives and Challenges

Despite their transformative potential, foundation models in bioinformatics face several significant challenges that must be addressed for widespread clinical and research adoption [21] [16] [10].

Critical Challenges:

- **Data Quality and Availability:** Limited access to large-scale, high-quality multimodal datasets, coupled with issues of batch effects, technical noise, and inconsistent processing steps across studies [16] [10].
- **Computational Intensity:** Substantial computational resources required for training and fine-tuning, creating barriers for resource-limited institutions [10].
- **Interpretability and Transparency:** Difficulty in interpreting the biological relevance of latent embeddings and model representations, particularly in clinical settings where decision-making processes must be explainable [16] [10].
- **Hallucination and Reliability:** Risks of hallucinated findings in generative models, requiring robust validation frameworks and human oversight [16].
- **Security and Privacy:** Ensuring data security, patient privacy, and ethical use of sensitive clinical information, particularly when using cloud-based resources [21].

Future Directions:

- **Region-Grounded Reasoning:** Developing models that can link outputs to specific regions in images or genomic sequences to enhance interpretability [16].
- **Robust Foundation Models:** Creating larger models pre-trained on comprehensive medical datasets specifically designed for healthcare applications [16].
- **Integration Strategies:** Establishing safe and effective protocols for integrating FMs into clinical practice, including validation standards and regulatory frameworks [16].
- **Cross-Modal Generalization:** Enhancing model capabilities to transfer knowledge across different biological modalities and species [10].

As these challenges are systematically addressed, foundation models are poised to become indispensable tools in bioinformatics, enabling deeper insights into cellular function, disease mechanisms, and therapeutic development [10]. Their continued evolution will likely establish the theoretical and practical foundation for ongoing innovation in molecular biology and precision medicine [9].

Foundation Models (FMs) are revolutionizing bioinformatics by providing a unified framework to interpret complex biological systems. Trained on massive-scale datasets through self-supervised learning, these models learn fundamental principles that can be adapted to a wide range of downstream tasks with minimal fine-tuning. Their capacity to capture intricate patterns and relationships within biological data makes them particularly suited for the high-dimensional, sparse, and heterogeneous nature of modern biological datasets. This technical guide details the core biological data types—DNA, RNA, proteins, and single-cell multi-omics—that form the corpus for these powerful models, outlining the specific FMs developed for each, their operational mechanisms, and their transformative applications in biomedical research and drug development.

DNA and Genomic Foundation Models

Genomic FMs are trained on DNA sequence data to understand the regulatory language of the genome and predict the functional impact of genetic variation.

Core Data Characteristics and Pretraining

Genomic FMs process DNA sequences as strings of nucleotides (A, C, G, T). Tokenization typically involves dividing long sequences into shorter k-mers (contiguous sequences of k nucleotides), which serve as the basic input tokens analogous to words in a sentence [3]. These models are pretrained on large-scale datasets from public repositories like the NCBI Sequence Read Archive (SRA) using self-supervised objectives, most commonly Masked Language Modeling (MLM), where random tokens in the input sequence are masked and the model is trained to predict them based on context [3].

Representative Models and Methodologies

Table 1: Foundation Models for DNA Sequence Analysis

Model Name	Architecture	Pretraining Data	Key Functionalities
DNABERT [6]	Transformer Encoder (BERT)	DNA sequences	Predicts regulatory regions (promoters, transcription factor binding sites), splice sites, and non-coding variant effects.
Enformer [6]	Deep Learning (CNN + Transformer)	DNA sequences with long-range context	Predicts the effects of non-coding DNA on gene expression, incorporating long-range interactions up to 100 kilobases.
DeepSEA [6]	Deep Learning (CNN)	Non-coding genomic variants	Predicts the chromatin and epigenetic effects of non-coding genomic variants.

Experimental Protocol for Variant Effect Prediction

A standard protocol for using DNA FMs to predict the pathogenicity of non-coding variants involves:

- **Sequence Input:** Extract a reference DNA sequence window (e.g., 1 kb) centered on the variant of interest.
- **Variant Representation:** Generate an alternative sequence by introducing the specific nucleotide change.
- **Model Inference:** Pass both reference and alternative sequences through the FM (e.g., Enformer or DeepSEA) to obtain predictions for chromatin features (e.g., histone marks, transcription factor binding) or gene expression levels.
- **Effect Scoring:** Calculate a prediction difference score (e.g., Euclidean distance between the model's outputs for the two sequences). A larger difference indicates a higher predicted functional impact [6].

RNA and Transcriptomic Foundation Models

Transcriptomic FMs focus on gene expression data, particularly from single-cell RNA sequencing (scRNA-seq), to characterize cellular heterogeneity, states, and functions.

Core Data Characteristics and Pretraining

The input data is a cell-by-gene matrix of expression counts. A significant challenge is the non-sequential nature of this data; genes have no inherent order. To apply transformer architectures, models use various tokenization strategies:

- **Rank-based:** Genes are ordered by their expression level within each cell, and the top-k genes form the input sequence [10] [22].
- **Binning:** Genes are partitioned into bins based on their expression values [10].
- **Normalized Counts:** Some models simply use normalized counts without complex ranking [10]. Each gene is typically represented by a token embedding that combines a gene identifier and its expression value. Special tokens for cell identity or batch information can be prepended [10].

Representative Models and Methodologies

Table 2: Foundation Models for Single-Cell Transcriptomics

Model Name	Architecture	Pretraining Data	Key Functionalities
Geneformer [22] [6]	Transformer Encoder	~30-95 million single-cell transcriptomes	Predicts tissue-specific gene network dynamics, cell type annotation, response to perturbation.
scGPT [10] [22] [6]	Transformer Decoder (GPT-style)	~30 million cells from scRNA-seq, scATAC-seq, CITE-seq	Cell type annotation, gene network inference, multi-omics data integration, and generative modeling.
scVI [6] [23]	Variational Autoencoder (VAE)	Single-cell gene expression data	Dimensionality reduction, visualization, clustering, and differential expression analysis.

Experimental Protocol for Cell Type Annotation

A common application of transcriptomic FMs is annotating cell types in a new scRNA-seq dataset:

- **Data Preprocessing:** The query dataset is filtered and normalized. Gene expression profiles are converted into a sequence of tokens according to the FM's specific tokenization scheme (e.g., top 2,048 ranked genes for Geneformer) [22].
- **Latent Embedding:** The tokenized sequences are passed through the pretrained FM to generate a low-dimensional latent embedding for each cell.
- **Annotation:**
  - **Zero-shot:** The embedding's position is compared to a reference atlas of known cell types.
  - **Fine-tuning:** A small set of labeled cells from the query dataset is used to fine-tune a classifier on top of the FM's embeddings, which then predicts labels for all unlabeled cells [10] [22].

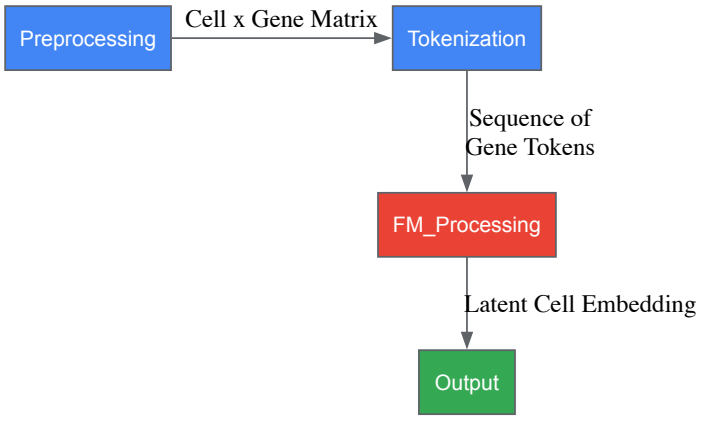

Figure 1: Workflow for processing single-cell RNA-seq data through a foundation model for tasks like cell type annotation.

Protein and Structural Foundation Models

Protein FMs interpret amino acid sequences and, in some cases, predict their three-dimensional structures, which is critical for understanding function and enabling drug design.

Core Data Characteristics and Pretraining

These models use amino acid sequences (strings of 20 standard letters) as primary input. Tokenization is straightforward, with each amino acid treated as a discrete token. Pretraining leverages large public databases like UniProt and can involve multiple self-supervised tasks, including MLM and predicting whether a protein is native-like [3]. Advanced models also incorporate physical and evolutionary constraints.

Representative Models and Methodologies

- **AlphaFold:** A deep learning model that uses neural networks to predict 3D protein structures from amino acid sequences with near-experimental accuracy. Its successor, AlphaFold3, extends this capability to complexes of proteins with other biomolecules [6].
- **ESM (Evolutionary Scale Modeling):** A series of transformer-based protein language models pretrained on millions of protein sequences. They generate informative representations that can be used for predicting structure, function, and the effects of mutations without an explicit multiple sequence alignment [3].

Experimental Protocol for Protein Structure Prediction

Using AlphaFold as a paradigm:

- **Input Preparation:** The target amino acid sequence is provided. A multiple sequence alignment (MSA) is generated for the target using databases of known protein sequences to infer evolutionary constraints.
- **Model Inference:** The target sequence and its MSA are fed into AlphaFold's neural network architecture.
- **Structure Generation:** The model outputs a 3D atomic coordinates file (PDB format) representing the most likely structure, along with per-residue confidence estimates (pLDDT scores) [6].

Single-Cell Multi-omics Foundation Models

This frontier in biological FMs aims to achieve a holistic view of the cell by jointly analyzing multiple molecular modalities measured within the same cell.

Core Data Characteristics and Pretraining

Single-cell multi-omics data encompasses simultaneously measured transcriptomic (RNA), epigenomic (e.g., scATAC-seq for chromatin accessibility), and proteomic (e.g., antibody-derived tags, ADTs) profiles [24] [25]. The key challenge for FMs is the "diagonal integration" of these distinct feature sets. Models are trained on datasets from technologies like CITE-seq (RNA + protein) and SHARE-seq (RNA + chromatin accessibility) [24] [25]. They must handle weak or limited correlations between some modalities (e.g., mRNA levels and protein abundance) and data sparsity [24].

Representative Models and Integration Strategies

Table 3: Foundation Models and Deep Learning Approaches for Single-Cell Multi-omics Integration

Model Name	Architecture	Key Integration Strategy	Modalities Handled
scMODAL [24]	Neural Networks + Generative Adversarial Networks (GANs)	Aligns cell embeddings using known feature links and adversarial learning.	scRNA-seq, scATAC-seq, Proteomics
scGPT [10] [25]	Transformer	Uses modality-specific tokens and a unified transformer to process and integrate data.	scRNA-seq, scATAC-seq, CITE-seq, Spatial
scMaui [23]	Variational Autoencoder (Product-of-Experts)	Combines marginal distributions of different modalities into a joint latent representation. Handles batch effects and missing data.	Multiple, flexible assays
Nicheformer [6]	Transformer	Trained on both dissociated and spatial transcriptomics data to contextualize cells within their tissue microenvironment.	scRNA-seq, Spatial Transcriptomics

Experimental Protocol for Multi-omics Data Integration

A standard workflow using a model like scMODAL or scMaui involves:

- **Input Data Preparation:** Provide cell-by-feature matrices for each modality (e.g., `X_rna`, `X_adt`). Compile prior knowledge of linked features (e.g., a gene and its protein product) into paired matrices [24].
- **Model Training:** The FM (e.g., an encoder network) is trained to project all cells into a shared latent space. Training employs several techniques:
  - **Adversarial Learning:** A discriminator network ensures the latent distributions of different modalities are indistinguishable, removing unwanted technical variation [24].
  - **Anchor-Based Guidance:** Mutual nearest neighbors (MNN) calculated on linked features are used to pull corresponding cells together in the latent space [24].
  - **Topology Preservation:** Regularization terms ensure the intrinsic geometric structure of each dataset is preserved [24].
- **Downstream Analysis:** The output is a unified, low-dimensional embedding of all cells. This embedding is used for clustering to identify novel cell states, trajectory inference, and cross-modality feature imputation [24] [23].

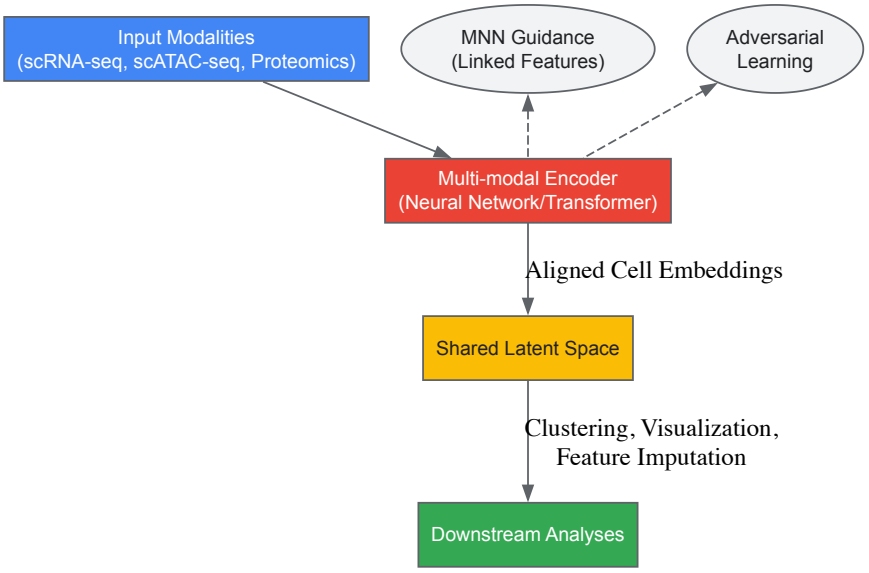

Figure 2: A generalized architecture for single-cell multi-omics integration using foundation models, showing key components like adversarial learning and mutual nearest neighbor (MNN) guidance.

Table 4: Essential Resources for Working with Biological Foundation Models

Resource Category	Item	Function & Utility
Data Repositories	NCBI GEO / SRA [10]	Primary archives for raw and processed genomic and transcriptomic sequencing data used for pretraining and fine-tuning.
	CZ CELLxGENE [10] [22]	Curated platform providing unified access to millions of annotated single-cell datasets.

Resource Category	Item	Function & Utility
Software & Tools	scvi-tools [6]	A Python package providing scalable implementations of probabilistic models for single-cell omics data, including scVI.
	Seurat [24] [23]	An R toolkit widely used for single-cell multi-omics analysis, often used as a baseline for comparison.
Pretrained Models	Hugging Face / Model Zoos	Platforms where pretrained FMs (e.g., scGPT, Geneformer) are often hosted for community use, enabling transfer learning.
Benchmarking Studies	Comparative Reviews [22]	Independent evaluations that provide holistic rankings of FM performance across diverse tasks, guiding model selection.

Foundation models represent a paradigm shift in computational biology, moving from task-specific algorithms to general-purpose models that learn the underlying language of DNA, RNA, proteins, and cellular systems. The integration of multiple data types, particularly through single-cell multi-omics FMs, is paving the way for a more comprehensive and predictive understanding of cellular function in health and disease. As these models continue to evolve, addressing challenges such as data quality, interpretability, and computational cost will be crucial. However, their current capacity to integrate heterogeneous data, generate novel biological hypotheses, and provide actionable insights for drug discovery already marks them as indispensable tools in the modern biologist's and drug developer's arsenal.

The emergence of foundation models has catalyzed a paradigm shift in bioinformatics, moving from task-specific models to general-purpose tools trained on massive, unlabeled datasets. This approach, known as pretraining, allows models to learn fundamental biological principles directly from data without manual annotation. By processing diverse biological sequences and structures through self-supervised learning, these models develop a comprehensive understanding of biological "grammar" that can be efficiently adapted to specialized tasks. The pretraining paradigm has demonstrated remarkable success across genomics, proteomics, transcriptomics, and bioimaging, enabling breakthroughs in disease prediction, drug discovery, and functional annotation. This technical guide examines the core methodologies, experimental protocols, and implementations of pretraining across biological domains, providing researchers with a comprehensive framework for leveraging unlabeled biological data.

Core Principles of Biological Pretraining

Conceptual Foundation and Biological Rationale

Pretraining in biological systems operates on the principle that biological sequences—whether DNA, RNA, proteins, or cellular images—contain inherent patterns and relationships that can be learned through self-supervised objectives without explicit labeling. This approach mirrors how large language models learn statistical regularities in human language, but applied to the "language of life." The fundamental insight is that the structure-function relationship in biology is encoded in sequences and images through evolutionary constraints and biophysical principles. By training on vast corpora of unlabeled data, models can internalize these constraints and develop representations that capture biologically meaningful features, from conserved protein motifs to regulatory DNA elements [26] [27].

The biological rationale stems from the observation that similar sequences often share similar functions across organisms, and that cellular images contain reproducible patterns corresponding to biological states. Pretraining enables models to capture these transferable principles, creating a foundation of biological knowledge that can be fine-tuned for specific applications. This is particularly valuable in biology where labeled data is scarce and expensive to generate, while unlabeled data exists in abundance through public repositories and large-scale sequencing initiatives [28] [10].

Key Architectural Components

Table 1: Core Architectural Components of Biological Foundation Models

Component	Function	Implementation Examples
Tokenization	Converts raw biological data into discrete units	k-mer splitting for DNA (e.g., "ATGCGA"), gene-level tokenization for expression data [26]
Embedding Layer	Maps tokens to dense vector representations	Learned embeddings for k-mers or gene identifiers [10]
Transformer Backbone	Captures long-range dependencies in sequences	Encoder-only (BERT-style), Decoder-only (GPT-style), or Encoder-Decoder architectures [26] [10]
Attention Mechanism	Models relationships between all positions in sequence	Multi-head self-attention with biological positional encodings [10]

Component	Function	Implementation Examples
Pretraining Head	Performs self-supervised learning objectives	Masked language modeling, next token prediction, contrastive learning [26]

The transformer architecture serves as the foundational backbone for most biological pretraining approaches, adapted to handle domain-specific challenges. Unlike natural language, biological sequences often lack inherent ordering—genes in a cell have no natural sequence, yet transformers require ordered input. Solutions include ranking genes by expression levels or using expression-based binning to create deterministic sequences from non-sequential data [10]. Positional encodings are adapted to represent biological context, such as genomic coordinates or cellular neighborhoods.

Domain-Specific Implementation Frameworks

Genomic Sequence Pretraining

Genomic pretraining treats DNA nucleotides as a vocabulary and sequences as texts to be understood. The standard workflow involves tokenizing DNA sequences into k-mers (typically 3-6 nucleotides), embedding these tokens, and processing through transformer layers. Pretraining tasks include masked language modeling where random nucleotide spans are masked and predicted based on context, or next nucleotide prediction in autoregressive frameworks [26]. Models like DNABERT and Nucleotide Transformer have demonstrated that this approach captures regulatory elements, conservation patterns, and functional annotations without supervision.

A critical advancement is the development of cross-species pretraining, where models trained on human genomic datasets transfer effectively to other organisms. Research has shown that models pretrained on human data then fine-tuned on diverse tissues and species achieve high prediction accuracy (Pearson correlation up to 0.8) while significantly reducing computational costs compared to training from scratch [28]. This demonstrates that fundamental genomic principles are conserved across species and can be transferred through pretraining.

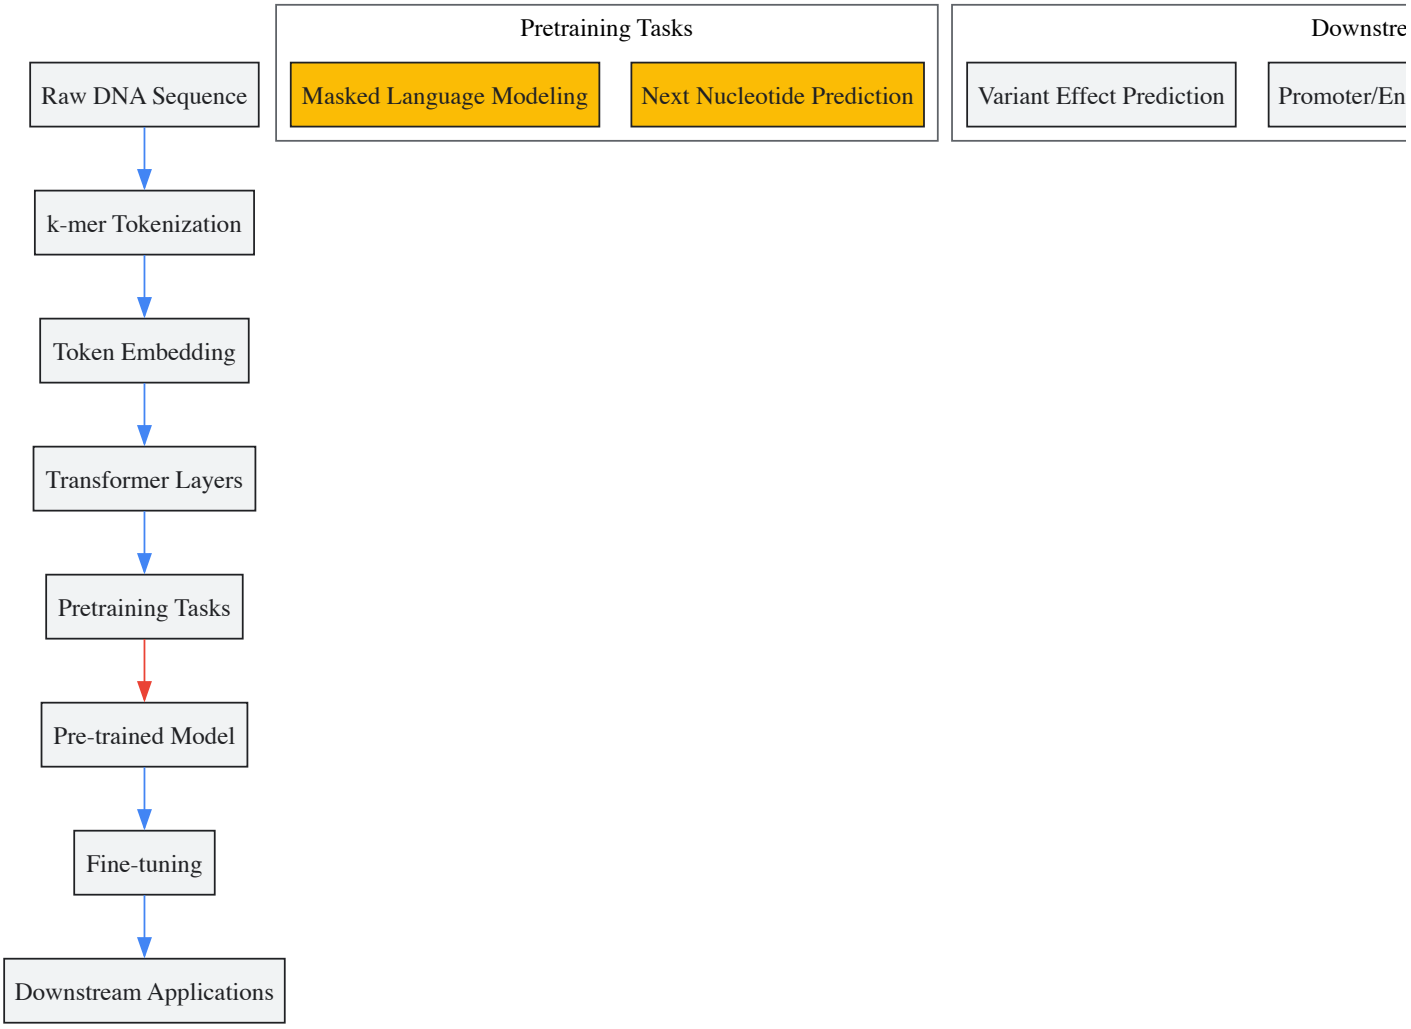

Single-Cell Multi-Omics Pretraining

Single-cell foundation models (scFMs) represent a transformative approach for analyzing cellular heterogeneity at scale. These models treat individual cells as "documents" and genes or genomic features as "words," learning to represent cellular states and regulatory relationships through pretraining on millions of single-cell transcriptomes [10]. The tokenization challenge is particularly acute in single-cell data, where gene expression profiles lack natural ordering. Solutions include ranking genes by expression value, binning expression levels, or using deterministic algorithms to create sequence structure from inherently non-sequential data.

Architecturally, scFMs employ both encoder-style models like scBERT for classification tasks and decoder-style models like scGPT for generation. Pretraining strategies include masked gene prediction, where random portions of the gene expression vector are masked and reconstructed based on context, enabling the model to learn gene-gene correlations and regulatory networks [10]. These models demonstrate remarkable transfer learning capabilities, adapting to novel cell types and conditions with minimal fine-tuning.

Table 2: Single-Cell Foundation Models and Their Applications

Model	Architecture	Pretraining Data	Key Applications
scBERT [10]	BERT-style Encoder	~30M single-cell transcriptomes	Cell type annotation, novel cell type discovery
scGPT [10]	GPT-style Decoder	Multi-omics single-cell data	Cell state generation, perturbation response prediction
GeneLLM [29]	Encoder-Decoder	cfDNA/cfRNA sequencing data	Preterm birth risk prediction, multi-omics integration

Bioimaging Pretraining

CytoImageNet exemplifies the pretraining paradigm for biological images, applying ImageNet-inspired approaches to microscopy data. This large-scale dataset contains ~890,000 weakly labeled images across 894 classes, compiled from 40 publicly available datasets [30]. Unlike natural images, microscopy images require specialized preprocessing, including channel normalization, artifact filtering, and multi-scale cropping to enhance diversity.

The pretraining workflow involves training convolutional networks (e.g., EfficientNetB0) to classify images using weak labels derived from metadata (e.g., organism, cell type, phenotype). Despite relatively low classification accuracy (11.32% validation), the learned features transfer effectively to downstream microscopy classification tasks, complementing ImageNet-pretrained features [30]. This demonstrates that domain-specific pretraining captures biologically relevant features not present in natural image distributions.

Experimental Protocols and Methodologies

Data Curation and Preprocessing

High-quality data curation is fundamental to successful biological pretraining. The CytoImageNet protocol illustrates comprehensive data processing: manual annotation of 65 datasets with dataset-specific processing to handle inconsistent metadata, standardization of file formats to PNG, conversion of RGB to grayscale where appropriate, normalization of fluorescent channels using 0.1th and 99.9th percentile pixel intensity, and quality control filtering to remove uniform images, binary masks, and dim/empty images [30].

For genomic sequences, the standard protocol involves:

- **Sequence acquisition** from reference genomes and sequencing repositories
- **Tokenization** using k-mer splitting with typical k values of 3-6
- **Sequence augmentation** through reverse complementation and random cropping
- **Quality filtering** to remove low-complexity and low-quality sequences

Single-cell data requires particularly careful preprocessing due to batch effects and technical noise. The standard workflow includes rigorous quality control, normalization using techniques like SCTransform, highly variable gene selection, and batch correction where necessary [10].

Pretraining Implementation

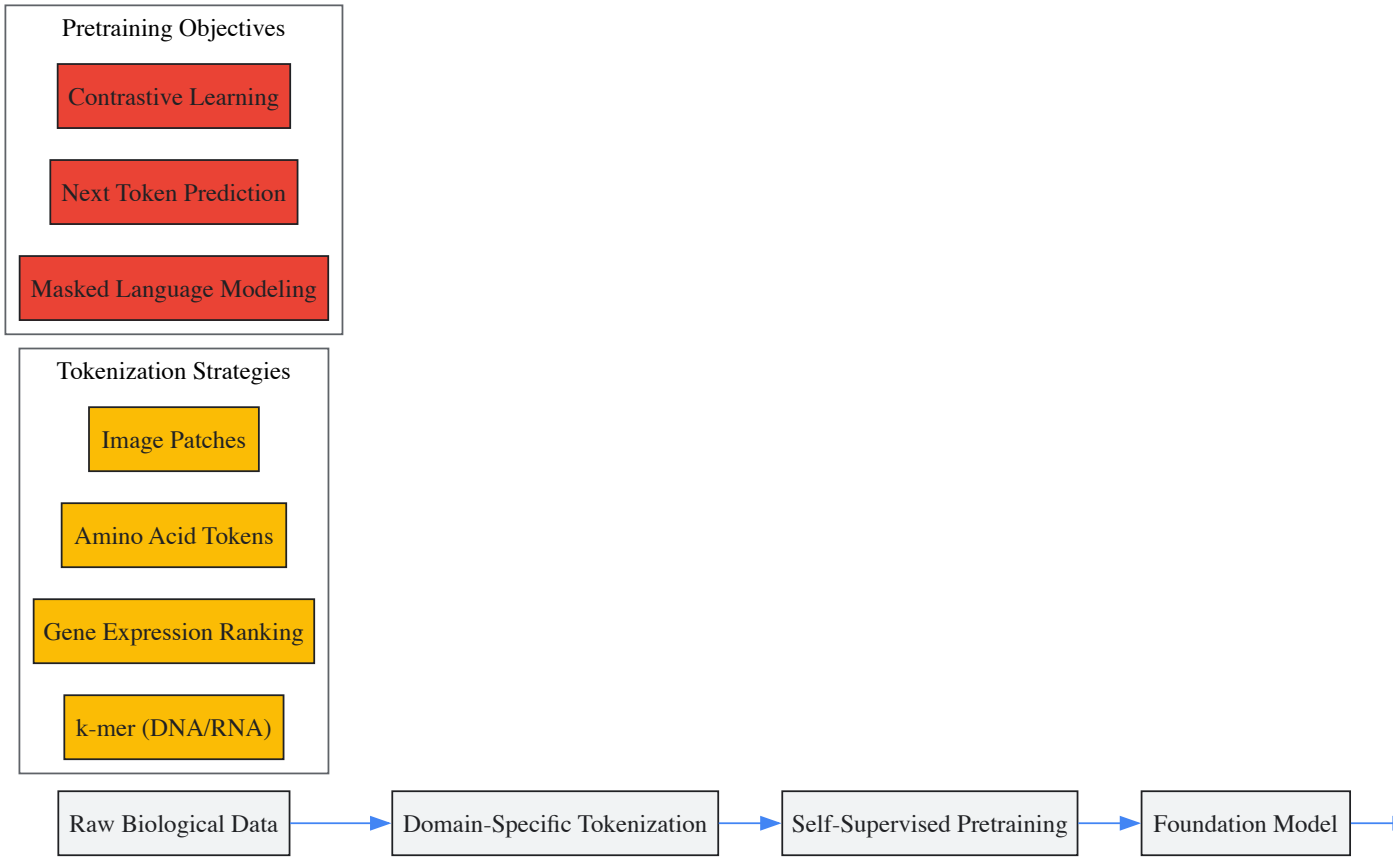

The technical implementation of pretraining follows a standardized workflow with domain-specific adaptations. The core process involves:

- **Model Architecture Selection**: Choosing appropriate transformer variants (encoder, decoder, or encoder-decoder) based on target tasks. Encoder models excel at classification and representation learning, while decoder models favor generation tasks.
- **Pretraining Objective Implementation**:
  - **Masked Language Modeling**: Randomly masking 15% of input tokens and training the model to reconstruct them
  - **Next Token Prediction**: Autoregressively predicting each token based on previous context
  - **Contrastive Objectives**: Learning representations by maximizing similarity between related biological sequences
- **Training Configuration**: Using large batch sizes (1024+ sequences), optimized learning rate schedules, and distributed training across multiple GPUs/TPUs. Training typically requires days to weeks on specialized hardware.
- **Validation and Checkpointing**: Monitoring performance on held-out validation sets and saving checkpoints for downstream fine-tuning.

The GeneLLM implementation for preterm birth prediction demonstrates a sophisticated multi-omics approach, processing cfDNA VCF files and cfRNA expression matrices through separate tokenization pipelines before integration in a transformer architecture [29]. This model achieved AUCs of 0.822 (cfDNA), 0.851 (cfRNA), and 0.890 (combined) for preterm birth prediction, significantly outperforming conventional machine learning approaches.

Transfer Learning and Fine-tuning

The true power of pretrained biological models emerges through transfer learning. The standard protocol involves:

- **Task-Specific Adaptation**: Replacing the pretraining head with task-specific layers (e.g., classification layers for disease prediction)
- **Progressive Unfreezing**: Gradually unfreezing layers during fine-tuning to prevent catastrophic forgetting
- **Differential Learning Rates**: Applying lower learning rates to earlier layers that capture general features and higher rates to task-specific layers

Cross-species transfer represents a particularly powerful application. Models pretrained on human genomic data successfully transfer to other organisms, achieving high accuracy with minimal fine-tuning. This approach dramatically reduces computational requirements compared to training from scratch for each new species or task [28].

Table 3: Research Reagent Solutions for Biological Pretraining

Resource Category	Specific Tools/Datasets	Primary Function	Access Information
Pretraining Datasets	CytoImageNet [30]	Bioimage pretraining	~890,000 images, 894 classes, available on Kaggle
Genomic Sequences	CZ CELLxGENE [10]	Single-cell transcriptomes	>100 million unique cells, standardized annotations
Model Architectures	DNABERT, Nucleotide Transformer [26]	Genomic sequence processing	Open-source implementations available
Training Frameworks	PyTorch, JAX, Hugging Face	Model development ecosystem	Open-source with biological extensions
Specialized Hardware	GPU/TPU clusters	Accelerated model training	Cloud providers and institutional resources

Advanced Applications and Validation Frameworks

Interpretability and Biological Discovery

Beyond predictive performance, pretrained models serve as discovery tools when combined with interpretability methods. Sparse autoencoders (SAEs) applied to protein language models have identified missing functional annotations in biological databases. For example, analysis of ESM-2 revealed a "Nudix box motif" feature that correctly identified annotations missing from Swiss-Prot [31]. Similarly, interpretability methods applied to genomic models have uncovered evolutionary relationships, including prophage regions and CRISPR-phage associations that reflect functional biological relationships rather than superficial sequence similarity [31].

The emerging methodology involves:

- **Feature Extraction**: Using SAEs to decompose model activations into interpretable directions
- **Concept Mapping**: Associating these directions with biological concepts through enrichment analysis
- **Experimental Validation**: Testing model-derived hypotheses through targeted experiments

This approach transforms black-box models into hypothesis generation engines, directly contributing to biological knowledge discovery.

Multi-Modal Integration

The most powerful biological pretraining approaches integrate multiple data modalities. The GeneLLM architecture for preterm birth prediction demonstrates this principle, processing both cfDNA variation profiles and cfRNA expression data through quantized representations that are combined in a transformer framework [29]. The integrated model significantly outperformed single-modality approaches, achieving an AUC of 0.890 compared to 0.822 (cfDNA only) and 0.851 (cfRNA only).

The technical implementation involves:

- **Modality-Specific Tokenization**: Converting each data type into appropriate token sequences
- **Cross-Attention Mechanisms**: Allowing information flow between modalities
- **Fusion Strategies**: Early (input-level), intermediate (layer-level), or late (prediction-level) integration

This multi-modal approach captures complementary biological signals, enabling more comprehensive modeling of complex phenotypes.

The pretraining paradigm represents a fundamental shift in biological computation, leveraging unlabeled data to build foundational knowledge that transfers across tasks and species. Through standardized methodologies for data curation, model architecture, and transfer learning, researchers can now develop powerful predictive models with dramatically reduced requirements for labeled data. As biological datasets continue to expand and model architectures evolve, pretrained foundation models will increasingly serve as essential tools for biological discovery, clinical translation, and therapeutic development. The integration of interpretability methods further transforms these models from black-box predictors to partners in scientific discovery, uncovering biological insights directly from data through carefully designed computational experiments.

From Sequence to Function: Methodologies and Applications of FMs in Biomedicine

The exponential growth of biological sequence data has necessitated a paradigm shift in bioinformatics, moving from specialized computational tools to general-purpose foundation models (FMs). These models, inspired by breakthroughs in natural language processing (NLP), leverage transformer architectures to develop a unified understanding of the "language of life" encoded in DNA, RNA, and proteins. This technical review examines the architecture, training methodologies, and experimental applications of biological FMs, with a focus on LucaOne, a pre-trained foundation model trained on nucleic acid and protein sequences from 169,861 species. We provide quantitative performance comparisons across diverse biological

tasks, detailed experimental protocols for benchmarking FM capabilities, and specialized visualization of core architectural principles. For researchers and drug development professionals, this review serves as both a technical reference and a strategic guide to leveraging FMs for decoding complex biological systems.

Biological information flows through three primary biopolymers—DNA, RNA, and proteins—each employing a linear sequence of molecular "letters" (4 nucleotides for DNA/RNA, 20 standard amino acids for proteins) that remarkably resembles human linguistic systems [32]. This parallel has motivated the application of natural language processing techniques to biological sequences. Foundation models represent the cutting edge of this convergence, utilizing semi-supervised learning on vast datasets to extract generalizable features that can be adapted to diverse downstream tasks with minimal fine-tuning [15].

Traditional computational methods often struggle to integrate information across DNA, RNA, and proteins, limiting comprehensive understanding of biological systems [32]. Single-modality models like DNABert2 [32] for nucleic acids and ESM2 [32] for proteins have demonstrated impressive results within their respective domains but fail to capture the interconnected nature of biological information flow. The emergence of unified foundational models like LucaOne [32] addresses this limitation through concurrent training on nucleic acid and protein sequences, enabling these models to inherently learn biological principles such as the central dogma of molecular biology without explicit instruction.

Unified Architectural Framework for Biological Sequences

Core Model Architecture

LucaOne implements an enhanced transformer encoder architecture specifically optimized for biological sequence processing [32]. The model incorporates several key modifications超越standard transformer designs:

- **Vocabulary and Tokenization**: Employs 39 unique tokens representing nucleotides and amino acids, with token-type encoding distinguishing nucleotides (assigned 0) from amino acids (assigned 1) [32].
- **Normalization and Positioning**: Utilizes pre-layer normalization instead of post-layer normalization to stabilize training of deep networks, and replaces absolute positional encoding with rotary position embedding for improved inference on longer sequences [32].
- **Scale and Parameters**: Comprises 20 transformer-encoder blocks with an embedding dimension of 2,560, totaling 1.8 billion parameters [32].

The model was pre-trained on massive datasets including RefSeq for nucleic acids (DNA and RNA) and UniProt, ColabFoldDB, and RCSB-PDB for protein sequences and structures [32]. This semi-supervised approach incorporated eight foundational sequence-based annotation categories alongside fundamental self-supervised masking tasks [32].

Visualization of Unified Model Architecture

The following diagram illustrates LucaOne's integrated processing approach for nucleic acid and protein sequences:

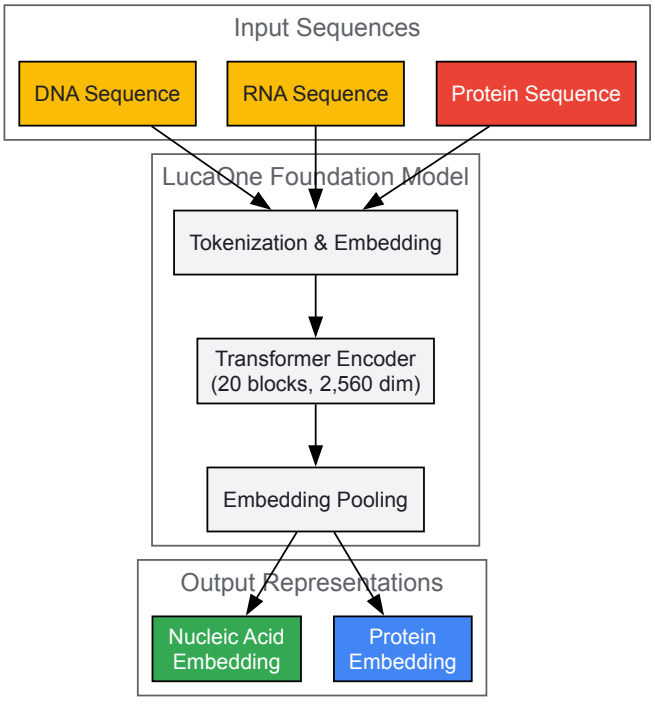

Unified Biological Sequence Processing

This unified architecture enables LucaOne to process and analyze data from nucleic acids and proteins simultaneously, facilitating extraction of complex patterns and relationships inherent in gene transcription and protein translation processes [32]. The model's ability to jointly represent these molecular modalities in a shared embedding space demonstrates emergent understanding of fundamental biological principles.

Experimental Validation and Performance Benchmarking

Central Dogma Recognition Task

A critical test for biological FMs is assessing their emergent understanding of the central dogma of molecular biology—the relationship between DNA sequences and their corresponding proteins [32]. To evaluate this capability, researchers designed an experimental task using DNA-protein matching pairs from NCBI RefSeq database with a 1:2 positive-to-negative sample ratio [32].

Experimental Protocol:

- **Dataset Construction**: Compiled DNA and protein matching pairs from RefSeq, with careful curation to ensure biological relevance [32].
- **Few-Shot Learning Validation**: Samples randomly allocated across training, validation, and testing sets in a ratio of 4:3:25 respectively to evaluate performance with limited training examples [32].
- **Downstream Network Architecture**: Implemented a simple classification network where LucaOne encoded nucleic acid and protein sequences into separate fixed embedding matrices, processed through pooling layers (max pooling or value-level attention pooling), concatenated, and passed through a dense layer for classification [32].

The following table summarizes quantitative performance comparisons across different modeling approaches:

Table 1: Performance Comparison on Central Dogma Recognition Task

Model Architecture	Pre-training Strategy	Key Performance Findings
One-hot + Transformer [32]	No pre-training	Unable to learn DNA-protein translation capability
Random initialization + Transformer [32]	No pre-training	Failed to acquire translation understanding
DNABert2 + ESM2-3B [32]	Separate nucleic acid & protein training	Substantially surpassed by unified LucaOne approach
LucaOne-Gene & LucaOne-Prot [32]	Independent nucleic acid & protein training	Inferior to unified training despite same architecture
LucaOne (unified) [32]	**Concurrent nucleic acid & protein training**	**Effectively learned DNA-protein correspondence with limited examples**

The findings demonstrate that modeling methods lacking pre-trained elements were unable to acquire DNA-protein translation capability, whereas LucaOne's unified embeddings effectively learned this relationship with limited training examples [32]. This suggests pre-trained foundational models provide additional information beyond specific task samples for biological computation tasks.

Embedding Space Analysis

Researchers utilized t-distributed stochastic neighbor embedding (t-SNE) to visualize embeddings from three distinct datasets: a nucleic acid dataset (S1) with sequences from 12 marine species, a protein dataset (S2) with sequences from 12 Pfam clans, and another protein dataset (S3) organizing sequences from the top 12 most prevalent Gene Ontology biological process terms [32].

Methodology:

- Compared LucaOne embeddings against MultiHot, DNABert2, and ESM2-3B embedding approaches [32].
- Evaluated embedding clustering density and biological relevance using standardized metrics [32].
- Examined correlation between nucleic acid and protein sequences of the same genes based on embedding proximity without explicit correspondence relationships during training [32].

Results demonstrated that sequences (nucleic acids and proteins) of the same gene exhibited convergence within the LucaOne embedding space, with more pronounced clustering compared to independently trained pre-trained models and sequence alignment methods [32]. This emergent property indicates the model developed an intrinsic understanding of biological relationships without explicit supervision.

Implementation Protocols for Foundation Model Applications

Workflow for Downstream Task Adaptation

The following diagram outlines a standardized protocol for applying foundation models like LucaOne to specific biological research questions:

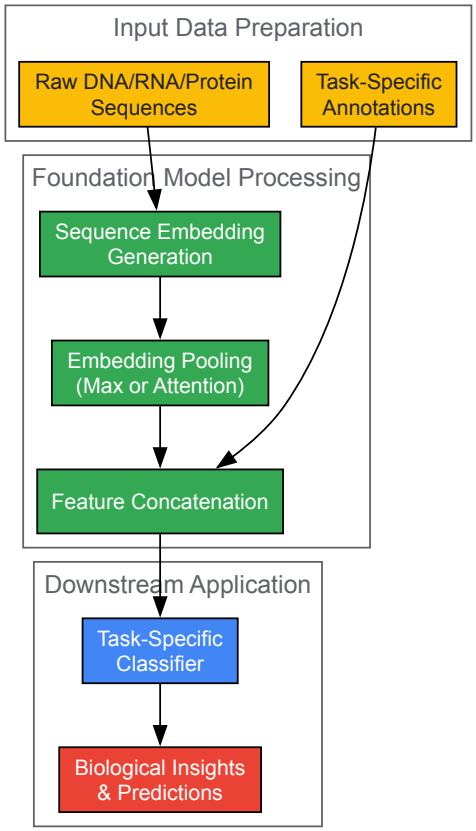

Foundation Model Application Workflow

Table 2: Foundation Model Research Reagent Solutions

Resource Category	Specific Tools & Databases	Function in FM Research
Sequence Databases [32]	RefSeq, UniProt, ColabFoldDB	Provide curated sequence data for pre-training and fine-tuning foundation models
Structural Databases [32]	RCSB-PDB, AlphaFold2 DB	Offer protein tertiary structure information for multi-modal training
Annotation Resources [32]	InterPro, Gene Ontology (GO)	Supply functional annotations for semi-supervised learning tasks
Analysis Platforms [33]	UniProt BLAST, Align tools	Enable sequence comparison and functional annotation expansion
Specialized Software [34]	Geneious Prime	Provides molecular biology and sequence analysis tools for validation
Workflow Systems [35]	Nextflow, Snakemake	Support reproducible analysis pipelines for large-scale FM applications
Variant Databases [36]	ClinVar, COSMIC, OncoKB	Enable clinical interpretation of genomic findings through integration

Future Directions and Implementation Challenges

As biological foundation models evolve, several critical challenges and opportunities emerge. Future development trajectories include improved multi-modal integration combining sequence, structure, and expression data; enhanced interpretability methods for understanding model decisions; and specialized architectures for particular biological domains [15]. The field must also address computational resource requirements, data privacy concerns for clinical applications, and standardized benchmarking approaches [35].

For precision medicine applications, foundation models face the additional challenge of integrating seamlessly with existing clinical workflows and laboratory information management systems (LIMS) [36]. Next-generation platforms are addressing these needs through automated data pipelines that connect sequencing outputs with clinical information, variant databases, and therapeutic guidelines while maintaining complete data provenance [36].

Foundation models represent a transformative approach to decoding the language of DNA, RNA, and proteins. By leveraging unified architectures trained on massive-scale biological datasets, models like LucaOne demonstrate emergent understanding of fundamental biological principles, including the central dogma of molecular biology. The experimental protocols and performance benchmarks outlined in this review provide

researchers with practical frameworks for applying these powerful tools to diverse biological questions. As the field advances, biological FMs promise to accelerate drug discovery, enhance diagnostic precision, and fundamentally expand our understanding of life's molecular machinery.

The field of bioinformatics is undergoing a paradigm shift with the adoption of foundation models (FMs), large-scale artificial intelligence (AI) systems trained on extensive datasets that can be adapted to a wide range of downstream tasks [9]. These models demonstrate notable proficiency in managing large-scale, unlabeled biological datasets, addressing the critical challenge of costly and labor-intensive experimental procedures [9]. Within this broader context of foundation models in bioinformatics, the specific domain of 3D molecular structure prediction has emerged as a particularly transformative application. Accurate prediction of molecular structures represents a fundamental challenge in computational biology and drug design, where understanding the precise spatial arrangement of atoms directly informs our comprehension of molecular function, interaction, and therapeutic potential.

Traditional drug development is a marathon process characterized by 10-15 year timelines, approximately $2 billion in operational costs, and a 90% failure rate in clinical trials [37]. AI-driven approaches are rewriting this narrative by analyzing vast datasets of genomic sequences, chemical libraries, and clinical records to dramatically accelerate discovery timelines [37]. The ability to reliably predict and generate 3D molecular architectures serves as the cornerstone for this transformation, enabling researchers to explore chemical space with unprecedented efficiency and precision. This whitepaper provides an in-depth technical examination of the methodologies, experimental protocols, and applications of AI-driven modeling of 3D molecular architectures, framed within the context of foundation models in bioinformatics.

Core Methodologies in AI-Driven 3D Molecular Modeling

Physics-Informed Generative Foundation Models

Recent advances have introduced sophisticated generative foundation models specifically designed for 3D molecular editing. MolEdit, a pre-trained multimodal molecular GenAI, exemplifies this approach by combining physics-informed and data-driven learning to effectively model the distribution of 3D molecular structures [38]. Unlike SMILES or graph-based approaches, MolEdit leverages 3D atomic coordinates as a unified representation of both isomeric and conformational variations, eliminating ambiguities inherent in discrete representations while offering a direct route to modeling continuous chemical and conformational spaces [38].

A critical innovation in MolEdit is its handling of molecular symmetries through **Group-Optimized (GO) labeling**, which reformulates the training labels of standard denoising diffusion probabilistic models (DDPMs) to respect translational, rotational, and permutation symmetries [38]. This non-invasive strategy effectively reduces degeneracy caused by symmetries while remaining model-agnostic and incurring minimal overhead. Additionally, the model employs an **Asynchronous Multimodal Diffusion (AMD) schedule** that decouples the diffusion of molecular constituents from that of atomic positions, resulting in a two-stage generation strategy that probabilistically decomposes the discrete and continuous variables in molecules [38].

To address the challenge of physically implausible structures, MolEdit incorporates a **Boltzmann-Gaussian Mixture (BGM) kernel** that aligns the diffusion process with physical constraints such as force-field energies [38]. This integration resembles preference alignment in other generative AI systems but uses a physics critic for guiding molecular structures, adding a Boltzmann factor to the forward diffusion transitions that emphasizes physical criteria like free energy.

Table 1: Comparative Analysis of 3D Molecular Generation Platforms

Platform/Model	Core Methodology	Molecular Representation	Key Innovation	Applicable Scope
3D-Scaffold Framework [39]	Reinforcement Learning with 3D Scaffold	3D atomic coordinates	Atom-by-atom placement guided by reward function	Drug-like molecules with high binding affinity
MolEdit [38]	Physics-Informed Diffusion Model	3D atomic coordinates	Group-optimized labeling & asynchronous multimodal diffusion	Small molecules to large bioactive compounds (up to 100 heavy atoms)
Transformer-based AP Prediction [40]	Deep Learning with Self-Attention	Amino acid sequences	Aggregation propensity prediction from sequence data	Decapeptides with tunable aggregation properties

Reinforcement Learning and 3D Scaffold Frameworks

Complementary to diffusion models, reinforcement learning approaches have demonstrated significant capability in molecular design. Pacific Northwest National Laboratory's 3D-Scaffold framework utilizes a novel deep learning model that efficiently generates 3D coordinates of molecules starting from a given scaffold while preserving its structure [39]. This technology iteratively adds atoms to a molecule under construction, guided by a reward function that optimally tunes the molecular structure for high-binding affinity and synthetic accessibility [39].

The model was trained using 5,000 molecules from the Food and Drug Administration and Enamine databases, consisting of six unique scaffolds that yielded a 3D scaffold framework capable of placing atoms one at a time in 3D space with high fidelity [39]. This approach addresses the limitations of traditional drug discovery methods by accelerating the discovery process and enhancing the precision of drug-target interactions, achieving atomic-level precision in generating drug-like molecules evaluated through binding affinity and interaction forces scores [39].

Multi-Scale Modeling for Peptide Aggregation Prediction

For peptide and protein structures, AI-driven approaches must span multiple spatial-temporal scales. Recent research has combined deep learning strategies—including genetic algorithms and reinforcement learning—to generate decapeptides with tunable aggregation propensities [40]. In this methodology, coarse-grained molecular dynamics (CGMD) simulations evaluate solvent-accessible surface area to define aggregation propensity (AP), with a Transformer-based prediction model achieving high accuracy in AP prediction with only a 6% error rate [40].

The aggregation propensity is defined as the ratio of the solvent-accessible surface area (SASA) of peptides aggregate before and after CGMD simulation, with AP = 1.5 serving as the threshold to distinguish between high aggregation propensity peptides (HAPP) and low aggregation propensity peptides (LAPP) [40]. This approach demonstrates how integrating AI with molecular modeling can guide the rational design of peptides with controlled assembly behavior, providing a scalable strategy for applications in biotechnology and medicine.

Experimental Protocols and Workflows

Workflow for De Novo Molecular Generation with MolEdit

The following diagram illustrates the complete workflow for the MolEdit foundation model, from pre-training to downstream applications:

Diagram 1: MolEdit Foundation Model Workflow

The experimental protocol for MolEdit begins with **multimodal pre-training** over large amounts of available molecular data, endowing the scalable model with emerging capabilities [38]. This pre-training uses a decomposition of molecular representations, inspired by the success of Text2Image GenAI, and is performed through an unsupervised objective for molecular AI focused on 3D molecular reconstruction [38].

For symmetry handling, the **group-optimized (GO) labeling** strategy reformulates training labels of standard DDPMs during pre-processing to respect translational, rotational, and permutation symmetries [38]. This process is non-invasive, introducing merely a plug-and-play modification to the training protocol that can be executed efficiently in practice without architectural changes.

Physics integration is achieved through the **Boltzmann-Gaussian Mixture (BGM) kernel**, which incorporates a Boltzmann factor to the forward diffusion transitions, emphasizing physical criteria like free energy [38]. This alignment with physical constraints such as force-field energies helps suppress undesired model hallucinations and prioritizes more realistic configurations during training and inference [38].

Protocol for Aggregation Propensity Prediction in Peptides

For peptide aggregation prediction, researchers have developed a comprehensive protocol combining computational simulations with AI-based prediction:

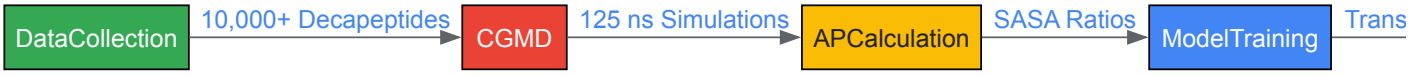

Diagram 2: Peptide Aggregation Prediction Pipeline

The experimental protocol begins with **data collection and preparation**, compiling a dataset of over 10,000 decapeptides for CGMD simulation [40]. In the initial simulation state, peptides are randomly distributed in the simulation box with a minimum inter-peptide distance constraint of 0.4 nm to prevent pre-aggregation, ensuring the initial SASA of peptides represents the maximum value [40].

Coarse-grained molecular dynamics simulations are then performed for 125 ns using the Martini CGMD strategy, which has been demonstrated as sufficient to identify differences in AP between HAPP and LAPP for decapeptide systems [40]. During these simulations, the SASA of the peptide system is tracked throughout the simulation timeframe.

Aggregation propensity calculation is performed as AP = SASA*initial / SASA*final, where peptides with high aggregation propensity will gradually approach and aggregate during the simulation, reducing the SASA of the peptide system and increasing the AP from 1.0 to 2.0 [40]. Peptides with AP > 1.5 are classified as HAPP, while those with AP < 1.5 are classified as LAPP [40].

Model training employs a deep learning model based on Transformer blocks with a self-attention mechanism, taking index encoding of amino acid sequences as input to predict AP [40]. The model uses existing MD simulation results as training datasets, stored in the weight parameters of the neural network, achieving a mean square error of approximately 0.004 on the validation set with no significant overfitting or outliers [40].

For **sequence optimization**, genetic algorithms are employed starting with 1000 randomly sampled initial sequences, allowing them to freely crossover with a limited mutation rate of 1% (meaning each residue has a 1% probability of being replaced by another residue) [40]. After 500 iterations, this approach increases average AP from 1.76 to 2.15, demonstrating effective optimization of aggregation propensity [40].

Performance Metrics and Quantitative Evaluation

Performance Benchmarks Across Molecular Generation Platforms

Table 2: Quantitative Performance Metrics of AI-Driven Structure Prediction Methods

Model/Method	Primary Task	Accuracy/Validity Metric	Computational Efficiency	Key Limitations
AlphaFold 3 [37]	Protein-ligand interaction modeling	High accuracy for protein structures	Predicts structures in hours vs. months/years experimentally	Limited accessibility and comprehensive validation across all biomolecules
3D-Scaffold Framework [39]	Drug candidate generation	High fidelity in 3D coordinate generation	Accelerates discovery timeline significantly	Limited to scaffold-based generation approach
Transformer AP Prediction [40]	Peptide aggregation prediction	6% error rate in AP prediction	Reduces assessment from hours to milliseconds	Limited to decapeptides in current implementation
MolEdit [38]	3D molecular generation	High validity and stability across scales	Supports zero-shot lead optimization	Computational intensity for large molecules

Genetic Algorithm Performance in Peptide Optimization

The performance of genetic algorithms in optimizing peptide sequences for aggregation propensity demonstrates remarkable efficiency. Starting with 1000 randomly sampled initial sequences with an average AP of 1.76, the algorithm achieved an average AP of 2.15 after 500 iterations with a mutation rate of 1% [40]. Validation through CGMD simulation confirmed the accuracy of these predictions, with example sequences showing close alignment between predicted and simulated AP values [40].

Specific sequence examples illustrate this performance:

- **LAPP Example**: VMDNAELDAQ with predicted AP of 1.14 (close to CGMD simulation)
- **HAPP Example**: WFLFFFLFFW with predicted AP of 2.24 (close to CGMD simulation)

Visualization of simulation snapshots confirmed that the LAPP example remained uniformly distributed in the simulation box without SASA decrease, while the HAPP example aggregated into large cluster structures within the same 125 ns simulation time [40].

Essential Research Reagents and Computational Tools

Table 3: Research Reagent Solutions for AI-Driven Molecular Modeling

Reagent/Tool	Type/Category	Function in Research	Example Implementation
Coarse-Grained Molecular Dynamics (CGMD) [40]	Simulation Method	Evaluates solvent-accessible surface area and defines aggregation propensity	Martini CGMD strategy for 125 ns simulations of decapeptides
Transformer Architecture [40] [38]	Neural Network Model	Base architecture for prediction and generation tasks	Self-attention mechanisms for AP prediction and molecular generation
Genetic Algorithms [40]	Optimization Algorithm	Generates and optimizes molecular sequences toward desired properties	Evolves decapeptide sequences with 1% mutation rate over 500 iterations
Denoising Diffusion Probabilistic Models (DDPMs) [38]	Generative Framework	Creates novel molecular structures through iterative denoising	MolEdit's foundation model with group-optimized labeling

Reagent/Tool	Type/Category	Function in Research	Example Implementation
Reinforcement Learning [39]	Machine Learning Paradigm	Guides molecular generation through reward functions	3D-Scaffold framework with atom-by-atom placement rewards
Monte Carlo Tree Search [40]	Search Algorithm	Enables targeted optimization while preserving functional features	Peptide sequence optimization with constrained search space

AI-driven modeling of 3D molecular architectures represents a paradigm shift in structural bioinformatics and drug discovery. The integration of foundation models with physics-informed principles has enabled unprecedented capabilities in generating valid, diverse molecular structures with desired properties. Approaches such as MolEdit's group-optimized diffusion, the 3D-Scaffold framework's reinforcement learning, and transformer-based aggregation prediction demonstrate the versatility and power of these methods across different molecular classes and design objectives.

As these technologies continue to mature, several challenges remain, including the computational intensity of training and inference, the need for improved interpretability of model predictions, and the integration of more comprehensive physical and biological constraints. Nevertheless, the rapid progress in AI-driven structure prediction suggests a future where de novo molecular design becomes increasingly routine, accelerating the discovery of novel therapeutics and functional materials. The convergence of foundation models with domain-specific knowledge in chemistry and biology promises to unlock deeper insights into molecular function and interaction, ultimately establishing these approaches as pivotal tools in advancing molecular science.

The exponential growth of biological sequence data has far outpaced the capacity of traditional experimental methods for functional characterization. Within the broader context of foundation models in bioinformatics, the development of scalable, accurate computational methods for predicting protein functions and reconstructing Gene Regulatory Networks (GRNs) has become a critical research frontier [41] [42]. Foundation models, pre-trained on vast datasets, are revolutionizing these domains by providing powerful representations that can be adapted to specific prediction tasks with limited annotated data [41].

Function annotation encompasses two primary domains: predicting the specific roles of proteins and elucidating the complex regulatory interactions between genes and their regulators. For proteins, this involves assigning functional descriptors such as Gene Ontology (GO) terms that describe molecular functions, biological processes, and cellular components [43] [42]. For GRNs, the challenge lies in reconstructing the directed networks of transcription factors (TFs), their target genes, and cis-regulatory elements (CREs) that collectively control cellular processes and responses [44] [45]. This technical guide examines the methodologies, experimental protocols, and computational tools that are advancing these interconnected fields.

Protein Function Prediction

Methodological Landscape

Protein function prediction has evolved from sequence-based homology inference to sophisticated frameworks that integrate multimodal biological data. Traditional methods relied heavily on sequence similarity and structural motifs, but contemporary approaches now incorporate protein-protein interactions, domain architectures, subcellular localization, and expression patterns to achieve more accurate and comprehensive annotations [43] [42].

Table 1: Key Computational Methods for Protein Function Prediction

Method	Approach	Data Modalities Integrated	Key Innovation
MIF2GO [46]	Deep Learning & Multimodal Fusion	Sequence, Domain, Subcellular Localization, Pathway, Interaction, Homology	Self-supervised pretraining with Siamese Contrastive Autoencoder (SCA) and hierarchical language model
GOHPro [47]	Heterogeneous Network Propagation	PPI networks, Domain profiles, Protein complexes, GO semantic relationships	Integrates protein functional similarity with GO hierarchical relationships in a two-layer network
DeepGOPlus [42]	Deep Learning	Protein sequence, homology	Combines sequence homology with deep convolutional neural networks scanning for predictive motifs
PANNZER [42]	Functional Ranking	Protein sequence	Combines motif scanning with functional annotation transfer from similar proteins

The MIF2GO Framework: A Case Study in Multimodal Integration

The MIF2GO framework exemplifies the trend toward sophisticated multimodal integration, sequentially fusing up to six biological modalities through three dedicated steps [46]:

- **Stage 1 - Siamese Contrastive Autoencoder (SCA)**: Encodes domain, subcellular localization, and pathway modalities into interaction and homology relationships using self-supervised learning. Contrastive learning brings the representation spaces of interaction and homology modalities closer together.
- **Stage 2 - Language Model with Hierarchical Adaptive Weighting (LM-HAW)**: Processes sequence modality using a self-supervised language model that extracts features from shallow, medium, and deep layers to capture different granularities of protein sequence information.
- **Stage 3 - Modal Hypernode Pooling (MHP)**: Fuses the embeddings from the SCA module and hierarchical sequence features from LM-HAW to generate unified protein representations for function prediction.

When evaluated on human protein datasets, MIF2GO achieved state-of-the-art performance with M-AUPR = 0.624 ± 0.002, m-AUPR = 0.804 ± 0.002, and F-max = 0.758 ± 0.001 for Molecular Function GO terms [46]. The method also demonstrated remarkable generalizability across species, including fruit fly, mouse, rat, S. cerevisiae, and B. subtilis, proving particularly effective for GO terms with few protein samples.

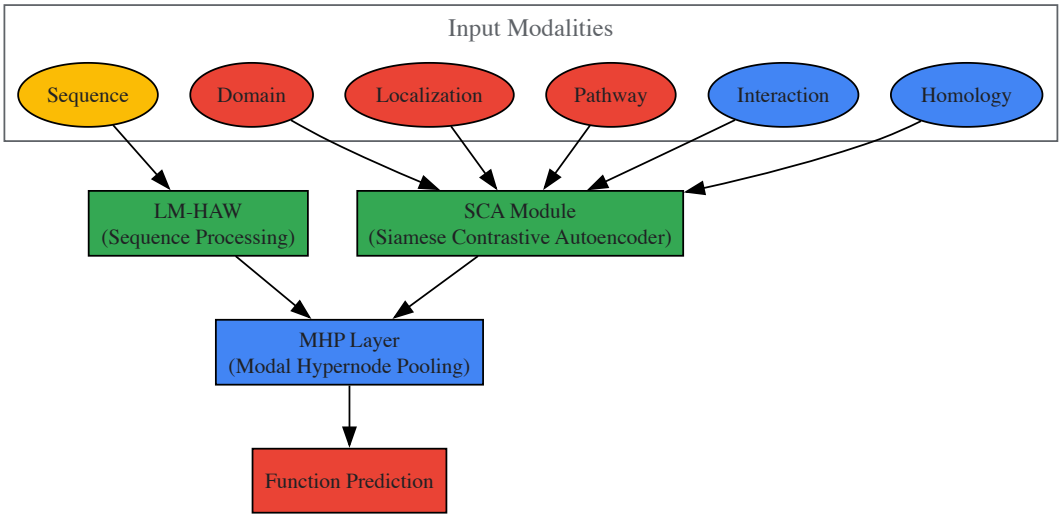

MIF2GO Multimodal Fusion Pipeline: The workflow illustrates the integration of six biological modalities through three specialized processing stages.

Gene Regulatory Network Inference

Methodological Foundations

GRN inference aims to reconstruct the directed regulatory relationships between transcription factors and their target genes. The methodological landscape has evolved significantly with advances in sequencing technologies, particularly with the advent of single-cell multi-omics data [44].

Table 2: Core Methodological Approaches for GRN Inference

Approach	Principle	Strengths	Limitations
Correlation-based [44]	Guilt-by-association via co-expression	Simple implementation, detects linear (Pearson) and non-linear (Spearman) relationships	Cannot distinguish directionality; confounded by indirect relationships
Regression Models [44]	Models gene expression as function of TFs/CREs	Interpretable coefficients indicate regulatory strength; handles multiple predictors	Prone to overfitting with high-dimensional predictors; unstable with correlated TFs
Probabilistic Models [48] [44]	Graphical models capturing dependence between variables	Provides confidence measures for interactions; handles uncertainty	Often assumes specific distributions (e.g., Gaussian) that may not fit biological data
Dynamical Systems [44]	Differential equations modeling system evolution over time	Captures temporal dynamics; interpretable parameters	Requires time-series data; computationally intensive; limited scalability
Deep Learning [49] [44] [45]	Neural networks learning complex hierarchical patterns	Captures non-linear relationships; integrates heterogeneous data	Requires large datasets; computationally intensive; limited interpretability

Hybrid Machine Learning for GRN Prediction

Recent research demonstrates that hybrid approaches combining convolutional neural networks (CNNs) with traditional machine learning consistently outperform individual methods. In plant systems, these hybrid models achieved over 95% accuracy on holdout test datasets by integrating prior knowledge with large-scale transcriptomic data from Arabidopsis thaliana, poplar, and maize [49] [45].

The typical hybrid architecture involves a two-step process:

- **Feature Extraction**: CNNs learn high-level representations from transcriptomic data
- **Classification**: Machine learning models (e.g., Random Forests, Extremely Randomized Trees) predict regulatory relationships using the CNN-derived features

These models successfully identified known master regulators of the lignin biosynthesis pathway (MYB46, MYB83) and upstream regulators (VND, NST, SND families) with higher precision than traditional methods [50] [45].

Transfer Learning for Cross-Species GRN Inference

A significant challenge in GRN inference is the limited availability of training data for non-model species. Transfer learning addresses this by enabling cross-species inference, where models trained on data-rich species (e.g., Arabidopsis) are adapted to species with limited data (e.g., poplar, maize) [45]. This approach enhances model performance and demonstrates the feasibility of knowledge transfer across evolutionary related species, particularly when considering conserved transcription factor families and orthologous genes.

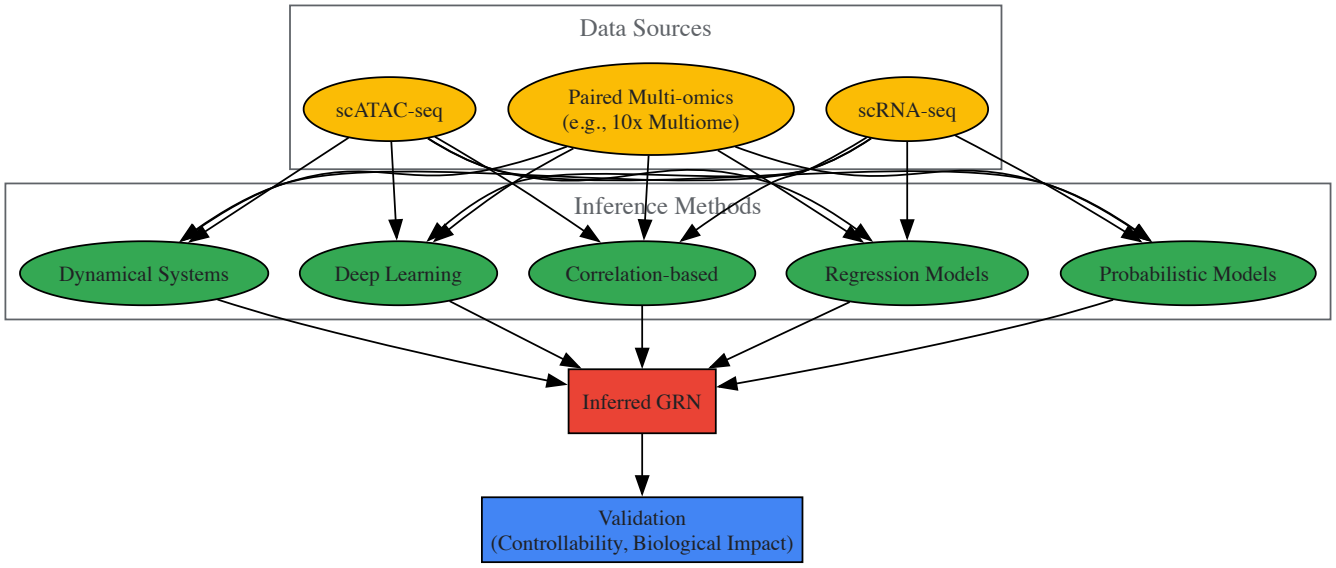

GRN Inference Workflow: From multi-omics data acquisition through computational inference to experimental validation.

Experimental Protocols and Validation Frameworks

GRN Inference Validation Protocol

Validating inferred GRNs presents unique challenges, particularly given the typical absence of complete ground truth networks. The concept of objective inferential validity provides a application-focused validation framework [48].

Protocol: Controllability-Based Validation

- **Network Inference**: Apply inference algorithms to gene expression data generated from a known GRN
- **Control Policy Design**: Design intervention strategies (stationary control) based on the inferred network to minimize undesirable phenotypic states
- **Performance Assessment**: Evaluate how well the control policy performs when applied to the true network based on steady-state mass shift toward desirable states
- **Comparative Analysis**: Compare inference methods based on control efficacy rather than topological accuracy alone

This approach recognizes that networks with different topologies may yield similar control performance, and prioritizes operational utility over strict structural fidelity [48].

Single-Cell Multi-omics GRN Inference Protocol

The emergence of single-cell multi-omics technologies has enabled unprecedented resolution for GRN inference. The following protocol leverages paired scRNA-seq and scATAC-seq data [44]:

Sample Preparation

- Perform single-cell multi-ome sequencing (e.g., 10x Multiome, SHARE-seq) to simultaneously profile gene expression and chromatin accessibility
- Target: 5,000-10,000 cells per condition to adequately capture cellular heterogeneity

Data Preprocessing

- Quality control: Filter cells based on mitochondrial percentage, unique molecular identifiers (UMIs), and gene counts
- Normalization: Normalize RNA counts using SCTransform and ATAC data using term frequency-inverse document frequency (TF-IDF)
- Integration: Harmonize RNA and ATAC modalities using weighted nearest neighbors (WNN)

Network Inference

- Identify candidate regulatory regions: Link accessible peaks to genes based on cis-regulatory domains (e.g., ±500kb from TSS)
- Calculate associations: Correlate TF expression with both target gene expression and TF motif accessibility in linked peaks
- Construct network: Integrate evidence using regression-based (e.g., LASSO) or machine learning methods to infer direct regulatory relationships

Validation

- Benchmark against known regulatory interactions from literature and databases
- Perform functional validation through perturbation experiments (CRISPRi/a) on top predicted regulators

The Scientist's Toolkit: Research Reagent Solutions

Table 3: Essential Research Reagents and Platforms for Function Annotation Studies

Reagent/Platform	Function	Application Context
10x Multiome [44]	Simultaneous profiling of gene expression and chromatin accessibility in single cells	Single-cell multi-omics GRN inference; identification of cell-type specific regulatory programs
SHARE-seq [44]	Joint measurement of chromatin accessibility and gene expression in single cells	Alternative to 10x Multiome for capturing regulatory relationships at single-cell resolution
DAP-seq [49] [45]	High-throughput identification of transcription factor binding sites in vitro	Validation of TF-binding specificities; prior knowledge for GRN inference algorithms
UniProtKB [46] [43]	Comprehensive protein knowledgebase with functional annotation	Benchmark dataset for protein function prediction methods; source of training data
Gene Ontology (GO) [43] [42] [47]	Standardized vocabulary for protein functions across three domains: MF, BP, CC	Gold-standard for protein function prediction evaluation; framework for annotation transfer
neXtProt [42]	Human protein-centric knowledge platform with detailed functional annotation	Validation resource for human protein function prediction methods; source of curated annotations
Complex Portal [47]	Manually curated resource of macromolecular complexes	Source of protein complex information for functional similarity networks in methods like GOHPro
Pfam [47]	Comprehensive collection of protein domains and families	Domain annotation for protein function prediction; input for domain structural similarity networks

The fields of protein function annotation and GRN inference are undergoing rapid transformation driven by advances in multimodal data integration, sophisticated computational frameworks, and the emergence of foundation models in bioinformatics. The integration of multiple biological modalities—from sequence and structure to interactions and expression—has proven essential for overcoming the limitations of individual data types and achieving robust predictive performance.

Hybrid approaches that combine the representational power of deep learning with the interpretability and efficiency of traditional machine learning are demonstrating remarkable success in both protein function prediction and GRN inference. Furthermore, transfer learning strategies are addressing the critical challenge of data scarcity in non-model organisms by enabling knowledge transfer from well-characterized species.

As foundation models continue to mature, their ability to learn generalizable representations from massive biological datasets promises to further accelerate progress in these domains. However, challenges remain in model interpretability, integration of diverse data modalities, and validation of predictions in biologically meaningful contexts. The continued development of sophisticated computational methods, coupled with rigorous experimental validation frameworks, will be essential for unraveling the complex functional landscapes of proteins and gene regulatory networks.

The traditional drug discovery paradigm is notoriously protracted, expensive, and prone to failure, often requiring over a decade and costing approximately $2.6 billion to bring a new drug to market, with a success rate of less than 10% [51]. This inefficiency underscores a critical need for innovative computational approaches that can enhance predictive reliability while reducing cost and time constraints. Artificial intelligence (AI), particularly deep learning and foundation models, has emerged as a transformative force, offering a paradigm shift from conventional computational methods. Unlike traditional techniques, AI excels in handling high-dimensional biological data, uncovering complex molecular patterns, and optimizing biochemical interactions with minimal human intervention [52] [51].

Foundation models, pre-trained on vast datasets of biological sequences, structures, and interactions, provide a powerful starting point for a multitude of downstream tasks in bioinformatics. These models capture fundamental principles of biology, from evolutionary constraints on protein sequences to the physical rules governing molecular folding and binding. This technical guide explores how these AI-driven methodologies are being systematically integrated into the critical early stages of drug discovery—target identification and compound design—to accelerate the development of new therapeutics. By leveraging foundation models, researchers can now more accurately identify druggable targets and design optimized compounds with desired properties, thereby streamlining the entire preclinical pipeline [8].

Foundation Models in Bioinformatics: A New Paradigm

Foundation models in bioinformatics are large-scale neural networks pre-trained on extensive corpora of biological data, such as genomic sequences, protein structures, and scientific literature. Their primary advantage lies in their ability to generate generalized representations of biological entities, which can then be fine-tuned for specific tasks with limited additional data. This represents a significant advancement over earlier models that were trained on narrow datasets for singular objectives [8].

At the sequence level, models like Ankh (for proteins) and MolFormer (for small molecules) learn to represent biological polymers and compounds as meaningful numerical embeddings. These embeddings encapsulate crucial functional and structural characteristics, enabling a model to, for instance, understand the functional implication of a specific protein domain or the reactivity of a molecular moiety directly from its sequence [53] [8]. For structural biology, breakthroughs such as AlphaFold have demonstrated that AI can predict protein 3D structures with atomic accuracy, a problem that had remained unsolved for decades. These structural predictions are invaluable for assessing target druggability by revealing well-defined binding pockets [8] [51]. The following diagram illustrates the general architecture of a foundation model processing biological data.

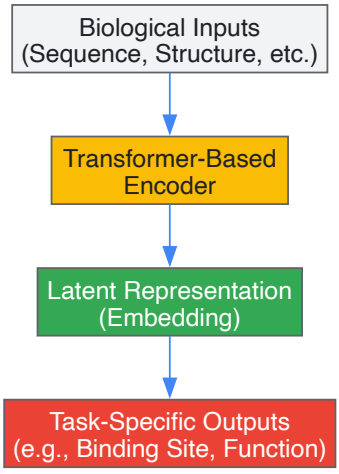

The workflow begins with diverse biological inputs, which are processed through a deep learning encoder—often a transformer architecture—to create a rich, latent representation. This representation serves as a foundational embedding that can be directed to various downstream prediction tasks, forming the core of the foundation model approach [53] [8].

AI-Driven Target Identification

Target identification is the foundational step in drug discovery, aiming to pinpoint biomolecules (typically proteins) whose modulation would produce a therapeutic effect in a specific disease. AI dramatically accelerates this process by analyzing multi-omics data (genomics, transcriptomics, proteomics) to uncover hidden patterns and novel oncogenic vulnerabilities that might be missed by traditional methods [51].

A key application is the prediction of protein-ligand binding sites. Methods like **LABind** exemplify the power of a ligand-aware, structure-based approach. LABind utilizes a graph transformer to capture binding patterns from the local spatial context of protein structures. It integrates protein sequence embeddings from a pre-trained language model (Ankh) with structural features, while simultaneously processing ligand information (from

SMILES sequences) via a molecular pre-trained model (MolFormer). A cross-attention mechanism then learns the distinct binding characteristics between the protein and the specific ligand, enabling the model to predict binding sites not just for ligands seen during training, but also for unseen ligands. This generalizability is a hallmark of a robust foundation model [53].

Experimental Protocol for Binding Site Prediction

For researchers seeking to implement a state-of-the-art binding site prediction method, the following protocol outlines the key steps based on the LABind methodology [53]:

- **Input Preparation**:
  - **Protein Data**: Obtain the protein's amino acid sequence and 3D structure (experimental from PDB or predicted via ESMFold/AlphaFold).
  - **Ligand Data**: For a ligand-aware prediction, provide the ligand's SMILES string.
- **Feature Extraction**:
  - **Protein Sequence Representation**: Generate protein residue embeddings using a pre-trained protein language model (e.g., Ankh).
  - **Protein Structure Representation**: Convert the 3D structure into a graph where nodes represent residues. Node features include angles, distances, and directions derived from atomic coordinates. Edge features encode spatial relationships between residues.
  - **Ligand Representation**: Process the ligand's SMILES string through a molecular pre-trained model (e.g., MolFormer) to obtain a latent ligand representation.
- **Model Inference**:
  - The model processes the protein graph and ligand representation through a graph transformer and cross-attention layers to learn interaction features.
  - A multi-layer perceptron (MLP) classifier uses these features to perform a per-residue classification, predicting whether each residue is part of a binding site.
- **Output and Validation**:
  - The output is a binary classification for every residue in the protein.
  - Predictions are typically evaluated using metrics such as AUC (Area Under the Curve), AUPR (Area Under the Precision-Recall Curve), and F1-score on benchmark datasets (e.g., DS1, DS2, DS3). LABind demonstrated superior performance on these benchmarks, outperforming other advanced methods [53].

The following workflow chart visualizes this multi-step process from data input to functional prediction.

Performance Benchmarking of AI Target Identification Methods

Table 1: Performance Metrics of AI Models in Target Identification and Assessment

Model / Framework	Primary Task	Key Performance Metric	Reported Result
LABind [53]	Protein-ligand binding site prediction	AUPR (Area Under Precision-Recall Curve)	Outperformed competing methods on DS1, DS2, DS3 benchmarks
AlphaFold [8]	Protein 3D structure prediction	Median Accuracy (CASP14)	0.96 Å (median)
optSAE + HSAPSO [52]	Drug classification & target identification	Accuracy	95.52%
AI-based Binding Predictor [8]	Protein-ligand interaction	AUC (Area Under ROC Curve)	0.93
Graph-based Deep Learning [52]	Protein sequence analysis for drug-target prediction	Accuracy	95%

AI-Accelerated Compound Design and Optimization

Once a target is identified, the focus shifts to designing and optimizing compounds that can effectively and selectively modulate it. AI methodologies are revolutionizing this space by enabling rapid virtual screening of ultra-large chemical libraries and the *de novo* design of novel molecular entities with optimized properties.

A prominent approach involves using **Stacked Autoencoders (SAE)** for robust feature extraction from chemical data. For instance, the **optSAE + HSAPSO** framework integrates an SAE with a Hierarchically Self-Adaptive Particle Swarm Optimization algorithm for hyperparameter tuning. This combination has demonstrated a high accuracy of 95.52% in drug classification tasks on DrugBank and Swiss-Prot datasets, with significantly reduced computational complexity (0.010 seconds per sample) and high stability (± 0.003) [52]. This efficiency is critical for screening millions of compounds in silico.

Furthermore, deep learning models can now generate novel molecular structures from scratch. These generative models learn the underlying probability distribution of known drug-like molecules and can sample new compounds from this distribution, which are then optimized through reinforcement learning feedback for specific properties such as high binding affinity, synthesizability, and low toxicity [8] [51].

Experimental Protocol for AI-Driven Compound Optimization

The following protocol details the process for implementing an optimized autoencoder framework for compound design and classification, as exemplified by the optSAE + HSAPSO model [52]:

- **Data Collection and Preprocessing**:
  - Curate a large dataset of chemical structures with known biological activities or properties (e.g., from DrugBank, ChEMBL).
  - Standardize molecular structures and encode them into numerical descriptors or fingerprints (e.g., ECFP, molecular weight, logP).
- **Model Construction and Pre-training**:
  - Construct a Stacked Autoencoder (SAE) comprising multiple layers of encoders and decoders. The encoder compresses the input features into a lower-dimensional latent representation, and the decoder attempts to reconstruct the input from this representation.
  - Pre-train the SAE in an unsupervised manner to learn efficient feature representations of the chemical space.
- **Hyperparameter Optimization**:
  - Employ an optimization algorithm like Hierarchically Self-Adaptive PSO (HSAPSO) to fine-tune the hyperparameters of the SAE (e.g., learning rate, number of layers, nodes per layer). HSAPSO dynamically adapts parameters during training to balance exploration and exploitation, leading to faster convergence and avoidance of local minima.
- **Model Training and Validation**:
  - Use the optimized SAE to extract features from the molecular data.
  - Train a classifier (e.g., a softmax layer) on these features for the specific task, such as predicting target binding or classifying drug likeness.
  - Validate the model using hold-out test sets or cross-validation, reporting metrics like accuracy, AUC, and Matthews Correlation Coefficient (MCC).

The iterative optimization process is visualized in the diagram below.

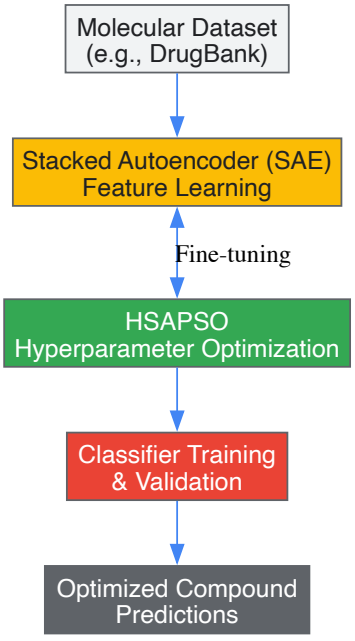

Performance Benchmarking of AI-Based Compound Design

Table 2: Performance Metrics of AI Models in Compound Design and Optimization

Model / Framework	Primary Task	Key Advantage	Reported Result / Metric
optSAE + HSAPSO [52]	Drug classification & target ID	High accuracy & computational efficiency	95.52% Accuracy, 0.010s/sample
Generative AI Models [8]	*De novo* drug design	High design success rate	Success rate up to 92%
AI-powered Predictive Modeling [51]	Binding affinity prediction	Improved lead optimization	Enhanced predictive accuracy vs traditional methods
Graph-based Deep Learning [52]	Drug-target interaction prediction	Utilizes complex structural data	95% Accuracy

The practical application of these AI methodologies relies on a ecosystem of software tools, databases, and computational resources. The table below catalogues key resources cited in this guide.

Table 3: Key Research Reagents and Computational Tools for AI-Driven Drug Discovery

Resource Name	Type	Primary Function in Research
LABind [53]	Software Tool	Predicts protein binding sites for small molecules and ions in a ligand-aware manner using graph transformers.
Ankh [53]	Foundation Model	A protein language model used to generate informative sequence representations and embeddings for proteins.
MolFormer [53]	Foundation Model	A molecular language model that generates representations of ligands from their SMILES strings.
AlphaFold [8] [51]	Software Tool	Predicts highly accurate 3D protein structures from amino acid sequences, aiding druggability assessment.
ESMFold / OmegaFold [53]	Software Tool	Provides protein structure predictions, used as input for binding site prediction when experimental structures are unavailable.
Stacked Autoencoder (SAE) [52]	Algorithm	Used for unsupervised feature learning and dimensionality reduction of high-dimensional molecular data.
Particle Swarm Optimization (PSO) [52]	Algorithm	An optimization technique used for efficient hyperparameter tuning of deep learning models.
DrugBank / Swiss-Prot [52]	Database	Curated repositories of drug, target, and protein information used for training and benchmarking AI models.

The integration of foundation models and other AI technologies into drug discovery represents a fundamental shift in how researchers approach target identification and compound design. By leveraging pre-trained knowledge on vast biological datasets, these methods offer unprecedented accuracy and speed, as evidenced by the performance of tools like LABind for binding site prediction and optSAE+HSAPSO for molecular classification. They are not merely incremental improvements but are transformative tools that can decipher the complex language of biology and chemistry. As these models continue to evolve, incorporating larger and more diverse datasets, and as their interpretability improves, their role in de-risking the drug discovery pipeline and delivering novel therapeutics to patients will only become more profound.

The exponential growth of single-cell RNA sequencing (scRNA-seq) data has revolutionized our understanding of cellular heterogeneity but simultaneously created formidable computational challenges. Traditional analytical pipelines, designed for lower-dimensional data, struggle with the high dimensionality, sparsity, and technical noise characteristic of modern single-cell datasets [12] [10]. Within this context, single-cell foundation models (scFMs) have emerged as transformative tools capable of learning universal biological representations from massive datasets through self-supervised learning. These models adapt the transformer architecture—originally developed for natural language processing—to decode the complex "language" of biology, where cells are treated as sentences and genes as words [10]. This technical guide examines the architecture, performance, and practical implementation of scFMs for atlas-level data integration, providing researchers and drug development professionals with the comprehensive framework needed to leverage these powerful tools in pushing the boundaries of single-cell biology.

Core Architectural Principles of scFMs

Tokenization Strategies for Non-Sequential Biological Data

Unlike natural language, gene expression data lacks inherent sequential ordering, presenting unique tokenization challenges. scFMs address this through several innovative approaches that convert raw expression values into model-processable tokens. The dominant strategy involves **gene ranking by expression levels**, where genes within each cell are ordered by their expression magnitude, creating a deterministic sequence for transformer processing [10]. Alternative approaches include **value binning**, which partitions expression values into discrete bins (scGPT), and **genomic position ordering**, which orders genes based on their chromosomal coordinates (UCE) [12] [22]. Each gene token typically combines a learnable embedding representing the gene identity with a separate value embedding capturing its expression level in the specific cell. Special tokens are often incorporated to enrich biological context, including cell-level metadata tokens, modality indicators for multi-omic models, and batch-specific tokens to mitigate technical variations [10].

Model Architectures and Pretraining Strategies

Most scFMs leverage transformer architectures with specific adaptations for biological data. The field primarily utilizes two variants: **BERT-like encoder models** with bidirectional attention that learn from all genes simultaneously (Geneformer, scBERT), and **GPT-like decoder models** with masked self-attention that iteratively predict masked genes conditioned on known expression patterns (scGPT) [10]. The pretraining process employs self-supervised objectives, most commonly **masked gene modeling** (MGM), where the model learns to reconstruct randomly masked portions of the input expression profile [12] [10]. Additional pretraining tasks include **contrastive learning** to align similar cellular states and **generative pretraining** for predicting gene expression distributions. These approaches enable models to learn fundamental biological principles from massive datasets—scGPT trained on 33 million cells, scFoundation on 50 million, and Nicheformer on 110 million cells—creating representations that capture universal features of gene regulation and cellular identity [54] [10].

Quantitative Benchmarking of scFM Performance

Comprehensive Evaluation Across Biological Tasks

Rigorous benchmarking studies have evaluated scFMs against traditional methods across diverse tasks. A 2025 benchmark assessed six leading scFMs (Geneformer, scGPT, UCE, scFoundation, LangCell, and scCello) alongside established baselines using 12 metrics spanning unsupervised, supervised, and knowledge-based approaches [12] [22]. The evaluation encompassed two gene-level tasks (tissue specificity prediction, Gene Ontology term prediction) and four cell-level tasks (batch integration, cell type annotation, cancer cell identification, and drug sensitivity prediction) across multiple datasets and cancer types [12]. Performance was assessed using both traditional metrics and novel biology-informed measures like scGraph-OntoRWR, which evaluates consistency with known biological relationships, and Lowest Common Ancestor Distance (LCAD), which quantifies the severity of cell type misclassification errors [12] [22].

Table 1: Performance Overview of Leading scFMs Across Key Tasks

Model	Pretraining Scale	Batch Integration	Cell Type Annotation	Gene Function Prediction	Clinical Task Performance
scGPT	33M cells	Excellent	Superior (zero-shot)	Strong	Robust across cancer types
Geneformer	30M cells	Good	Good	Excellent	Variable
scFoundation	50M cells	Good	Moderate	Strong	Limited data
scBERT	Limited	Moderate	Weaker	Weaker	Not assessed
Traditional SVM	N/A	Good (with Harmony)	Excellent (with training)	Limited	Good for specific datasets

Task-Specific Strengths and Limitations

The benchmarking reveals that no single scFM consistently outperforms all others across diverse applications, emphasizing the need for task-specific model selection [12] [22]. For **batch integration**, scFMs demonstrate remarkable robustness in removing technical variations while preserving biological signals, particularly in challenging cross-tissue and cross-species scenarios. In **cell type annotation**, scGPT excels in zero-shot settings where models identify novel cell types without retraining, while traditional support vector machines (SVM) remain competitive for within-dataset classifications when sufficient training data exists [12] [55]. For **gene-level tasks**, Geneformer and scFoundation demonstrate exceptional performance in predicting gene functions and relationships, benefiting from their specialized pretraining strategies [12] [56]. In **clinically relevant applications** such as cancer cell identification and drug sensitivity prediction, scFMs show promise but with greater performance variability across different cancer types and therapeutics [12].

Table 2: Comparative Performance Metrics Across Model Types

Evaluation Metric	Leading scFMs	Traditional ML (SVM)	Baseline Methods (Seurat, Harmony)
Batch Integration (kBET)	0.72-0.89	0.68-0.85	0.71-0.88
Cell Annotation Accuracy (Zero-shot)	0.75-0.92	Not applicable	Not applicable
Cell Annotation Accuracy (Supervised)	0.85-0.96	0.88-0.98	0.65-0.82
Gene Ontology Prediction (AUROC)	0.81-0.90	0.70-0.78	Limited capability
Drug Sensitivity Prediction (r)	0.45-0.62	0.40-0.58	Limited capability
Computational Resources	High	Low to moderate	Low

Experimental Protocols for scFM Implementation

Protocol 1: Zero-Shot Cell Type Annotation

Purpose: To identify cell types in novel datasets without task-specific training. **Materials**: Preprocessed scRNA-seq dataset (count matrix), pretrained scFM (scGPT or Geneformer recommended), reference cell type embeddings if available. **Methodology**: Begin with standard preprocessing—quality control, normalization, and filtering. For scGPT, implement the following key steps: (1) Map gene identifiers to the model's vocabulary, padding or trimming to match the model's input dimensions (typically 1,200-2,000 genes); (2) Generate cell embeddings through a forward pass of the pretrained model; (3) Compute similarity scores between query cell embeddings and reference cell type embeddings using cosine similarity in the latent space; (4) Assign cell types based on highest similarity scores, applying a confidence threshold to flag low-confidence predictions [12] [56]. **Validation**: Compare annotations with marker gene expression and assess using the LCAD metric to ensure biologically plausible misclassifications when perfect accuracy isn't achieved [12] [22].

Protocol 2: Atlas-Level Batch Integration

Purpose: To integrate multiple single-cell datasets into a unified embedding space while preserving biological variation. **Materials**: Multiple scRNA-seq datasets with batch metadata, pretrained scFM (scGPT or scFoundation recommended), computational resources with adequate GPU memory. **Methodology**: Process each dataset independently through the pretrained model to obtain batch-specific cell embeddings. Apply the model's built-in integration capabilities—scGPT uses attention masking and batch tokenization, while scFoundation employs a read-depth-aware masked gene modeling approach [12] [22]. The key innovation in scFMs is their ability to learn batch-invariant representations during pretraining, which enables effective integration without explicit batch correction algorithms. For evaluation, utilize the scGraph-OntoRWR metric to verify preservation of biological relationships and standard metrics like kBET to assess batch mixing [12]. The roughness index (ROGI) can predict integration performance by measuring the smoothness of the cell-property landscape in the latent space [12] [22].

Visualization Framework for scFM Workflows

The following diagrams illustrate key experimental workflows and conceptual relationships in scFM applications.

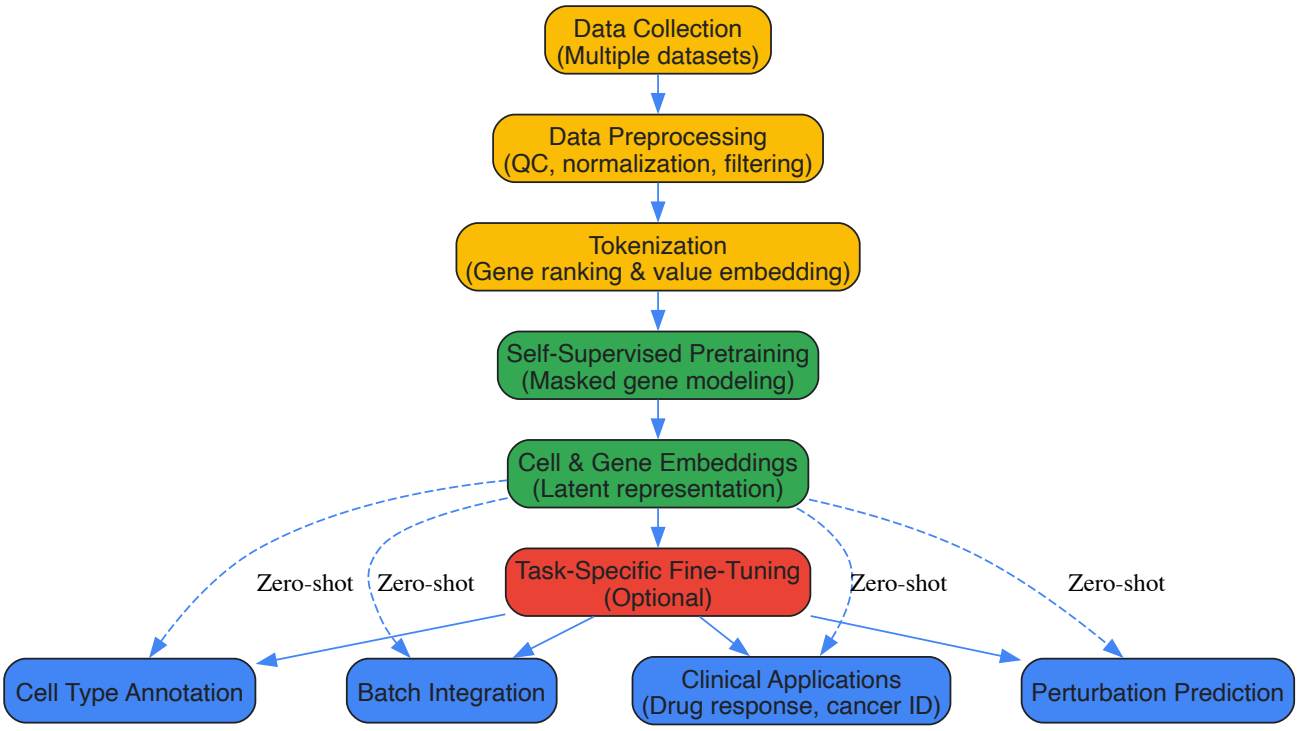

Diagram 1: Comprehensive scFM Workflow

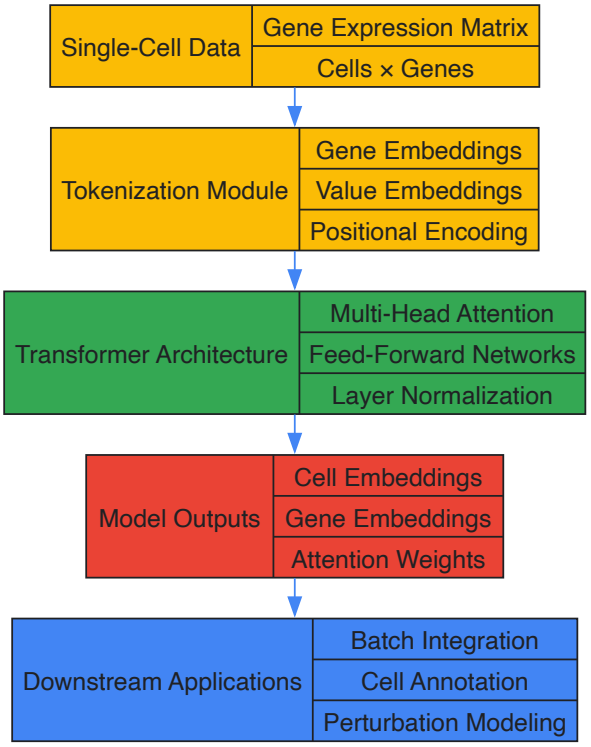

Diagram 2: scFM Architecture Overview

Table 3: Essential Resources for scFM Implementation

Resource Category	Specific Tools/Platforms	Function/Purpose	Access Method
Model Frameworks	scGPT, Geneformer, scFoundation, BioLLM	Core model architectures and pretrained weights	GitHub, Hugging Face, BioLLM unified API
Data Repositories	CELLxGENE Discover, DISCO, Human Cell Atlas	Curated single-cell datasets for training and validation	Public portals (cellxgene.cziscience.com)

Resource Category	Specific Tools/Platforms	Function/Purpose	Access Method
Benchmarking Platforms	BioLLM, scRNAseq_Benchmark	Standardized evaluation of model performance	GitHub, published pipelines
Computational Infrastructure	GPU clusters (NVIDIA A100/H100), Cloud computing (AWS, GCP)	Hardware acceleration for training and inference	Institutional HPC, commercial cloud
Visualization Tools	CELLxGENE Explorer, Scanpy, Seurat	Interactive exploration of model outputs	Python/R packages, web applications

Future Directions and Clinical Translation

The rapid evolution of scFMs points toward several transformative directions. **Multimodal integration** represents the frontier, with models like CellWhisperer demonstrating the power of combining transcriptomic data with textual annotations through contrastive learning, enabling natural language queries of cellular data [57]. **Spatial context awareness** is advancing through architectures like Nicheformer, which models cellular niches across millions of spatially resolved cells [54]. For clinical translation, key challenges include improving model interpretability to build trust in predictive outputs and enhancing robustness across diverse patient populations and experimental conditions [54] [10]. The development of federated learning frameworks will enable model refinement across institutions while preserving data privacy, accelerating the incorporation of scFMs into biomarker discovery, therapeutic target identification, and personalized treatment stratification in clinical practice [54].

Single-cell foundation models represent a paradigm shift in computational biology, offering unprecedented capabilities for atlas-level data integration and biological discovery. While benchmarking reveals that traditional methods remain competitive for specific tasks with sufficient training data, scFMs provide superior generalization, zero-shot capabilities, and multimodal integration potential. Their ability to learn universal representations from massive datasets positions them as indispensable tools for constructing comprehensive cell atlases, unraveling disease mechanisms, and accelerating therapeutic development. As the field matures, standardized benchmarking frameworks like BioLLM and biologically informed evaluation metrics will guide researchers in selecting appropriate models for specific applications, ultimately bridging the gap between computational innovations and biological insights that transform our understanding of cellular systems.

The advent of high-throughput technologies has revolutionized biology, enabling the generation of vast amounts of molecular data across multiple layers, including the genome, epigenome, transcriptome, proteome, and metabolome [58]. While each omic dataset provides valuable insights individually, in concert, they can reveal new and valuable insights into cellular heterogeneity, developmental trajectories, and disease mechanisms [59] [60]. This integrated approach, known as multi-omics or multimodal integration, aims to holistically understand biological systems by simultaneously analyzing data from these different molecular levels. The primary challenge lies in the computational integration of these complex, high-dimensional datasets, which often differ in scale, noise characteristics, and biological meaning [59]. Successfully overcoming this challenge enables researchers to assess the flow of information from one omic level to another, thereby bridging the critical gap from genotype to phenotype [58].

The era of single-cell and spatial omics technologies has further intensified the need for sophisticated integration strategies. These technologies produce data that captures molecular states across millions of individual cells, offering unprecedented resolution but also introducing new complexities related to data sparsity and technical variability [60]. More recently, foundation models (FMs), originally developed for natural language processing, have emerged as transformative tools for decoding this cellular complexity. These large, pretrained neural networks learn universal representations from massive and diverse datasets, demonstrating exceptional capabilities in cross-task generalization and multimodal alignment [60] [3]. This review explores the current landscape of multimodal omics integration, with a specific focus on the strategies, computational tools, and emerging foundation models that are driving the field toward a more complete, holistic understanding of biological systems.

Defining the Integration Landscape

The computational tools and strategies for multimodal integration can be meaningfully categorized based on the nature of the input data. A principal distinction is whether the multi-omics data is **matched** (profiled from the same cell) or **unmatched** (profiled from different cells) [59]. This distinction fundamentally shapes the integration approach.

- **Matched (Vertical) Integration**: This strategy merges data from different omics layers within the same set of cells or samples. The cell itself serves as a natural anchor to bring these modalities together [59]. This is typically the most straightforward integration scenario, as the direct correspondence between measurements in the same cell provides a strong biological link. Technologies that concurrently profile RNA and protein (e.g., CITE-seq) or RNA and epigenomic information (e.g., scATAC-seq) are prime candidates for this approach.

- **Unmatched (Diagonal) Integration**: This form addresses the more substantial challenge of integrating omics data drawn from distinct populations of cells. Since the cell cannot be used as an anchor, the methodology must instead project cells from different modalities into a co-embedded space or non-linear manifold to find commonality [59]. This approach is technically demanding but valuable, as it allows the combination of datasets generated in separate experiments.

- **Mosaic Integration**: An alternative to diagonal integration, this method is applicable when the experimental design includes various combinations of omics that create sufficient overlap. For example, if one sample is assessed for transcriptomics and proteomics, another for transcriptomics and epigenomics, and a third for proteomics and epigenomics, the commonalities between these samples can be leveraged for integration [59] [60]. Tools like COBOLT and MultiVI are designed for this type of analysis [59].

Dominant Computational Methodologies

The computational frameworks employed for integration are as varied as the data types themselves. They range from classical statistical models to advanced deep-learning architectures.

Table 1: Core Computational Methodologies for Multimodal Integration

Methodology	Core Principle	Representative Tools	Typely Applied Data
Matrix Factorization	Decomposes high-dimensional data matrices into lower-dimensional representations (factors) shared across omics.	MOFA+ [59]	mRNA, DNA methylation, chromatin accessibility
Deep Learning (Autoencoders)	Uses neural networks to compress data into a latent space, forcing the integration of different modalities.	scMVAE[cite:1], DCCA[cite:1], totalVI[cite:1]	mRNA, chromatin accessibility, protein
Variational Autoencoders	A probabilistic variant of autoencoders that learns a distribution of the latent space, often providing better generalization.	GLUE[cite:1], Cobolt[cite:1]	Chromatin accessibility, DNA methylation, mRNA
Manifold Alignment	Projects different datasets onto a common low-dimensional manifold, preserving the intrinsic structure of each.	UnionCom[cite:1], Pamona[cite:1]	mRNA, DNA methylation, chromatin accessibility
Bayesian Models	Employs probabilistic frameworks to integrate data and quantify uncertainty in the results.	BREM-SC[cite:1]	mRNA, protein
Network-Based Methods	Leverages graph theory to connect entities from different omics layers based on prior knowledge or data-derived correlations.	citeFUSE[cite:1], Seurat v4[cite:1]	mRNA, protein, accessible chromatin

The Emergence of Foundation Models in Omics Integration

Foundation models represent a paradigm shift in bioinformatics. These are large-scale models pretrained on vast datasets using self-supervised learning objectives, which can then be adapted (via fine-tuning) to a wide range of downstream tasks with minimal task-specific data [3]. In the context of single-cell omics, FMs leverage architectures like transformers to learn universal representations of cells and genes.

These models excel through pretraining strategies such as **masked gene modeling** (where parts of the input data are hidden and the model learns to predict them), **contrastive learning** (which teaches the model to identify similar and dissimilar data pairs), and **multimodal alignment** (which explicitly learns the relationships between different data types) [60]. A key strength is their **cross-modal generalization capability**. For instance, a model pretrained on transcriptomic data can often make accurate predictions on epigenomic data or even align histology images with spatial gene expression, as demonstrated by PathOmCLIP [60].

Table 2: Notable Foundation Models for Single-Cell and Multi-Omics Analysis

Model	Year	Key Features	Reported Performance / Application
scGPT [60]	2024	Generative pretrained transformer; trained on over 33 million cells.	Superior performance in zero-shot cell type annotation, multi-omic integration, and gene network inference.
scPlantFormer [60]	2024	Lightweight foundation model pretrained on 1 million *Arabidopsis thaliana* cells.	92% cross-species annotation accuracy in plant systems.
Nicheformer [60]	2024	Employs graph transformers to model spatial cellular niches; trained on 53 million spatially resolved cells.	Enables spatial context prediction and integration.
PathOmCLIP [60]	2024	Uses contrastive learning to align tumor histology images with spatial gene expression.	Validated across five tumor types for gene expression prediction from histology.
EpiAgent [60]	-	Epigenomic foundation model focused on ATAC-seq data.	Capable of cis-regulatory element (cCRE) reconstruction in a zero-shot manner.

Practical Protocols for Multi-Omics Integration

A Web-Based Workflow for Multi-Omics Analysis

For researchers without extensive computational backgrounds, web-based suites like the **Analyst** software provide an accessible entry point. A typical integrative workflow using these tools can be executed in approximately two hours and involves three key components [61]:

- **Single-Omics Data Analysis**: The first step involves processing and analyzing each omics dataset individually to identify significant features (e.g., differentially expressed genes or metabolites). This can be done using tools like **ExpressAnalyst** (for transcriptomics/proteomics) and **MetaboAnalyst** (for lipidomics/metabolomics) [61].
- **Knowledge-Driven Integration**: The lists of significant features from the previous step are used as inputs to a tool like **OmicsNet**. This platform allows researchers to project these features onto biological networks (e.g., protein-protein interaction, gene regulatory networks) to understand their interrelationships and identify key regulatory nodes [61].
- **Data-Driven Integration**: The normalized matrices from all omics datasets, along with sample metadata, are loaded into a tool like **OmicsAnalyst**. This software performs joint dimensionality reduction (e.g., using PCA or t-SNE) to create a unified view of the samples, revealing patterns that might be invisible in any single omics dataset [61].

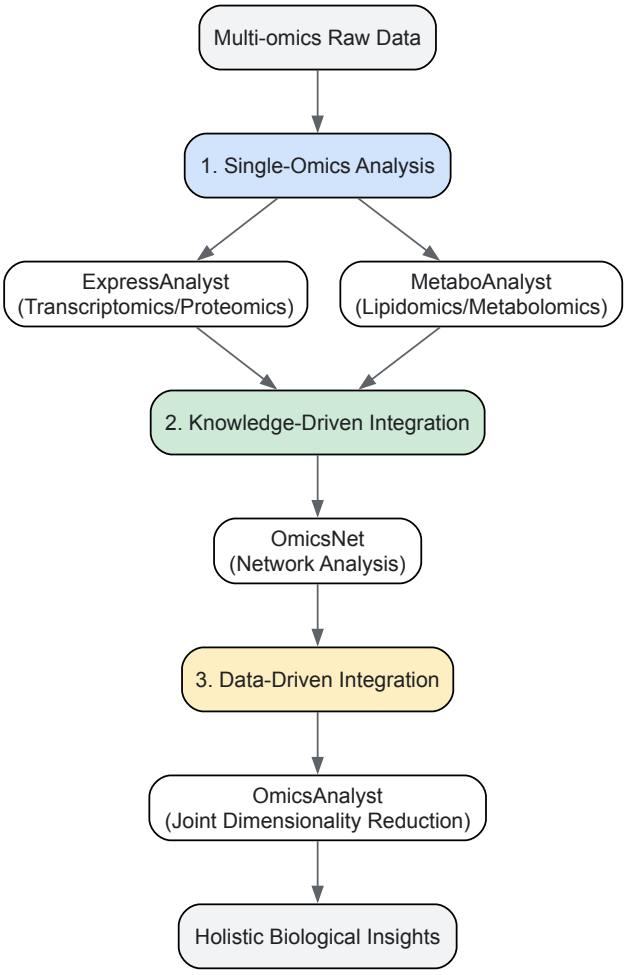

Experimental Considerations and Reagent Solutions

Successful multi-omics studies depend on careful experimental design and the selection of appropriate reagents and platforms. The following table details key materials and their functions, particularly for single-cell and spatial multi-omics approaches.

Table 3: Essential Research Reagents and Platforms for Multi-Omics

Research Reagent / Platform	Function in Multi-Omics Workflow
10x Genomics Single Cell Multiome ATAC + Gene Expression	Enables concurrent profiling of chromatin accessibility (scATAC-seq) and transcriptome (scRNA-seq) from the same single nucleus, generating matched data for vertical integration.
CITE-seq (Cellular Indexing of Transcriptomes and Epitopes by Sequencing)	Antibodies conjugated to oligonucleotide barcodes allow for the simultaneous measurement of surface protein abundance and transcriptome in single cells.

Research Reagent / Platform	Function in Multi-Omics Workflow
Antibodies	
Visium Spatial Gene Expression Slide & Reagents (10x Genomics)	Captures the whole transcriptome from tissue sections while retaining spatial location information, forming a key dataset for spatial integration with histology or other omics.
Cell Hashing Antibodies	Allows for sample multiplexing, where cells from different donors or conditions are tagged with unique barcoded antibodies, enabling pooled processing and later demultiplexing, reducing batch effects.
Single-Cell Barcoded Beads	Microbeads containing unique oligonucleotide barcodes are used to label individual cells' RNA or DNA, enabling downstream sequencing and attribution of reads to their cell of origin.
DISCO Database	A public data repository that aggregates single-cell and spatial omics data, providing a foundational resource for pretraining foundation models and benchmarking integration tools [60].

Challenges and Future Directions

Despite significant progress, the field of multimodal omics integration continues to face several formidable challenges. **Technical variability** across different experimental platforms and batches remains a major obstacle, often introducing noise that can confound true biological signals [60]. This "batch effect" can propagate through analysis pipelines and is particularly problematic when using transfer learning with foundation models [60]. A second critical challenge is **limited model interpretability**; while deep learning models can achieve high predictive accuracy, understanding the underlying biological mechanisms for their predictions—the "why" behind the result—is often difficult [60] [3]. Finally, there is a significant gap in **translating computational insights into clinical applications**. Bridging this gap requires models that are not only accurate but also robust, reliable, and clinically actionable.

Future progress will likely be driven by several key initiatives. There is a pressing need for **standardized benchmarking** of integration methods and foundation models to allow for fair comparisons and guide tool selection [60]. The development of **multimodal knowledge graphs** that systematically incorporate prior biological knowledge can help ground model predictions in established biology and improve interpretability [60]. Furthermore, fostering **collaborative frameworks** that facilitate decentralized data analysis and model sharing, similar to the Hugging Face platform in natural language processing, will be crucial for accelerating innovation and ensuring reproducibility [60]. As these technical and collaborative hurdles are overcome, multimodal integration, powered by sophisticated foundation models, will increasingly bridge the gap between cellular omics data and actionable biological understanding, ultimately paving the way for new discoveries in precision medicine.

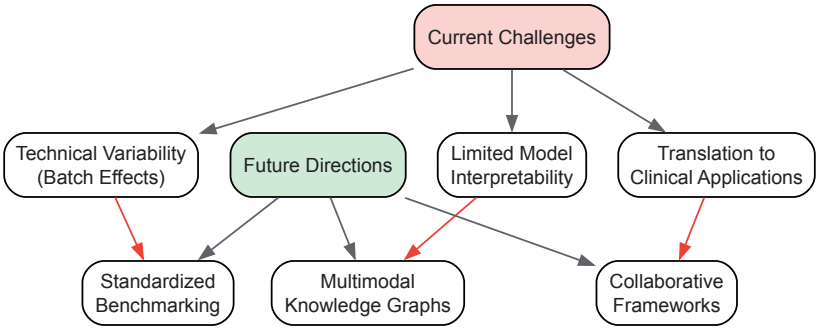

Navigating the Challenges: Data, Interpretability, and Model Biases

Addressing Data Scarcity and Noise in High-Throughput Biological Data

The advent of high-throughput sequencing and multi-omics technologies has generated biological data at an unprecedented scale and complexity, creating new opportunities for foundation models in bioinformatics [8] [9]. However, two fundamental challenges persistently hinder model development and biological discovery: data scarcity and technical noise. Data scarcity arises because experimental functional characterization remains laborious, with many critical datasets encompassing only hundreds to thousands of curated examples—insufficient for training data-hungry deep learning models [62] [63]. Simultaneously, technical noise from library preparation, sequencing stochasticity, and batch effects obscures biological signals, particularly for low-abundance molecules [64] [65].

Within the context of foundation model development, these challenges become particularly acute. Foundation models require massive, high-quality datasets for pre-training, yet biological data often exhibits sparsity, non-uniform distribution across protein families, and high noise-to-signal ratios

[8] [9]. This technical brief examines cutting-edge computational strategies that address these dual challenges, enabling more robust biological insights and facilitating the development of more accurate predictive models in bioinformatics.

Physics-Informed Machine Learning to Overcome Data Scarcity

Core Methodology and Implementation

Physics-informed machine learning represents a paradigm shift for problems with limited functional data. This approach integrates physical principles with data-driven modeling, creating hybrid systems that leverage both domain knowledge and statistical learning. The methodology involves:

- **Feature Engineering via Physical Modeling**: Instead of relying solely on sequence-based features, physics-informed approaches compute energetic and dynamic properties from molecular modeling and simulation. For instance, in studying big potassium (BK) channels, researchers quantified mutational effects on both open and closed states using molecular dynamics simulations and Rosetta mutation calculations [62].
- **Integration with Statistical Learning**: These physics-derived descriptors serve as input features for machine learning models. Physical properties including energy changes, solvent accessibility, and dynamic fluctuations complement sequence-based features such as evolutionary conservation and amino acid physicochemical properties [62].
- **Model Training and Validation**: Random forest models have demonstrated particular success in this domain, capable of learning transferable relationships between physical descriptors and functional outcomes even with sparse data [62].

Table 1: Quantitative Performance of Physics-Informed ML for BK Channel Gating Prediction

Model Type	Training Data Size	Correlation Coefficient (R)	RMSE	Key Features
Physics-Informed Random Forest	473 mutations	~0.7	~32 mV	Energetic effects, dynamic properties from MD simulations [62]
Physics-Informed Model (Novel Mutations)	4 novel mutations	0.92	18 mV	Physical descriptors from open/closed state calculations [62]

Experimental Protocol and Workflow

The typical workflow for implementing physics-informed machine learning involves:

- **Structure Preparation**: Obtain high-resolution structures of relevant biological states (e.g., open/closed conformations of ion channels).
- **Molecular Dynamics Simulations**: Perform all-atom MD simulations to sample conformational dynamics and derive dynamic properties.
- **Energetic Calculations**: Compute mutational effects on stability and interactions using tools like Rosetta or FoldX.
- **Feature Compilation**: Combine physics-derived features with sequence-based and evolutionary features.
- **Model Training**: Train machine learning models (e.g., random forests) using the compiled features and available functional data.
- **Validation**: Test model predictions on held-out data and novel experimental mutations [62].

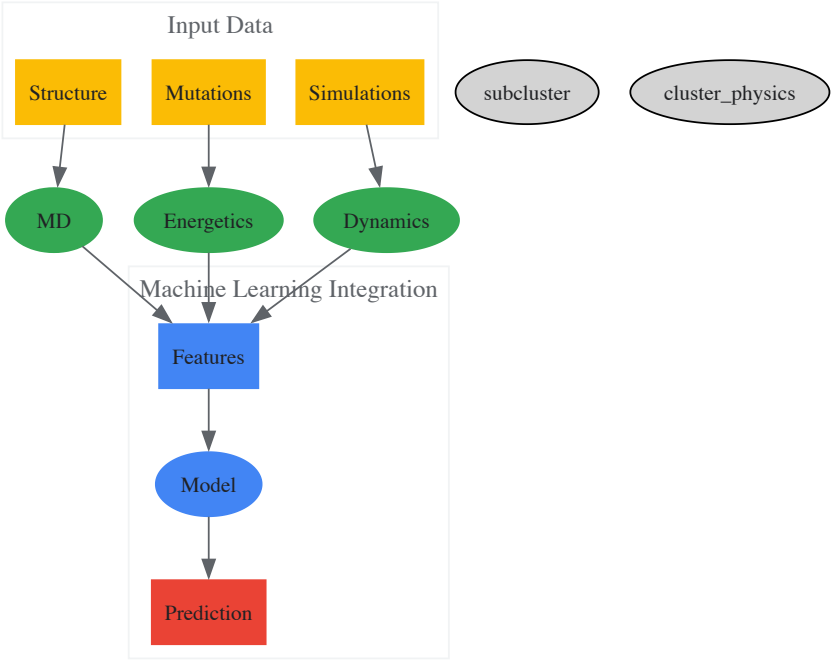

Figure 1: Physics-Informed ML Workflow for Scarce Data

Advanced Denoising Algorithms for Biological Data

Network Filtering Approaches

Network filters provide a powerful methodology for reducing technical noise in large-scale biological data by leveraging interaction networks to identify groups of correlated measurements. The core principle involves using known biological relationships to distinguish signal from noise [66].

Implementation Framework:

- **Assortative Filtering**: For correlated measurements, smoothing filters replace node values with the mean or median of neighboring values: (f{\bullet, 1}[i,{\mathbf{x }}, G] = \frac{1}{1+k{i}} \left({xi + \sum{j \epsilon \nu{i}}x{j}w_{ij}} \right)) [66]
- **Disassortative Filtering**: For anti-correlated measurements, sharpening filters enhance contrast: (f{\circ}[i,{\mathbf{x }}, G] = \alpha(x{i} - f_{\bullet, 1}[i, {\mathbf{x }}, G]) + \bar{\mathbf{x}}) [66]
- **Modular Filtering**: Networks are first partitioned into communities using detection algorithms, then appropriate filters are applied to each module based on its correlation pattern [66].

Table 2: Network Filter Performance on Biological Data Tasks

Application Domain	Filter Type	Performance Improvement	Key Metric
Protein Expression Prediction	Patchwork Filter	Up to 43% accuracy increase	Prediction accuracy vs. unfiltered data [66]
Bulk RNA-seq Analysis	Correlation-based	Improved DE detection consistency	Convergence of DE calls across methods [64]
Single-cell RNA-seq	Modular Network Filter	Enhanced rare cell type detection	Cluster separation and marker gene identification [66]

Comprehensive Noise Reduction in Single-cell Data

Single-cell technologies introduce unique noise challenges due to low starting material and amplification biases. RECODE (Resolution-Enhancement via Computational DEnoising) represents a state-of-the-art approach that simultaneously reduces technical and batch noise while preserving full-dimensional data [65].

Experimental Protocol for Single-cell Denoising:

- **Data Preprocessing**: Quality control, normalization, and initial dimensionality reduction.
- **Noise Modeling**: Characterize technical variance using spike-in controls or UMIs.

- **Batch Effect Correction**: iRECODE function identifies and removes systematic biases across datasets.
- **Signal Reconstruction**: Rebuild expression matrices with reduced noise while maintaining biological variability.
- **Validation**: Assess performance via clustering improvement, marker gene detection, and trajectory inference accuracy [65].

The algorithm has been extended to diverse single-cell modalities including single-cell Hi-C and spatial transcriptomics, demonstrating broad applicability across epigenomic and spatial domains [65].

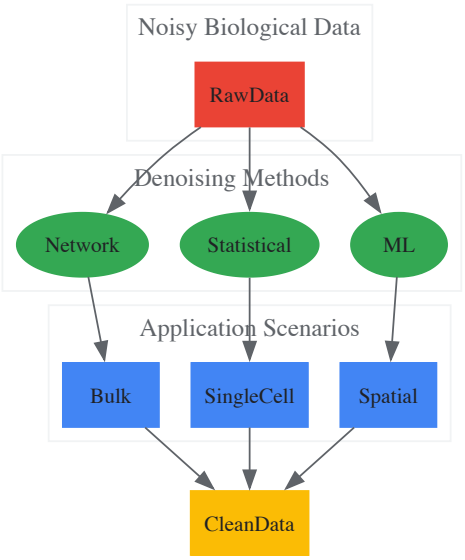

Figure 2: Biological Data Denoising Framework

Additional Computational Strategies for Limited Data

Transfer Learning and Data Augmentation

In drug discovery and functional genomics, several specialized techniques address data scarcity:

Transfer Learning: Pre-training models on large, general biomolecular datasets followed by fine-tuning on specific tasks with limited data. This approach is particularly valuable for predicting molecular properties and de novo drug design [63].

Data Augmentation: Generating modified versions of training examples to artificially expand datasets. In image-based biological data, this includes rotations, blurs, and contrast adjustments. For molecular data, careful structure-preserving transformations maintain biological validity [63].

Active Learning: Iteratively selecting the most valuable unlabeled data points for experimental characterization to maximize model improvement with minimal additional data [63].

Foundation Model Integration

Foundation models pretrained on massive biological corpora provide powerful starting points for downstream tasks with limited data. These models capture fundamental biological principles during pre-training, which can be transferred to specific applications through fine-tuning [8] [9]. Key advantages include:

- **Multi-scale Representations**: Capturing sequence, structure, and functional relationships across biological scales
- **Few-shot Learning Capabilities**: Adapting to new tasks with minimal examples
- **Multi-modal Integration**: Combining information across genomic, transcriptomic, and proteomic domains [9]

Essential Research Reagents and Computational Tools

Table 3: Research Reagent Solutions for Data Scarcity and Noise Challenges

Tool/Reagent	Type	Primary Function	Application Context
Rosetta	Software Suite	Protein energy calculation & design	Physics-based feature generation for ML [62]
GROMACS	Molecular Dynamics	Simulation of biomolecular systems	Deriving dynamic properties for features [62]

Tool/Reagent	Type	Primary Function	Application Context
noisyR	R Package	Comprehensive noise filtering	Bulk and single-cell RNA-seq denoising [64]
RECODE/iRECODE	Algorithm	Technical and batch noise reduction	Single-cell RNA-seq, Hi-C, spatial data [65]
AlphaFold2	AI System	Protein structure prediction	Structural feature generation for functional prediction [8]
DeepVariant	AI Tool	Genetic variant calling	Accurate mutation detection from NGS data [67]

Addressing data scarcity and noise in high-throughput biological data requires integrative approaches that combine physical modeling, network biology, and sophisticated machine learning. Physics-informed feature generation enables predictive modeling even with limited functional data, as demonstrated by the accurate prediction of BK channel gating properties using only 473 mutational measurements [62]. Simultaneously, network filtering and specialized denoising algorithms like RECODE significantly enhance signal-to-noise ratios across diverse data modalities [66] [65].

For foundation models in bioinformatics, these strategies are particularly crucial. They enable more effective pre-training on noisy real-world data and facilitate adaptation to specialized domains with limited fine-tuning examples. As biological data continues to grow in scale and complexity, the integration of physical principles, network-based denoising, and foundation model architectures will play an increasingly vital role in extracting meaningful biological insights and advancing drug discovery efforts.

The integration of **artificial intelligence (AI)**, particularly machine learning (ML) and deep learning (DL), into bioinformatics and drug discovery has revolutionized the analysis of complex biological data, from genomics to medical imaging [68]. However, this revolution has been accompanied by a significant challenge: the **inherent opacity** of high-performing AI models, often termed "black boxes" [69]. This opacity creates a critical trust gap, especially in sensitive and high-stakes domains like healthcare and drug development, where understanding the rationale behind a decision is as important as the decision itself [70] [71]. Clinicians may hesitate to rely on an AI's diagnosis without understanding its reasoning, and researchers cannot easily extract testable biological hypotheses from model outputs [72]. Consequently, **Explainable AI (XAI)** has emerged as an essential field focused on developing methods that make AI models transparent, interpretable, and trustworthy [73] [69].

Framed within the broader context of **foundation models (FMs)** in bioinformatics, the interpretability problem becomes even more pressing. Foundation models, which are large-scale models pre-trained on vast datasets and adaptable to a wide range of downstream tasks, are increasingly being applied to biological data [9] [10]. For instance, single-cell foundation models (scFMs) use transformer architectures to interpret millions of single-cell transcriptomes [10]. While these models show remarkable promise, their complexity and scale can further deepen the black-box problem. Therefore, developing strategies for "white-box" biological AI—models whose internal workings and decision-making processes are transparent—is a fundamental prerequisite for their reliable adoption in scientific discovery and clinical practice. This guide provides an in-depth examination of these strategies, catering to the needs of researchers, scientists, and drug development professionals.

The Interpretability Challenge in Bioinformatics

The Black-Box Problem and Its Consequences

The "black-box" problem refers to the difficulty in understanding the internal logic and the reasoning behind the predictions made by complex AI models such as deep neural networks and ensemble methods [69]. In bioinformatics, this lack of transparency poses several concrete risks:

- **Impeded Scientific Discovery**: The primary goal of bioinformatics is to gain biological insights. A black-box model might accurately predict, for example, a protein's function or a patient's disease risk, but if it cannot reveal *which features* (e.g., genes, genomic variants, or structural motifs) drove that prediction, it offers little value for hypothesis generation and mechanistic understanding [68] [72].
- **Erosion of Clinical Trust**: In medicine, a clinician is ultimately responsible for diagnostic and treatment decisions. As noted by one clinician, there is a preference for older, less accurate but more trustworthy models over newer, more accurate black boxes when making life-or-death decisions like prioritizing patients for organ transplantation [72]. Regulatory frameworks, such as those from the US FDA, also emphasize the need for healthcare professionals to independently review the basis of an algorithm's recommendations [71].
- **Perpetuation of Bias**: Without transparency, it is challenging to identify and mitigate biases that an AI model may have learned from the training data. This could lead to skewed predictions that perform poorly on underrepresented populations or specific biological conditions [72].

Foundational Concepts: Interpretability vs. Explainability

While often used interchangeably, the terms "interpretability" and "explainability" have nuanced distinctions that are important in a technical context. **Interpretability** is the ability to understand the cause-and-effect relationships within a model's inputs and outputs, often without necessarily comprehending its internal mechanics. It is concerned with the intuition behind the model's decisions [69]. **Explainability**, on the other hand, involves providing a deeper understanding of the internal logic and processes of the AI model, often through post-hoc techniques that elucidate how the model arrived at a specific output [69]. In essence, an interpretable model allows you to see *what* features were important, while an explainable model helps you understand *why* they were important. White-box models aim to achieve both.

White-Box Strategy 1: Inherently Interpretable Model Design

The most straightforward strategy for achieving interpretability is to use models that are transparent by design, known as **ante-hoc** or **self-explainable** models [70] [74]. These models provide intrinsic interpretability due to their simple structures.

Key Model Types and Architectures

- **Linear/Logistic Regression**: These are among the most interpretable models. The prediction is a weighted sum of the input features. The **coefficients** of the model directly indicate the **direction and magnitude** of each feature's influence on the output, making the reasoning process transparent [74].
- **Decision Trees and Rule-Based Models**: These models make predictions through a series of hierarchical, human-readable `IF-THEN` rules. The path from a root node to a leaf node provides a clear explanation for any given prediction, showing which features and thresholds were used [70] [69].
- **Self-Explaining Neural Networks**: Emerging architectures are designed to be inherently interpretable. For example, some models use **knowledge-primed neurons (KP Neurons)** where specific neurons are designed to correspond to biologically meaningful concepts, allowing researchers to trace a prediction back to known biological pathways or entities [68].

Table 1: Comparison of Inherently Interpretable (White-Box) Models

Model Type	Interpretability Mechanism	Advantages	Limitations	Typical Bioinformatics Applications
Linear/Logistic Regression	Feature coefficients	Simple, global interpretability, fast inference	Assumes linear relationship, cannot capture complex interactions	Preliminary feature selection, clinical risk score development
Decision Trees	IF-THEN rule paths	Intuitive visual representation, handles mixed data types	Can become large and complex, prone to overfitting	Classifying cell types from gene expression, patient stratification
Rule-Based Systems	Symbolic logic rules	Highly transparent, easily validated by experts	Requires expert knowledge to define rules, inflexible	Diagnostic rule engines, knowledge bases for molecular interactions
Self-Explaining Neural Networks	Concept-activation vectors	Balances performance and interpretability, integrates domain knowledge	Complex to design and train	Linking model predictions to known biological pathways [68]

Experimental Protocol for Validating Interpretable Models

When employing an inherently interpretable model, validation must go beyond predictive accuracy to assess the quality of the explanations.

- **Model Training**: Train the chosen white-box model (e.g., a decision tree) on the biological dataset (e.g., gene expression data for cancer subtype classification).
- **Feature Importance Extraction**: Directly extract the model's internal explanation. For a decision tree, this involves analyzing the features used at the splits near the root. For linear regression, it involves ranking the absolute values of the coefficients.
- **Biological Plausibility Assessment**: This is a critical, domain-specific step. The top features identified by the model must be evaluated by domain experts against established biological knowledge (e.g., checking if genes identified as important are known oncogenes or are involved in relevant pathways).
- **Hypothesis Generation and Testing**: Use the model's explanations to form new biological hypotheses. For instance, if an under-studied gene is consistently identified as important, this could warrant further *in vitro* or *in vivo* experiments to validate its role.

The following workflow diagram illustrates this protocol for validating an interpretable model:

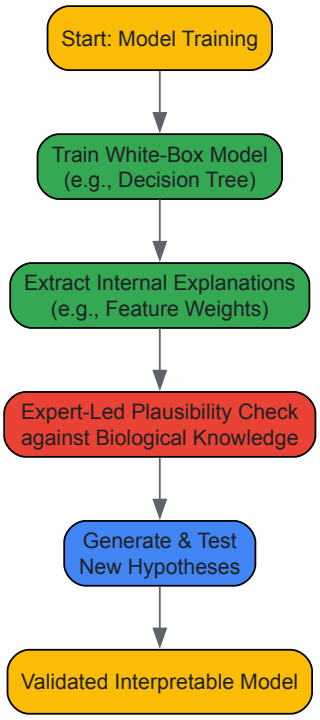

White-Box Strategy 2: Post-Hoc Explanation Techniques

For situations where complex, high-performance models like deep learning are necessary (a common scenario with foundation models), **post-hoc explanation techniques** are required. These methods analyze a trained model after the fact to approximate and explain its behavior [70]. They are categorized as either **model-specific** (designed for a particular model architecture) or **model-agnostic** (can be applied to any model) [68].

Model-Agnostic Methods

These methods treat the underlying model as a black box and probe it by analyzing input-output pairs.

- **SHAP (SHapley Additive exPlanations)**: SHAP is a unified framework based on cooperative game theory that assigns each feature an **importance value** for a particular prediction [73] [68]. It is widely used in bioinformatics for tasks like identifying critical genes in transcriptomic data [68] or important molecular features in drug discovery [73]. SHAP provides both local (per-prediction) and global (whole-model) interpretability.
- **LIME (Local Interpretable Model-agnostic Explanations)**: LIME explains individual predictions by creating a local, interpretable model (like a linear model) that approximates the complex model's behavior in the vicinity of the instance being explained [73] [68]. It is particularly useful for explaining "edge cases" in a dataset.

Model-Specific Methods

These methods leverage the internal architecture of specific models, particularly deep neural networks.

- **Attention Mechanisms**: In transformer-based foundation models, the **attention scores** can be directly interpreted to reveal which parts of the input sequence the model "paid attention to" when making a prediction [10] [68]. For example, in a single-cell foundation model, attention scores can indicate which genes were most influential in assigning a cell type. In a protein sequence model, they can highlight functionally critical residues [68].
- **Gradient-based Methods (e.g., Grad-CAM)**: Grad-CAM (Gradient-weighted Class Activation Mapping) uses the gradients of a target concept (e.g., a specific disease class) flowing into the final convolutional layer of a neural network to produce a coarse localization map highlighting important regions in the input image [70]. This is extensively used in bioimaging to highlight regions in a medical scan (e.g., chest X-ray) that led to a diagnosis [70] [68].

Experimental Protocol for Applying Post-Hoc XAI

A standardized protocol for applying post-hoc explanations to a foundation model in bioinformatics ensures robust and reliable interpretations.

- **Model and Data Preparation**: Train or fine-tune a foundation model (e.g., a transformer model for single-cell RNA-seq data) on a specific downstream task like cell type annotation or disease state prediction [10].

- **Explanation Generation**: Apply one or more XAI techniques (e.g., SHAP, Attention Scores) to the model's predictions. For a given cell's predicted type, use SHAP to get a list of gene importances or extract the model's attention weights over the input gene tokens.
- **Explanation Aggregation and Visualization**: Aggregate explanations across many instances to identify globally important features. Visualize SHAP summary plots or generate attention heatmaps across sequences to discern patterns.
- **Biological Validation and Interpretation**: Correlate the top features identified by the XAI method with known biological databases (e.g., gene ontology, protein-protein interaction networks) to assess plausibility and generate insights.

The following workflow visualizes this multi-stage protocol:

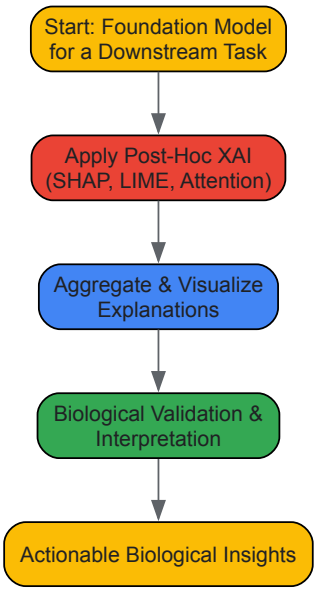

Table 2: Summary of Prominent Post-Hoc XAI Techniques in Bioinformatics

Method	Type	Core Principle	Output	Example Applications in Bioinformatics
SHAP [73] [68]	Model-Agnostic	Shapley values from game theory assigns each feature a fair contribution to the prediction.	Local and global feature importance scores.	Identifying key genes in transcriptomic data [68], predicting antiviral peptides [73], risk prediction for diabetic retinopathy [70].
LIME [73] [68]	Model-Agnostic	Fits a local surrogate model around a single prediction to approximate the black box.	Local feature importance for a single instance.	Explaining individual image classifications in bioimaging [68].
Attention Scores [10] [68]	Model-Specific	Uses the internal attention weights of transformer models.	Heatmaps showing input tokens (e.g., genes, residues) the model focused on.	Interpreting single-cell foundation models (scFMs) [10], identifying critical motifs in biological sequences [68].
Grad-CAM [70] [68]	Model-Specific	Uses gradients from a target class flowing back to a convolutional layer.	Heatmaps highlighting salient regions in an input image.	Visualizing decision process in chest X-ray classification [70], breast cancer segmentation [70].
Layer-wise Relevance Propagation (LRP) [68]	Model-Specific	Propagates the prediction backward through the network using conservation rules.	Relevance scores assigned to each input feature.	Interpreting models for gene expression data analysis [68].

The Scientist's Toolkit: Essential Reagents for XAI Experiments

To effectively implement the strategies outlined above, researchers require a suite of computational tools and resources. The following table details key "research reagents" for conducting XAI experiments in bioinformatics.

Table 3: Research Reagent Solutions for XAI in Biology

Reagent / Tool	Type / Category	Primary Function	Relevance to White-Box AI
SHAP Python Library	Software Library	Computes Shapley values for any ML model.	The go-to tool for model-agnostic feature attribution, applicable to everything from linear models to complex foundation models [73] [68].
Captum	Software Library	A PyTorch library for model interpretability.	Provides a wide range of model-specific and model-agnostic algorithms (including Grad-CAM, LIME, LRP) for interpreting PyTorch deep learning models.
Transformer Models (e.g., scBERT, scGPT)	Foundation Model	Pre-trained architectures for biological sequence (e.g., gene, protein) analysis.	Serves as the base model for many bioinformatics tasks. Their inherent attention mechanisms provide a direct path for model-specific interpretability [10].
Annotated Biological Databases (e.g., CZ CELLxGENE, GO, PDB)	Data Resource	Provide ground-truth annotations for genes, proteins, cells, and pathways.	Critical for the "biological validation" step. XAI outputs (important genes/features) are cross-referenced with these databases to assess plausibility and generate meaning [10].
TensorBoard	Visualization Tool	A suite of web applications for inspecting and understanding ML model runs.	Enables visualization of model graphs, feature embeddings, and attention weights, which is crucial for debugging and interpreting model internals.

The journey from black-box to white-box biological AI is not merely a technical challenge but a foundational requirement for the credible and productive integration of AI into the life sciences. As foundation models become more prevalent in bioinformatics, the strategies discussed—ranging from the use of inherently interpretable models to the sophisticated application of post-hoc explanation techniques like SHAP and attention analysis—provide a robust toolkit for researchers. By systematically implementing these strategies and rigorously validating explanations against biological knowledge, scientists and drug developers can bridge the trust gap. This will transform AI from an inscrutable predictor into a powerful, collaborative partner that not only makes accurate predictions but also delivers profound, actionable insights into the complex mechanisms of biology and disease.

Mitigating Model Hallucinations and Ensuring Output Reliability

The integration of foundation models into bioinformatics and drug discovery represents a paradigm shift, offering unprecedented capabilities for analyzing complex biological systems. However, the propensity of these models to generate hallucinations—content that is factually incorrect or unfaithful to source data—poses a significant risk to scientific integrity and therapeutic development. In clinical and research contexts, model hallucinations can manifest as fabricated laboratory values, invented medical conditions, or incorrect biological associations, potentially misleading research directions and compromising drug safety [75]. As foundation models become increasingly embedded in bioinformatics workflows, from target identification to clinical trial design, establishing rigorous frameworks for mitigating hallucinations and ensuring output reliability becomes paramount for maintaining scientific rigor and accelerating responsible innovation.

Defining and Classifying Hallucinations in Scientific AI

In the context of foundation models for bioinformatics, hallucination refers to the generation of content that deviates from established scientific knowledge or contradicts input data. Researchers have established precise categorizations to understand and address this phenomenon:

- **Intrinsic vs. Extrinsic Hallucinations**: Intrinsic hallucinations directly contradict information provided in the input source, while extrinsic hallucinations contain factual errors that cannot be verified against the input but may be incorrect based on external knowledge [76]. In bioinformatics, an intrinsic hallucination might involve misrepresenting experimental data presented in a query, while an extrinsic hallucination could involve inventing non-existent protein-protein interactions.
- **Factuality vs. Faithfulness Hallucinations**: Factuality hallucinations involve inconsistencies with real-world scientific facts, including factual contradictions (conflicting with known facts) and factual fabrications (inventing unsupported facts) [76]. Faithfulness hallucinations demonstrate inconsistencies with user instructions, context, or internal logic, including instruction inconsistency (deviating from original scientific query), context inconsistency (contradicting provided experimental context), and logical inconsistency (containing internal scientific contradictions) [76].

The susceptibility of models to hallucination varies significantly. Recent studies testing six large language models with physician-validated clinical vignettes containing fabricated details found hallucination rates ranging from 50% to 82% across different models and prompting conditions [75]. These findings highlight the pervasiveness of the challenge in scientific domains.

Quantifying the Hallucination Risk: Experimental Evidence

Rigorous empirical studies provide critical insights into the prevalence and patterns of hallucination across different models and conditions. A comprehensive 2024 study evaluated six LLMs using 300 physician-validated clinical vignettes, each containing a single fabricated detail (laboratory test, physical/radiological sign, or medical condition) [75]. The study design involved presenting each vignette in short (50-60 words) and long (90-100 words) versions with identical medical content, testing models under default conditions, with mitigation prompts, and with temperature adjustments, generating 5,400 total outputs.

Table 1: Hallucination Rates Across Model Types and Conditions

Model Condition	Overall Hallucination Rate	Best Performing Model (GPT-4o)	Worst Performing Model (Distilled-DeepSeek-R1)
Default Settings	66%	53%	82%
With Mitigation Prompt	44%	23%	62%
Temperature = 0	No significant improvement	No significant improvement	No significant improvement

The experimental protocol employed automated classification with physician validation, defining hallucination as any response that elaborated on, endorsed, or treated the fabricated element as real [75]. Key findings demonstrated that prompt-based mitigation significantly reduced overall hallucination rates from 66% to 44% ($p < 0.001$), with the best-performing model (GPT-4o) showing reduction from 53% to 23% ($p < 0.001$) [75]. Temperature adjustments offered no statistically significant improvement, and short vignettes showed slightly higher odds of hallucination [75].

Theoretical Foundations of Hallucination Risk

Formal theoretical frameworks help contextualize the empirical findings on model hallucination. Learning-theoretic approaches, including PAC-Bayes and Rademacher complexity, allow researchers to derive bounds on hallucination risk by treating it as a generalization error [77]. This theoretical conceptualization defines a *hallucination risk* for models, distinguishing between the inherent capacity for hallucination (based on model architecture and training) and the realized instances of hallucination in outputs [77].

The theoretical perspective explains why models with similar performance on benchmark tasks can exhibit dramatically different hallucination rates in scientific applications. It further accounts for why increasing model size or training data does not automatically eliminate hallucinations, as the phenomenon stems from fundamental limitations in how models capture and represent knowledge, particularly in specialized domains with complex, structured relationships like biology and chemistry [77] [78].

Methodologies for Hallucination Mitigation

Retrieval-Augmented Generation (RAG)

Retrieval-augmented generation enhances foundation models by incorporating real-time information retrieval from external scientific databases and knowledge sources, reducing reliance on parametric knowledge alone [76]. The standard RAG workflow comprises two critical phases:

- **Retrieval Phase**: The input question is processed as a query to a retriever that accesses external resources such biomedical databases (e.g., PubChem, ChEMBL), document corpora, or knowledge bases using dense or sparse retrieval methods [76]. The retriever returns the top-k most relevant text segments based on query-document relevance.
- **Generation Phase**: The retrieved documents are combined with the original question and passed to the generation model to produce a scientifically-grounded response [76].

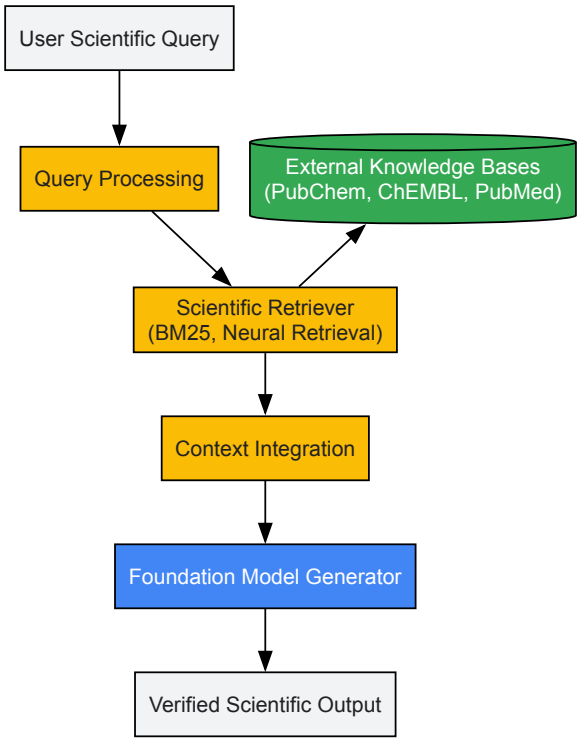

Advanced Prompt Engineering Strategies

Prompt engineering represents a crucial intervention point for reducing hallucinations in scientific applications. The 2024 clinical vignette study demonstrated that a targeted mitigation prompt could reduce hallucination rates by nearly half [75]. Effective prompt strategies include:

- **Instruction-Based Constraints**: Explicitly instructing models to use only clinically validated information and acknowledge uncertainty instead of speculating [75]. For example: "Based only on established scientific knowledge from validated sources, describe the mechanism of action. If information is incomplete or uncertain, explicitly state the limitations."
- **Few-Shot Learning with Counterexamples**: Providing examples of both correct and hallucinated responses during inference to establish boundaries for model behavior [76].
- **Chain-of-Thought Prompting**: Requiring models to articulate reasoning steps before providing final answers, making flawed logic more detectable [76].
- **Uncertainty Calibration**: Encouraging models to qualify confidence levels and identify knowledge boundaries in their responses [77].

Table 2: Prompt Engineering Effectiveness for Hallucination Mitigation

Prompt Strategy	Mechanism of Action	Effectiveness	Implementation Complexity
Instruction-Based Constraints	Directly limits speculation	High (23-44% reduction) [75]	Low
Few-Shot Learning with Counterexamples	Teaches model to distinguish valid/invalid responses	Moderate	Medium
Chain-of-Thought Reasoning	Makes reasoning transparent for validation	Moderate-High	Medium
Uncertainty Calibration	Encourages acknowledgment of knowledge limits	Moderate	Low

Hallucination-Aware Fine-Tuning

Specialized fine-tuning approaches can reduce hallucination propensity by incorporating scientific truthfulness as an optimization objective during training. Key methodologies include:

- **Contrastive Training**: Exposing models to both correct and hallucinated examples and training to maximize the difference in likelihood scores [77].

- **Factual Reinforcement Learning**: Using reward models that prioritize factual accuracy and scientific consistency over stylistic fluency [76].
- **Domain-Specific Adaptation**: Fine-tuning general foundation models on curated scientific corpora with high factual density and established verification mechanisms [78].

The Scientist's Toolkit: Research Reagent Solutions

Implementing effective hallucination mitigation requires specialized tools and frameworks tailored to scientific domains. The following research reagents represent essential components for reliable AI-assisted bioinformatics research:

Table 3: Essential Research Reagents for Hallucination Mitigation

Reagent Solution	Function	Application Context
Retrieval-Augmented Generation (RAG) Framework	Provides real-time access to current scientific knowledge	All stages of bioinformatics research
Biomedical Knowledge Bases (PubChem, ChEMBL, ZINC)	Source of validated chemical and biological information	Target identification, compound screening [78]
Structured Output Parsers	Enforces JSON/XML formatting for automated validation	Experimental data extraction and synthesis [75]
Fact-Verification Modules	Cross-references model outputs against trusted sources	Results validation and hypothesis generation
Uncertainty Quantification Tools	Measures model confidence and identifies low-probability assertions	Risk assessment for experimental decisions
Adversarial Test Suites	Evaluates model susceptibility to scientific misinformation	Model selection and deployment readiness [75]

Experimental Protocols for Hallucination Evaluation

Adversarial Hallucination Assessment Protocol

Rigorous evaluation of model susceptibility to hallucinations requires standardized experimental protocols. The following methodology, adapted from recent clinical studies, provides a framework for systematic assessment:

Objective: To quantify model propensity to adopt and elaborate on fabricated scientific details in domain-specific prompts.

Materials:

- 300+ scientific vignettes with single fabricated elements (e.g., fictitious laboratory tests, invented biological mechanisms, or non-existent compounds)
- Test set with balanced representation of subdomains (e.g., genomics, proteomics, metabolic pathways)
- Multiple foundation models for comparative assessment
- Automated classification pipeline with expert validation

Procedure:

- **Vignette Development**: Create scientific scenarios containing one deliberately fabricated element, verified by domain experts as biologically implausible.
- **Model Testing**: Present each vignette to test models under controlled conditions (default settings, mitigation prompts, temperature adjustments).
- **Response Classification**: Automatically classify outputs as "hallucination" if models elaborate on or endorse fabricated elements, and "non-hallucination" if they express uncertainty, reject, or omit the fabricated content.
- **Expert Validation**: Have domain experts independently review a subset of classifications to verify automated labeling accuracy.
- **Statistical Analysis**: Calculate hallucination rates across conditions and models using mixed-effects logistic regression to account for repeated measures [75].

Validation Metrics:

- Overall hallucination rate by model and condition
- Odds ratios for different intervention strategies

- Inter-rater reliability between automated and expert classifications

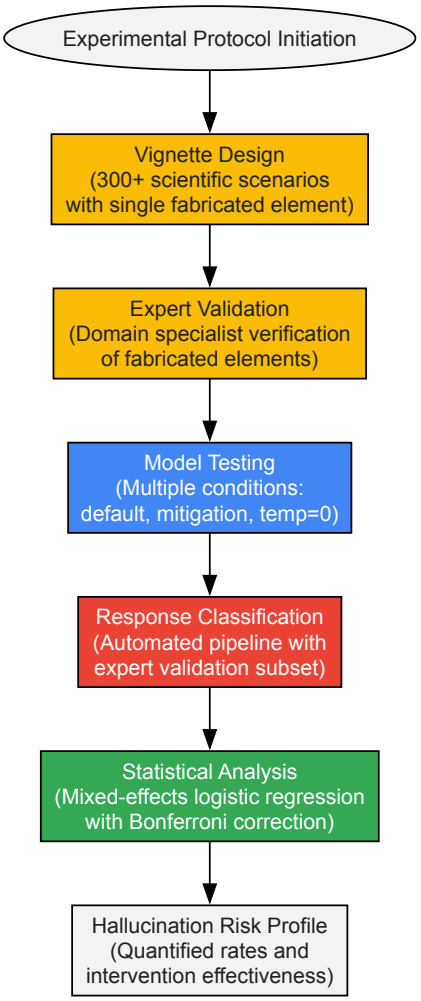

Mitigating model hallucinations and ensuring output reliability represents a critical frontier in the application of foundation models to bioinformatics and drug discovery. The experimental evidence demonstrates that while hallucination risks are substantial—affecting 50-82% of outputs in adversarial conditions—targeted interventions can reduce these rates by nearly half. A multifaceted approach combining retrieval-augmented generation, advanced prompt engineering, hallucination-aware fine-tuning, and rigorous evaluation protocols offers a pathway toward more trustworthy AI systems for scientific research. As foundation models become increasingly embedded in the drug development pipeline, from target identification to clinical trial design, establishing and maintaining rigorous standards for factual accuracy and reliability will be essential for realizing the transformative potential of AI in bioinformatics while safeguarding scientific integrity.

Overcoming Computational Intensity in Training and Fine-Tuning

In the burgeoning field of bioinformatics, foundation models have emerged as transformative tools for tasks ranging from single-cell genomics to drug discovery. These models, predominantly built on transformer architectures, leverage self-supervised learning on massive datasets to develop generalized representations that can be adapted to diverse downstream biological tasks [79] [80]. However, their remarkable performance comes with significant computational costs that present formidable barriers to widespread adoption, particularly in resource-constrained environments. The training and fine-tuning of these models demand exceptional computational resources, including high-performance GPUs with substantial memory, extensive storage systems, and sophisticated engineering infrastructure [80] [81].

The computational intensity stems from multiple factors: the massive scale of biological datasets encompassing millions of cells or genomic sequences, the inherent complexity of transformer architectures with their self-attention mechanisms, and the need for specialized preprocessing approaches to convert biological data into model-compatible formats [80]. For instance, single-cell foundation models (scFMs) process data from tens of millions of cells across diverse tissues and conditions, while genomic language models train on entire reference genomes and multi-species genomic collections [80] [81]. This article provides a comprehensive technical examination of strategies and methodologies to overcome these computational challenges, enabling more efficient and accessible implementation of foundation models across biological research domains.

Architectural and Algorithmic Strategies

Model Efficiency through Design Innovations

Innovations in model architecture represent the frontline approach to reducing computational demands. Researchers have developed several strategic modifications to standard transformer designs that substantially decrease memory requirements and computational complexity while maintaining performance.

Tokenization strategies for biological data significantly impact computational efficiency. In single-cell foundation models, rather than using all approximately 20,000 genes, models implement expression-based ranking to select the top 1,000-5,000 most informative genes per cell [80]. This selective approach reduces sequence length and consequently the computational load of the self-attention mechanism, which scales quadratically with sequence length. Similarly, DNA language models employ k-mer tokenization (typically k=3-6), where overlapping sequences of k nucleotides are treated as single tokens, dramatically shortening input sequences compared to base-by-base processing [81].

Parameter-efficient fine-tuning methods have emerged as crucial tools for adapting large pre-trained models to specific tasks without the prohibitive cost of full model retraining. Techniques including **low-rank adaptation (LoRA)**, adapter layers, and prefix tuning allow researchers to fine-tune only small subsets of parameters—often less than 1% of the total—while achieving performance comparable to full fine-tuning [82]. The VenusFactory platform implements such approaches, enabling task-specific adaptation of protein language models like ESM2 and ProtT5 with significantly reduced computational requirements [82].

Table 1: Computational Requirements of Representative Biological Foundation Models

Model	Parameter Count	GPU Memory	Primary Efficiency Strategy
DNABERT	110M	4GB+	Fixed k-mer tokenization (k=3,4,5,6)
ESM2-8M	8M	2GB+	Scalable architecture variants
ESM2-650M	650M	16GB+	Transfer learning from smaller models
Nucleotide Transformer (500M)	500M	16GB+	Multi-species pre-training
scBERT	110M	4GB+	Gene expression ranking
VenusPLM-300M	300M	12GB+	Efficient tokenization variants

Scaling and Distributed Training Methodologies

Managing the substantial computational load of training biological foundation models requires sophisticated distributed training approaches that partition the workload across multiple accelerators.

Data parallelism remains the foundational approach, where identical model replicas operate on different data batches across multiple GPUs, with gradients synchronized periodically. This method effectively scales training almost linearly with the number of available GPUs. For example, training the Nucleotide Transformer models required distributed data parallelism across dozens of GPUs to process its training corpus of 850 genomes from diverse species [81].

Model parallelism techniques address the challenge of models too large to fit within a single GPU's memory. **Tensor parallelism** splits individual model layers across devices, while **pipeline parallelism** distributes different model layers across multiple GPUs. The largest ESM2 variants with 15 billion parameters necessitate such approaches, as their memory footprint exceeds 40GB during training [82].

Mixed-precision training using 16-bit floating-point numbers (FP16) or Brain Floating Point (BF16) reduces memory usage by approximately 50% and accelerates computation by leveraging specialized tensor cores in modern GPUs. This approach maintains model accuracy through loss scaling techniques that preserve gradient information that might otherwise be lost in lower-precision representations [82].

Data-Centric Optimization Techniques

Efficient Data Representation and Preprocessing

The unique characteristics of biological data necessitate specialized preprocessing approaches that optimize computational efficiency while preserving critical biological information.

In single-cell genomics, **gene expression ranking** transforms non-sequential gene expression data into an ordered sequence that transformer architectures can process. Rather than relying on arbitrary gene ordering, models implement deterministic ranking based on expression magnitude, creating meaningful sequences while reducing dimensionality [80]. Some implementations further optimize by binning genes into expression-level categories or employing strategic downsampling of low-information genes [80].

For genomic sequences, **k-mer tokenization** with strategic overlap provides contextual information while controlling sequence length. The 6-mer approach used in fine-tuned DNA transformer models creates a vocabulary size of 4,096 possible tokens (4^6), effectively balancing contextual

information with manageable sequence length [81]. This approach reduces sequence length approximately 6-fold compared to single-nucleotide tokenization, dramatically decreasing computational requirements for attention mechanisms.

Selective sequence processing strategies further enhance efficiency. Models can implement attention mechanisms with restricted context windows or employ progressive training approaches that begin with shorter sequences before advancing to full-length processing. The ProSST protein model series exemplifies this approach with configurable sequence length support from 20 to 4,096 residues, enabling researchers to select appropriate capacity for their specific applications [82].

Data Augmentation and Quality Control

Careful data curation and augmentation significantly impact computational efficiency by reducing the need for repeated training runs and improving sample efficiency.

In clinical pathway modeling, researchers have implemented **synthetic data generation** through topic model-guided augmentation, effectively expanding training datasets and improving model robustness without additional data collection [83]. The LDA-BiLSTM framework demonstrated that strategically augmented data could improve accuracy by 22-25% while reducing training instability and the need for repeated epochs [83].

Quality filtering and **batch effect correction** are particularly crucial for single-cell data, where technical artifacts can significantly impact model performance. Implementing rigorous quality control pipelines during data preprocessing removes low-quality cells and genes, reducing noise and improving training efficiency [80]. The scFM development process emphasizes careful dataset balancing and composition to create maximally informative training corpora [80].

Table 2: Data Processing Techniques for Computational Efficiency

Technique	Application Context	Impact on Computational Efficiency
Gene expression ranking	Single-cell transcriptomics	Reduces sequence length by 75-95% (from ~20,000 to 1,000-5,000 genes)
K-mer tokenization (k=6)	Genomic sequences	6x reduction in sequence length compared to base-level processing
Strategic gene filtering	Single-cell omics	Removes low-information features, reducing dimensionality
Quality control pipelines	All biological data	Reduces noise, improves training stability and convergence
Topic model-guided augmentation	Clinical pathway data	Expands effective dataset size without additional collection costs

Experimental Protocols and Implementation Frameworks

Protocol for Efficient Fine-Tuning of Biological Foundation Models

This section provides a detailed methodology for computationally efficient adaptation of pre-trained biological foundation models to specific downstream tasks, based on established practices from recent literature.

Required Resources and Setup:

- Pre-trained foundation model (e.g., DNABERT, ESM2, scBERT)
- Task-specific labeled dataset
- GPU cluster with sufficient VRAM (minimum 12GB recommended)
- Deep learning framework (PyTorch/TensorFlow)
- Specialized libraries (Hugging Face Transformers, VenusFactory)

Procedure:

- **Data Preparation and Tokenization**
  - For genomic sequences: Implement k-mer tokenization with k=6, creating overlapping sequences with step size 1 [81]
  - For single-cell data: Apply gene ranking by expression value, selecting top 1,000-5,000 genes based on variance [80]
  - Split data into training (80%), validation (10%), and test (10%) sets maintaining biological group distributions
- **Model Configuration Setup**
  - Load pre-trained weights from established biological foundation models
  - Configure parameter-efficient fine-tuning method (adapter layers or LoRA)
  - For LoRA: Set rank parameter to 8-16, apply to query and value attention layers
  - Freeze all base model parameters except designated fine-tuning parameters

- **Training Loop Implementation**
  - Set batch size to maximum feasible given GPU memory constraints
  - Implement mixed-precision training (AMP) with FP16/BF16 precision
  - Use gradient accumulation (steps=4-8) to simulate larger batch sizes
  - Configure learning rate scheduler with linear warmup (10% of steps) followed by cosine decay
  - Employ early stopping with patience of 5-10 epochs based on validation performance
- **Evaluation and Deployment**
  - Assess model performance on held-out test set using domain-specific metrics
  - Compare against baseline models to verify efficiency gains
  - Deploy optimized model for inference with appropriate quantization if needed

This protocol was validated in DNA transformer fine-tuning experiments, where a naturally-trained sentence transformer adapted to DNA sequences achieved competitive performance with domain-specific models while requiring substantially less computational resources [81].

Workflow for Distributed Training of Large Biological Models

For scenarios requiring full model training rather than fine-tuning, the following distributed training protocol provides computational efficiency.

Procedure:

- **Infrastructure Configuration**
  - Set up multi-node GPU cluster with high-speed interconnects (InfiniBand preferred)
  - Configure distributed training environment (PyTorch DDP, Horovod)
- **Data Partitioning and Loading**
  - Implement sharded dataset loading to distribute storage I/O across nodes
  - Configure data sampling with appropriate balancing for biological classes
  - Set up data prefetching with multiple worker processes
- **Model Parallelism Setup**
  - For models >1B parameters: Implement tensor parallelism across available GPUs
  - Configure gradient checkpointing to trade computation for memory
  - Set up optimizer state partitioning (ZeRO Stage 2/3) for memory efficiency
- **Training Execution**
  - Launch distributed training job across all available nodes
  - Monitor training metrics and resource utilization across nodes
  - Implement automatic fault tolerance and recovery for long-running jobs

The VenusFactory platform exemplifies this approach, providing containerized implementations of these protocols for various protein language models including ESM series and ProtTrans models [82].

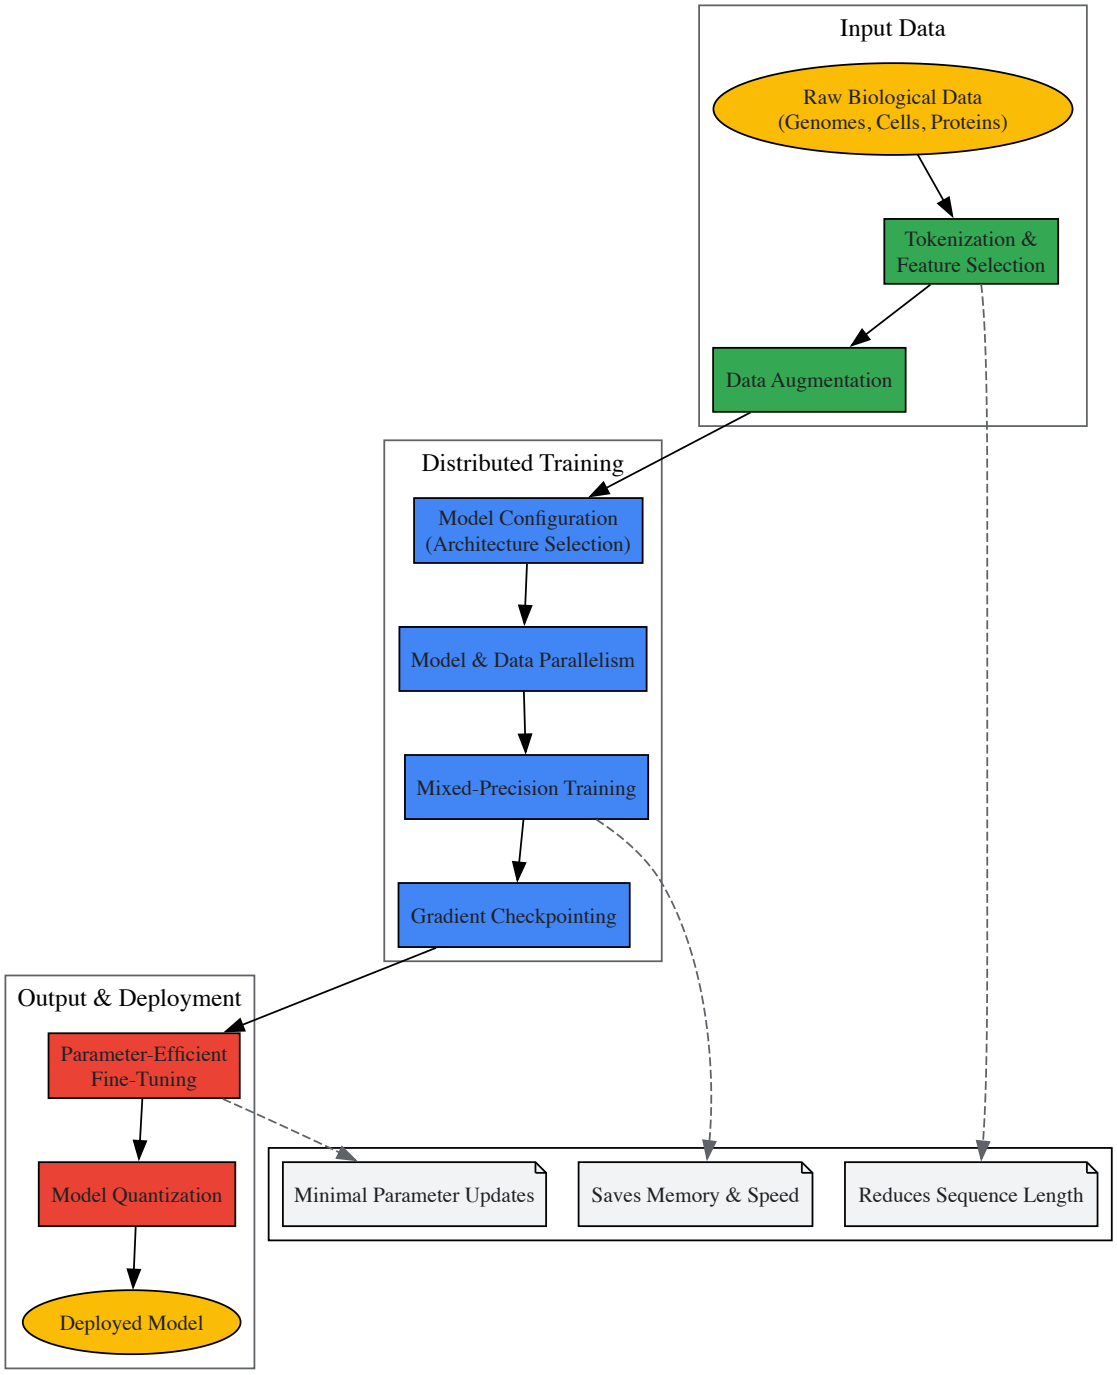

Efficient Training Workflow for Biological Foundation Models

The Scientist's Toolkit: Research Reagent Solutions

Table 3: Essential Computational Resources for Biological Foundation Models

Resource Category	Specific Tools & Platforms	Function in Computational Efficiency
Model Architectures	DNABERT, ESM2, scBERT, Nucleotide Transformer	Pre-designed architectures optimized for biological data
Training Frameworks	VenusFactory, Hugging Face, DeepSpeed	Provide optimized implementations of distributed training
Efficiency Libraries	LoRA, AdapterHub, AMP	Enable parameter-efficient fine-tuning and mixed-precision training
Data Processing Tools	CELLxGENE, Scanpy, Biopython	Standardize biological data preprocessing and tokenization

Resource Category	Specific Tools & Platforms	Function in Computational Efficiency
Computational Resources	High-memory GPUs (A100/H100), GPU clusters	Provide necessary hardware acceleration for large-scale training
Benchmarking Suites	ProteinGym, TissueNet, DNA benchmark tasks	Standardized evaluation to measure efficiency gains

The computational challenges inherent in biological foundation models are substantial but not insurmountable. Through strategic architectural modifications, sophisticated distributed training approaches, data-centric optimization techniques, and specialized experimental protocols, researchers can significantly reduce the computational burden of training and fine-tuning these powerful models. The continued development of parameter-efficient fine-tuning methods, model compression techniques, and specialized hardware for biological AI will further enhance accessibility.

As these efficiency-improving strategies mature, they promise to democratize access to cutting-edge bioinformatics tools, enabling researchers worldwide to leverage foundation models for diverse applications from single-cell analysis to drug discovery. The integration of these approaches into unified platforms like VenusFactory represents a promising direction for the field, lowering barriers to entry while maintaining state-of-the-art performance [82]. Through continued innovation in computational efficiency, foundation models are poised to become increasingly central to biological discovery and therapeutic development.

The emergence of foundation models in bioinformatics represents a paradigm shift, moving from task-specific algorithms to general-purpose models pre-trained on vast molecular datasets. These models promise to unlock profound biological insights by learning universal representations of cellular function and disease mechanisms. However, their development and application are fraught with two persistent and interconnected technical challenges: the tokenization of inherently non-sequential omics data and the management of pervasive batch effects. Tokenization, the process of converting raw molecular data into discrete model inputs, is complicated by the fact that biological features like genes lack a natural sequential order, unlike words in a sentence. Concurrently, batch effects—technical variations introduced by different experiments, protocols, or platforms—can confound biological signals and lead to misleading conclusions. This technical guide delves into the core of these challenges, providing a detailed analysis of current solutions, methodologies, and tools essential for researchers and drug development professionals working at the forefront of computational biology.

Tokenizing Non-Sequential Omics Data

The Core Challenge and Common Strategies

In natural language processing, words naturally form a sequence, providing a clear structure for tokenization. In contrast, omics data, such as a cell's gene expression profile, is fundamentally **non-sequential**. The order of genes on a microarray or in a single-cell RNA sequencing output is arbitrary and does not carry inherent biological meaning for most analytical tasks. This presents a fundamental hurdle for transformer-based architectures, which require an ordered input sequence [10].

To overcome this, researchers have developed several strategic workarounds to impose a functional order on the feature set:

- **Ranking by Expression:** A prevalent strategy involves ranking genes within each individual cell based on their expression levels. The top k highly expressed genes are then fed into the model as a deterministic sequence, creating a "sentence" representing the cell's state [10].
- **Expression Value Binning:** Another approach partitions genes into bins according to their expression values. The bin assignments and rankings are used to determine the order and position of each gene token in the input sequence [10].
- **Utilizing Gene Identifiers:** Some models forgo complex ranking and simply use normalized counts, relying on a fixed but arbitrary order of gene identifiers (e.g., alphabetically or by genomic coordinate) [10].

The table below summarizes and compares these primary tokenization strategies.

Table 1: Common Tokenization Strategies for Non-Sequential Omics Data

Strategy	Core Methodology	Key Advantages	Key Limitations
Ranking by Expression	Ranks genes within a cell by expression magnitude; uses top k genes as sequence.	Captures the most biologically relevant features for each cell dynamically.	Gene order is not consistent across cells, complicating cross-cell comparisons.
Expression Value Binning	Assigns genes to bins based on expression levels; uses binning for ordering.	Provides a structured method to handle continuous expression values.	The binning scheme is arbitrary and may not reflect true biological hierarchies.
Fixed Gene Identifier Order	Uses a predefined, fixed order of genes (e.g., alphabetical, genomic position).	Simple to implement and provides a consistent input order across all cells.	The imposed sequence is biologically meaningless and may hinder model

Strategy	Core Methodology	Key Advantages	Key Limitations
			learning.

Advanced Architectures and Tokenization Enhancements

Beyond these basic strategies, advanced models incorporate additional information to enrich the tokenization process and provide greater biological context:

- **Special Tokens:** Models often prepend special tokens to represent cell identity or metadata, enabling the model to learn cell-level context. For multi-omics integration, modality-specific tokens (e.g., `[RNA]` or `[ATAC]`) are included to distinguish the data source [10].
- **Incorporation of Metadata:** Gene-level metadata, such as Gene Ontology terms or chromosomal location, can be embedded alongside the expression value to provide a richer biological context for each token [10].
- **Dynamic Tokenization for Genomic Sequences:** For foundation models working directly with DNA sequence data, the tokenization challenge is different. DNA lacks inherent word boundaries. Innovative models like **MergeDNA** address this by using a hierarchical architecture with **differentiable token merging**. This allows the model to dynamically chunk adjacent DNA bases into variable-length "words" based on local context and information density, moving beyond fixed *k*-mer or byte-level strategies [84].

The following diagram illustrates a generalized workflow for tokenizing single-cell omics data, integrating the strategies discussed above.

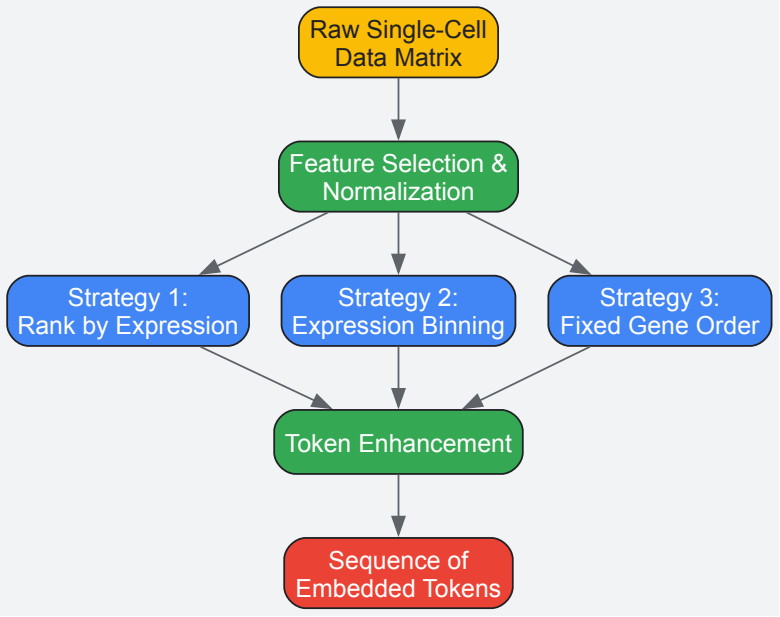

Managing Batch Effects

Batch effects are technical variations unrelated to the biological question that can severely compromise the integrity and reproducibility of omics studies. Their impact is profound:

- **Misleading Conclusions:** Batch effects can dilute biological signals, reduce statistical power, and, in the worst cases, lead to incorrect conclusions. A prominent example is a clinical trial where a change in RNA-extraction solution caused a shift in gene-based risk calculations, resulting in incorrect treatment regimens for 162 patients [85].
- **Irreproducibility:** Batch effects are a paramount factor contributing to the "reproducibility crisis" in science. They have been responsible for the retraction of high-profile articles and the failure to replicate key results, leading to economic losses and discredited research [85].

These effects originate at nearly every stage of a high-throughput study. Key sources include flawed or confounded study design, variations in sample preparation and storage, differences in reagent lots, personnel, protocols, and equipment, and inconsistencies in sequencing runs [85].

Batch Effect Correction Algorithms (BECAs)

A wide array of computational methods has been developed to mitigate batch effects. They can be broadly categorized as follows:

- **Location-Scale Adjustments:** These methods adjust the mean and/or variance of genes across batches. A seminal tool in this category is **ComBat**, which uses an empirical Bayes framework to correct for batch effects [85] [86]. Its regularized version, **reComBat**, was developed to handle the large, complex design matrices common in public data repositories by addressing the singularity problem in linear regression [86].

- **Matrix Factorization and Dimensionality Reduction:** Methods like **Harmony** and **LIGER** integrate datasets by projecting them into a shared low-dimensional space where batch differences are minimized while biological variance is preserved [87]. These are particularly powerful for complex single-cell data.
- **Deep Learning-Based Approaches:** Non-linear models, particularly variational autoencoders (VAEs) and generative adversarial networks (GANs), learn a batch-effect-free latent space representation of the data. Tools like **scGen** and **normAE** fall into this category. They offer great flexibility but often require large datasets and significant computational resources [85] [88].

The table below provides a comparative overview of these BECA categories.

Table 2: Categories of Batch Effect Correction Algorithms (BECAs)

Category	Representative Tools	Underlying Principle	Applicability
Location-Scale Adjustments	ComBat, reComBat [86]	Empirical Bayes adjustment of mean and variance per gene and batch.	Bulk and single-cell data; effective for moderate, linear batch effects.
Matrix Factorization	Harmony [87], LIGER [87], JIVE [88]	Dimensionality reduction to find a shared latent space where batches are aligned.	Single-cell and bulk data; handles complex, non-linear batch effects well.
Deep Learning Models	scGen, normAE [85], VAEs [88]	Neural networks learn non-linear mappings to a latent space invariant to batch.	Large, complex datasets; powerful but computationally intensive.
Reference-Based	Cross-platform normalization [86]	Aligns batches to a designated reference batch or set of "housekeeping" genes.	Limited by the availability of a reliable reference, especially for microbes.

Experimental Protocol: Applying the reComBat Algorithm

For researchers dealing with large-scale, multi-source public data (e.g., from GEO), where the design matrix of biological covariates can be large and highly correlated, the following protocol for applying the **reComBat** algorithm is recommended [86]:

- **Data Preprocessing and Preparation:**
  - Compile your multi-batch gene expression matrix (e.g., microarray or bulk RNA-seq).
  - Format the data into an `n x p` matrix, where `n` is the number of samples and `p` is the number of features (genes).
  - Define two design matrices:
    - **Model of Interest (X):** Contains the biological covariates and outcomes you wish to preserve (e.g., disease status, treatment group).
    - **Batch Matrix (C):** Specifies the batch membership for each sample (e.g., study ID, processing date).
- **Model Fitting and Standardization:**
  - Fit a linear model to the data: $Y = X\beta + C\alpha + \varepsilon$, where Y is the expression matrix, β and α are coefficients, and ε is the error.
  - Standardize the data by subtracting the model-predicted values and scaling by the residual standard deviation to obtain residuals with mean zero and variance one.
- **Empirical Bayes Adjustment:**
  - Estimate the prior distributions for the batch effect parameters (additive γ and multiplicative δ) from the standardized data across all genes.
  - Compute the posterior batch effect adjustments using these empirical Bayes estimates. This step "shrinks" the estimates toward the overall mean, making the correction more robust, especially for small batches.
- **Data Correction and Output:**
  - Apply the adjusted parameters to remove the estimated batch effects from the standardized data.
  - Reverse the standardization process to return the corrected data to its original scale.
  - The output is a batch-corrected gene expression matrix ready for downstream integrative analysis.

The following diagram visualizes the multi-faceted process of diagnosing and correcting for batch effects in a typical research workflow.

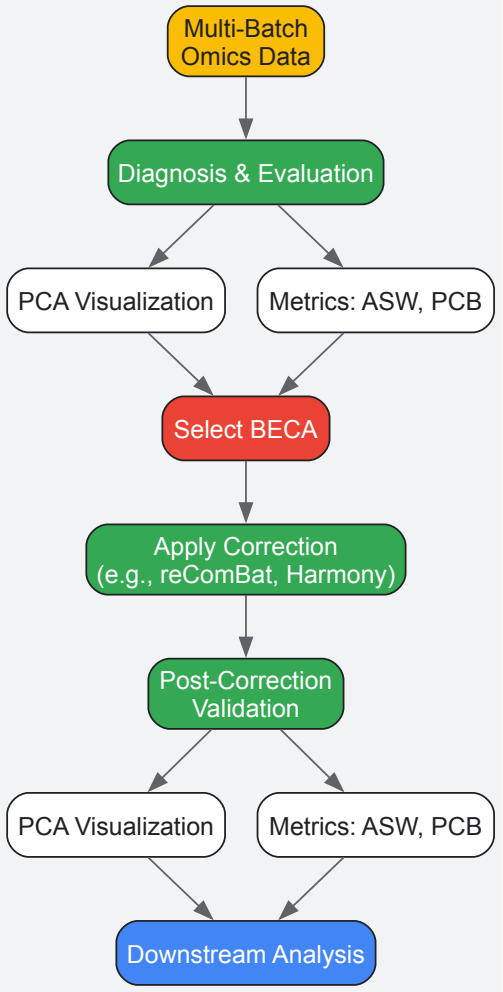

The Scientist's Toolkit

This section provides a curated list of key resources, tools, and datasets that form the essential toolkit for researchers tackling tokenization and batch effects in the development of bioinformatics foundation models.

Table 3: Research Reagent Solutions and Essential Resources

Item / Resource	Type	Function / Application	Example / Source
Curated Single-Cell Atlases	Data Source	Provides large-scale, annotated datasets essential for pre-training scFMs and benchmarking BECAs.	CZ CELLxGENE [10], Human Cell Atlas [10] [60]
Batch Effect Correction Tools	Software Tool	Algorithms to remove technical variation while preserving biological signal prior to or within model training.	Harmony [87], reComBat [86], Seurat Integration [87]
Foundation Model Hubs	Software Platform	Repositories for sharing, versioning, and deploying pre-trained foundation models, promoting reproducibility.	BioLLM [60] (interface for benchmarking >15 scFMs)
Dynamic Tokenization Models	Software Tool	Frameworks for adaptive genomic tokenization, moving beyond fixed k-mers to context-aware chunking.	MergeDNA [84]
Multi-Omics Integration Suites	Software Tool	Flexible frameworks that support multi-task learning and various integration strategies for complex data.	Flexynesis [89]

The successful implementation of foundation models in bioinformatics hinges on the effective resolution of the twin challenges of tokenization and batch effects. While significant progress has been made—evidenced by dynamic tokenization strategies like MergeDNA and robust batch-correction tools such as reComBat and Harmony—no universal solution exists. The choice of strategy is highly dependent on the specific data modality, the biological question, and the scale of integration. Future directions will likely involve the tighter coupling of these two aspects, perhaps through end-to-end trainable pipelines that jointly optimize data harmonization and representation learning. As the field marches toward ever-larger models and more ambitious multi-omic integrations, a deep and practical understanding of these technical hurdles is not just beneficial but essential for any researcher aiming to contribute to this transformative era in computational biology and precision medicine.

Benchmarking Foundation Models: Performance, Pitfalls, and Practical Selection

Rigorous Benchmarking Frameworks for Biological Task Evaluation

The rapid emergence of foundation models in bioinformatics has created an urgent need for rigorous, standardized benchmarking frameworks to evaluate their capabilities and limitations. These models, trained on massive biological datasets, promise to revolutionize everything from single-cell analysis to genomic sequence interpretation. However, without proper evaluation, their real-world utility remains uncertain. Recent studies reveal that **despite their theoretical promise**, many biological foundation models consistently **underperform well-established supervised methods** on specific tasks, highlighting the critical importance of robust benchmarking [90]. This whitepaper provides a comprehensive technical guide to current benchmarking frameworks, their methodologies, and their applications in validating foundation models for biological tasks.

The challenge lies not only in model architecture but in the fundamental nature of biological data. Unlike natural language, biological sequences lack inherent ordering, and experimental data often suffer from **batch effects, technical noise, and inconsistent processing** across studies [10]. Furthermore, the non-sequential nature of omics data presents unique challenges for transformer architectures originally designed for language, requiring specialized tokenization approaches where genes or features are treated as tokens and cells as sentences [10]. These complexities necessitate benchmarking suites that go beyond standard machine learning metrics to assess true biological relevance and practical utility.

Benchmarking Landscape for Biological Foundation Models

The table below summarizes major benchmarking frameworks currently available for evaluating foundation models across different biological domains:

Table 1: Comprehensive Benchmark Suites for Biological Foundation Models

Benchmark Name	Biological Domain	Core Tasks	Data Scale	Key Innovation
BioProBench [91]	Biological Protocols	Protocol QA, Step Ordering, Error Correction, Protocol Generation, Protocol Reasoning	27K protocols, 556K instances	Multi-task evaluation focusing on procedural understanding
DNALONGBENCH [92]	Genomics	Enhancer-target interaction, eQTL, 3D genome organization, Regulatory activity, Transcription initiation	Sequences up to 1M bp	Focus on long-range DNA dependencies
scFM Evaluation [10] [93]	Single-cell Biology	Cell type annotation, Perturbation response prediction, Gene expression prediction	10M+ single cells	Unified evaluation of single-cell foundation models

These frameworks address different aspects of a critical problem: many foundation models exhibit strong performance on surface-level tasks but **struggle significantly with deep reasoning** and structured generation. For instance, while top models achieve approximately 70% accuracy on protocol question answering, their performance drops to below 50% on tasks requiring temporal reasoning like step ordering [91]. This performance gap underscores the need for specialized benchmarks that can probe beyond superficial metrics.

Methodologies for Benchmark Construction and Evaluation

Data Curation and Processing

Robust benchmark construction begins with comprehensive data curation from diverse biological sources. For protocol-centric benchmarks like BioProBench, this involves collecting tens of thousands of full-text protocols from authoritative resources including Bio-protocol, Protocol Exchange, JOVE, Nature Protocols, and Protocols.io [91]. The curation pipeline requires **deduplication, cleaning of formatting artifacts** using regular expressions and NLP techniques, and structured extraction of key elements including title, identifiers, keywords, and operation steps. For complex nested structures with sub-steps, parsing rules based on indentation and symbol levels restore parent-child relationships, ensuring accurate representation of experimental workflows [91].

Similarly, for genomic benchmarks like DNALONGBENCH, data processing must handle extremely long sequences (up to 1 million base pairs) while maintaining biological significance. The selection criteria for such benchmarks emphasize: (1) **biological significance** - tasks must address realistic genomics problems; (2) **long-range dependencies** - tasks requiring contexts spanning hundreds of kilobase pairs; (3) **task difficulty** - posing significant challenges for current models; and (4) **task diversity** - spanning various length scales and task types including classification, regression, 1D, and 2D predictions [92].

Task Design and Instance Generation

Strategic task design creates benchmarks that comprehensively evaluate model capabilities across different cognitive levels:

- **Protocol Question Answering (PQA)** simulates information retrieval scenarios by querying critical dimensions like reagent dosages, parameter values, and operational instructions, addressing real-world ambiguities including inconsistent units and undefined concentration ranges [91].
- **Error Correction (ERR)** assesses the ability to identify and correct safety-critical errors related to reagents, parameters, and operations, simulating high-risk scenarios like volume overrides and incorrect concentrations [91].
- **Step Ordering (ORD)** evaluates understanding of protocol hierarchy and procedural dependencies through both top-step (ordering main stages) and child-step (ordering sub-steps within a stage) challenges [91].
- **Long-range Genomic Tasks** including enhancer-target gene interaction, expression quantitative trait loci (eQTL), and 3D genome organization assess model capabilities in capturing dependencies across extreme genomic distances [92].

The generation of high-quality task instances typically employs a combination of rule-based transformations and LLM-assisted generation with careful human validation. For example, in BioProBench, up to five easy, three standard, and one difficult task are generated per protocol, retaining original step numbering and prioritizing representative subtasks when source material is insufficient [91].

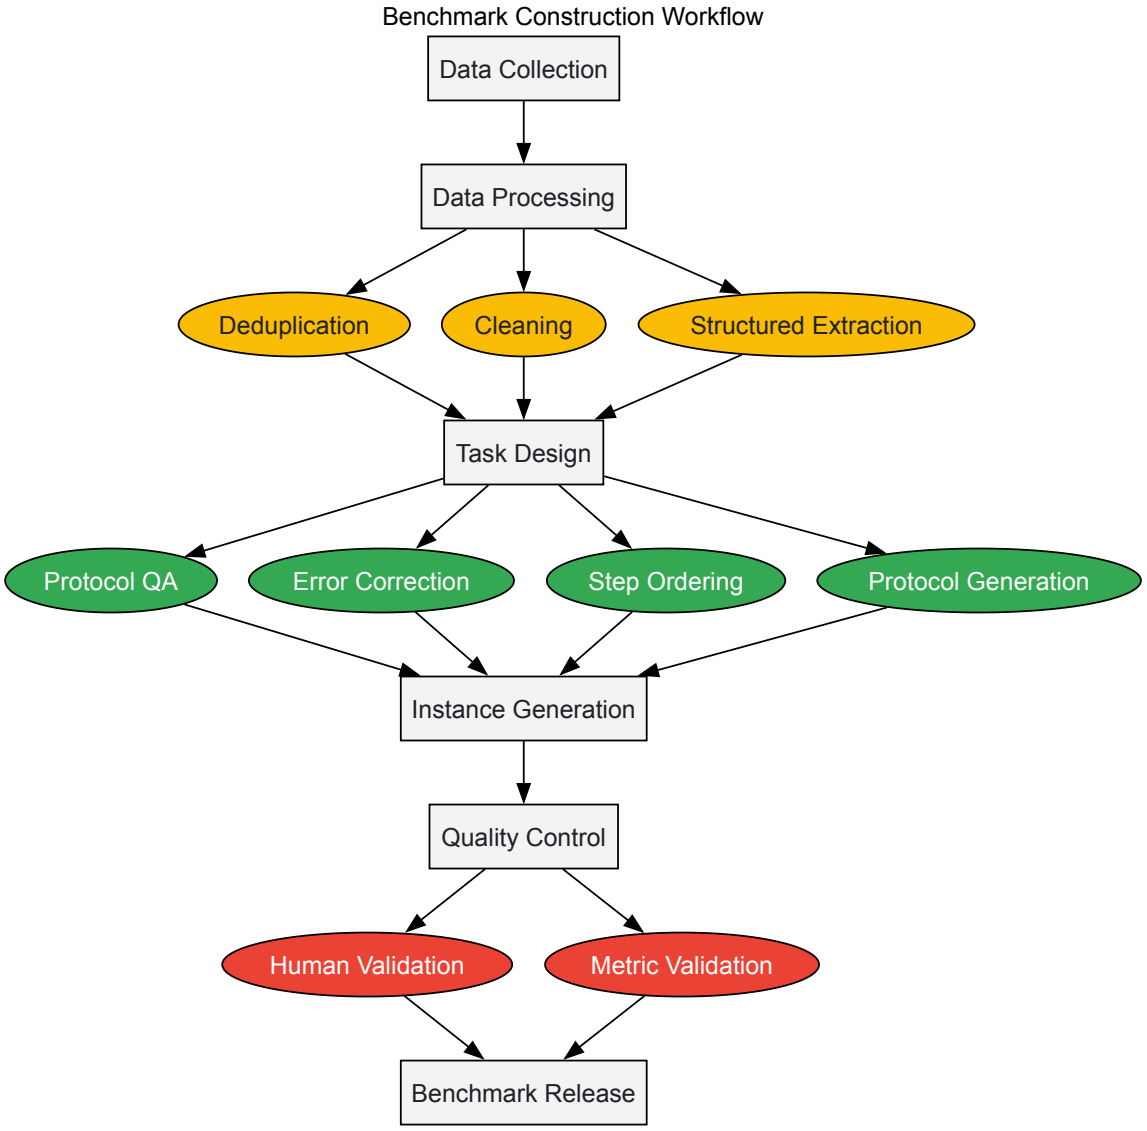

Evaluation Metrics and Baselines

Comprehensive benchmarking requires both standard NLP metrics and domain-specific measures:

Table 2: Evaluation Metrics for Biological Foundation Model Benchmarks

Metric Category	Specific Metrics	Application Context	Limitations
Standard NLP Metrics	BLEU, ROUGE, METEOR, BERTScore	Protocol generation, Text reconstruction	Capture lexical overlap but fail to assess executability

Metric Category	Specific Metrics	Application Context	Limitations
Statistical Metrics	Pearson Correlation, AUROC, SCC	Gene expression prediction, Interaction classification	May not reflect biological significance
Domain-Specific Metrics	Keyword-based content metrics, Embedding-based structural metrics	Protocol fidelity, Scientific correctness	Require careful calibration and validation
Execution-based Metrics	Step granularity, Action ordering, Semantic fidelity [94]	Protocol executability, Experimental reliability	Computationally intensive to assess

Effective benchmarking necessitates comparison against appropriate baseline models, including: (1) **Simple baselines** (e.g., training set mean); (2) **Task-specific expert models** that represent current state-of-the-art; (3) **Traditional ML models** (CNNs, Random Forests); and (4) **Other foundation models** for cross-architecture comparison [92] [93]. Surprisingly, recent studies have found that even simple baseline models sometimes outperform sophisticated foundation models. For instance, in post-perturbation RNA-seq prediction, a simple training mean baseline outperformed both scGPT and scFoundation, while Random Forest models with Gene Ontology features "outperformed foundation models by a large margin" [93].

Experimental Protocols for Benchmark Implementation

Implementing the BioProBench Evaluation Framework

The BioProBench framework provides a standardized approach for evaluating model capabilities on biological protocols. The implementation requires:

First, data preparation and partitioning according to the benchmark specifications. The 27K protocols should be split into training, validation, and test sets with approximately 70-15-15 ratio, ensuring that protocols from the same source are not disproportionately represented in any single split. For the Perturbation Exclusive (PEX) evaluation setting, ensure that specific perturbations are completely held out from training [91] [93].

Next, model adaptation to the five core tasks:

- For **Protocol Question Answering**, format input as multiple-choice questions with five options, incorporating both syntactic and semantic distractors.
- For **Error Correction**, present models with subtly modified protocol steps containing critical errors in reagents, parameters, or operations.
- For **Step Ordering**, provide shuffled steps (both top-level and child-level) and evaluate the ability to restore correct experimental temporal logic.
- For **Protocol Generation**, use the easy, standard, and difficult tasks to assess capability to generate executable protocols under increasing complexity.
- For **Protocol Reasoning**, evaluate performance on inferencing tasks requiring understanding of experimental cause-effect relationships.

Finally, execute the evaluation using the hybrid metrics framework, calculating both standard NLP metrics and domain-specific measures. The benchmark implementation should report performance disaggregated by task type and difficulty level to identify specific model weaknesses [91].

Implementing Genomic Long-Range Dependency Evaluation

The DNALONGBENCH framework evaluates model capabilities on long-range genomic dependencies using this experimental protocol:

Begin with data acquisition and preprocessing. Download the benchmark sequences from the designated repository, noting that input sequences are provided in BED format listing genome coordinates. This format allows flexible adjustment of flanking context without reprocessing. For each task, prepare the input sequences according to the specified lengths (450,000 bp for enhancer-target and eQTL tasks, 1,048,576 bp for contact map prediction, etc.) [92].

Next, implement the baseline models for comparison:

- **Lightweight CNN**: Train a three-layer CNN using cross-entropy loss for classification tasks and mean squared error loss for contact map prediction.
- **Expert models**: Reproduce the state-of-the-art specialized models for each specific task as reference points.
- **DNA foundation models**: Fine-tune pre-trained models including HyenaDNA and Caduceus variants on each specific task.

Execute the evaluation using the task-specific metrics: AUROC for classification tasks, Pearson correlation for regression tasks, and stratum-adjusted correlation coefficient for contact map prediction. Perform statistical significance testing on results across multiple runs to ensure robustness [92].

Critical Findings from Current Benchmarking Efforts

Recent benchmarking studies have revealed several critical patterns in foundation model performance:

Table 3: Performance Patterns of Foundation Models on Biological Tasks

Model Category	Strengths	Weaknesses	Representative Examples
General Foundation Models	Strong on surface-level understanding, Protocol QA	Struggle with deep reasoning, Scientific accuracy	GPT-5, Gemini-2.5-pro-exp [91]
Domain-Specific Foundation Models	Better biological knowledge, Gene embeddings	Lag on complex procedural dependencies, Limited scope	scGPT, scFoundation [10] [93]
Expert Models	State-of-the-art on specific tasks, Computational efficiency	Narrow focus, Limited transferability	Task-specific CNNs, GEARS [92]
Traditional ML with Biological Features	Strong performance, Interpretability	Feature engineering required, Limited to known biology	Random Forests with GO features [93]

The benchmarking results consistently show that foundation models face significant challenges with **structured generation and logical sequencing**. Even advanced models struggle with tasks like step ordering (achieving less than 50% exact match) and protocol generation (BLEU scores below 15), indicating fundamental limitations in procedural reasoning [91]. Furthermore, the pretraining-finetuning paradigm, while effective for general language tasks, often fails to surpass simpler biologically-informed approaches, suggesting that current foundation models may not be effectively capturing the underlying biological principles.

Another critical finding is the **disconnect between quantitative metrics and biological utility**. Models achieving high Pearson correlations on gene expression prediction (e.g., >0.95 in raw expression space) may fail to capture biologically meaningful signals, as these metrics are heavily influenced by baseline expression magnitudes rather than perturbation-specific responses [93]. This highlights the necessity for biologically-grounded evaluation metrics that better reflect real-world research needs.

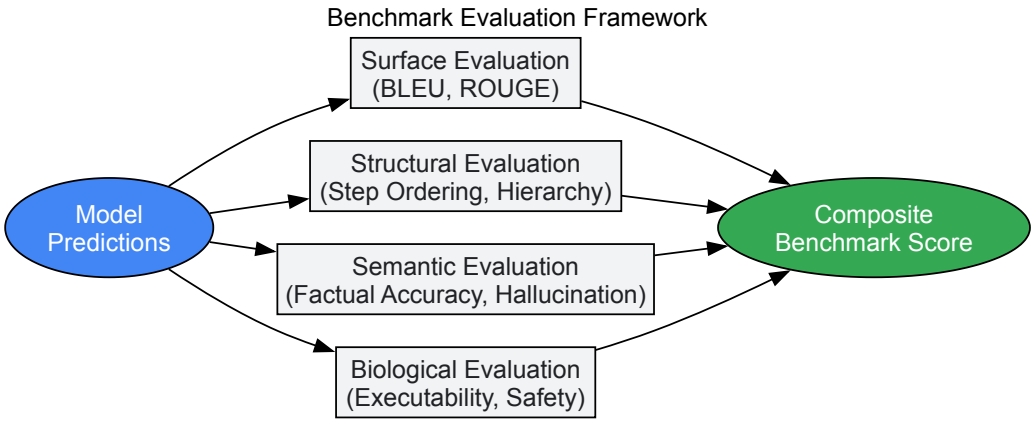

Table 4: Key Research Reagent Solutions for Benchmarking Biological Foundation Models

Resource Category	Specific Resources	Function in Benchmarking	Access Information
Protocol Repositories	Bio-protocol, Protocol Exchange, JOVE, Nature Protocols, Protocols.io	Source materials for protocol understanding and generation benchmarks	Publicly available with varying licensing
Single-cell Data Platforms	CZ CELLxGENE, Human Cell Atlas, PanglaoDB	Pretraining data for single-cell foundation models; evaluation benchmarks	Publicly accessible databases
Genomic Data Resources	ENCODE, GEO, SRA, Expression Atlas	Source data for genomic prediction tasks and long-range dependency benchmarks	Public repositories with standardized access
Benchmark Suites	BioProBench, DNALONGBENCH, BEND, LRB	Standardized evaluation frameworks for model comparison	Available via GitHub and academic repositories
Specialized Evaluation Tools	SCORE mechanism [94], Sketch-and-Fill paradigm	Structured evaluation of protocol quality and executability	Integrated into benchmark implementations

The experimental workflow for comprehensive benchmarking requires several critical computational "reagents": (1) **Structured data parsers** for processing biological protocols with nested step information; (2) **Tokenization schemes** adapted to biological sequences, such as gene ranking by

expression levels or binning by expression values; (3) **Domain-specific metrics** that go beyond textual similarity to assess biological plausibility; and (4) **Baseline implementations** including simple statistical baselines, traditional ML models, and expert models for comparison [10] [91] [92].

Emerging tools like the Structured COmponent-based REward (SCORE) mechanism provide specialized evaluation capabilities for assessing protocol quality across multiple dimensions: step granularity (controlling scale and avoiding redundancy), action ordering (ensuring logically consistent sequences), and semantic fidelity (verifying alignment between predicted and reference actions) [94]. Similarly, the "Sketch-and-Fill" paradigm separates analysis, structuring, and expression to ensure each protocol step is explicit and verifiable, addressing common issues of unordered steps and redundant operations in model-generated protocols [94].

Rigorous benchmarking frameworks are essential for advancing biological foundation models from theoretical constructs to practically useful research tools. Current benchmarks reveal significant limitations in model capabilities, particularly for tasks requiring deep reasoning, structured generation, and understanding of long-range dependencies. The emerging pattern from multiple benchmarking studies indicates that while foundation models show promise, they often underperform simpler approaches informed by biological domain knowledge.

Future benchmarking efforts should focus on: (1) Developing more biologically meaningful metrics that better reflect real-world research utility; (2) Creating benchmarks that assess model capabilities across multiple biological scales from molecular interactions to cellular systems; (3) Establishing standardized evaluation protocols that enable fair comparison across different model architectures; and (4) Addressing current limitations in benchmarking datasets, particularly the low perturbation-specific variance in some commonly used datasets [93]. As biological foundation models continue to evolve, rigorous benchmarking will be essential for guiding their development toward truly transformative applications in biological research and therapeutic development.

The "foundation model" (FM) paradigm—pretraining expansive models on massive, domain-specific datasets followed by fine-tuning on target tasks—has rapidly expanded beyond natural language processing and computer vision into specialized scientific domains, including bioinformatics [95] [9]. This approach promises a universal shift in artificial intelligence (AI) application, suggesting that large-scale pretraining is the key to unlocking superior performance on a wide array of downstream tasks. In bioinformatics, FMs are being aggressively developed and applied to genomics, transcriptomics, proteomics, and single-cell analysis, with claims of addressing long-standing challenges in computational biology [9] [10]. These models, often built on transformer architectures, are designed to learn fundamental biological principles from enormous corpora of unlabeled data, theoretically enabling them to generalize effectively to new problems with minimal task-specific data [10].

However, a critical examination of this emerging landscape reveals a surprising and consistent trend: simple, well-tuned supervised models frequently match or even exceed the performance of these complex, resource-intensive FMs [95] [96]. This phenomenon challenges the prevailing narrative of FM inevitability and dominance, suggesting that for many specialized biological problems, the benefits of large-scale pretraining have yet to be conclusively demonstrated. This technical review investigates this counter-narrative, synthesizing evidence from multiple domains within bioinformatics where traditional supervised learning not only remains competitive but occasionally renders FMs obsolete. We frame this analysis within a broader thesis on foundation models in bioinformatics review research, arguing that rigorous comparison against strong, well-tuned baselines is not merely an academic exercise but an essential practice for validating the true added value of any new FM. For researchers and drug development professionals, this insight is critical for allocating computational resources efficiently and avoiding unnecessary complexity in model deployment.

Evidence from the Bench: Case Studies in FM Underperformance

Genomics and the Nucleotide Transformer Benchmark

In genomics, several foundation models, such as the Nucleotide Transformer, have been developed using a "lift-and-shift" approach, where architectures like BERT are adapted for genomic sequences with specialized tokenizations and embeddings [95]. These models are pretrained on vast datasets of DNA sequences, sometimes encompassing entire genomes, with the objective of learning a generalizable representation of genomic function and regulation. The premise is that this deep pretraining will enable the model to excel at downstream tasks like predicting promoter regions, transcription factor binding sites, or variant effects.

Contrary to this expectation, empirical evaluations demonstrate that lightly adapted convolutional neural network (CNN) architectures, such as a wide ResNet or UNet, can attain state-of-the-art performance on the Nucleotide Transformer benchmark [95]. The key to this success lies not in architectural novelty or massive pretraining, but in targeted, automated model development. By leveraging tools like DASHA, a NAS-based pipeline, researchers can efficiently tune hyperparameters like kernel sizes and dilation rates in standard CNNs using only data from the target task. This supervised workflow, which forgoes pretraining entirely, consistently matches or surpasses the performance of FMs that have consumed orders of magnitude more data and computational resources. This result indicates that for many genomic recognition tasks, the inductive biases of CNNs—their innate ability to capture local spatial hierarchies—are sufficiently powerful, and that the purported general knowledge encoded in FMs may not provide a tangible performance benefit on these specific benchmarks [95].

Predicting Gene Perturbation Effects in Single-Cell Biology

The transcriptomics domain, particularly the prediction of single-cell transcriptional responses to genetic perturbations, presents another compelling case study. A recent benchmark published in *Nature Methods* revealed that simple linear models outperformed cutting-edge deep learning foundation models in forecasting these effects [96]. This surprising outcome can be attributed to several factors rooted in experimental biology. The benchmark datasets were primarily derived from genetically homogeneous cancer cell lines cultured under uniform, simplified laboratory conditions. This setup significantly reduces the biological complexity and variability typically encountered in heterogeneous tissues or *in vivo* environments.

Under these simplified conditions, the effects of most gene perturbations and their combinations were found to be largely independent or additive, with very few gene pairs eliciting true synergistic or buffering interactions within the transcriptome [96]. Consequently, a simple additive model was sufficiently complex to capture the underlying biological response pattern. The superior performance of simple linear baselines in this instance appears to be driven as much by the reductionist design of the biological experiments as by the limitations of the FM architectures. This suggests that for FM approaches to demonstrate clear superiority, they may need to be evaluated on more complex biological systems—such as heterogeneous co-cultures, primary tissues, or *in vivo* models—where non-linear and emergent interactions are more prevalent and demand greater model capacity [96].

Broader Trends Across Specialized Domains

This pattern of FM underperformance is not isolated to genomics and transcriptomics but appears to be a broader phenomenon across several specialized data modalities. A large-scale analysis across genomics, satellite imaging, and time series data—domains with at least five FMs each evaluated on nine or more standard tasks—found that it was consistently possible to train simple supervised models that matched or outperformed the latest FMs [95]. In time series forecasting, for instance, a well-tuned linear auto-regression (AR) model matched or outperformed every open-source time series FM on a standard suite of forecasting tasks, despite using four or more orders of magnitude fewer parameters and data [95].

A summary of key comparative results is presented in Table 1 below, quantifying the performance of simple baselines versus complex FMs across different domains.

Table 1: Performance Comparison of Simple Baselines vs. Foundation Models

Domain	Simple Baseline Model	Competing Foundation Model(s)	Performance Outcome	Key Factor
Genomics	Tuned CNN (e.g., Wide ResNet, UNet) [95]	Nucleotide Transformer & other genomics FMs [95]	Matched or outperformed on benchmark tasks [95]	Automated hyperparameter tuning (kernel size, dilation) [95]
Transcriptomics	Simple Linear/Additive Model [96]	Deep-learning foundation models for perturbation prediction [96]	Outperformed in predicting gene perturbation effects [96]	Simplified biological conditions (homogeneous cell lines, additive effects) [96]
Time Series	Tuned Linear Auto-regression (AR) [95]	Multiple open-source time series FMs [95]	Matched or outperformed on standard forecasting tasks [95]	Using >4 orders of magnitude fewer parameters & data [95]
Satellite Imaging	Lightly modified UNet [95]	SatMAE & other satellite FMs [95]	Matched downstream classification performance [95]	Strong, in-domain supervised model development [95]

These findings collectively demonstrate that fields like genomics, satellite imaging, and time series have not yet experienced their "BERT moment"—a reference to the point at which BERT-style models definitively supplanted previous supervised approaches in natural language processing [95]. The consistent success of simple baselines reinforces the critical importance of comparing new FMs against strong, well-tuned supervised models as a minimum standard for evaluation.

Underlying Causes: Why Simplicity Trumps Complexity

The Strong Baseline Hypothesis

A primary reason for the observed underperformance of FMs is the widespread failure to establish strong, well-tuned supervised baselines for comparison. Many FM studies fall into a "comparison echo chamber," benchmarking new models primarily against other existing FMs rather than against the best possible supervised model trained exclusively on target task data [95]. This creates a misleading impression of progress. A supervised workflow comprising thoughtful model development, rigorous hyperparameter tuning, and training on high-quality target data remains a formidable competitor. The recent success of tools like DASHA for architecture search and Auto-AR for time series underscores that a primary differentiator is often the care and sophistication applied to the tuning process, not the scale of pretraining [95]. When a simple AR model, a century-old technique, is rescued from obsolescence merely by considering longer lookback parameters and GPU-accelerated training, it highlights that many basic methods have not been fully explored or optimized before being abandoned for more complex alternatives [95].

Data Simplicity and the Lack of Emergent Complexity

In some biological contexts, the data themselves may not possess the degree of complexity required to justify an FM approach. As seen in the gene perturbation study, when cellular systems are reduced to their simplest form—homogeneous cell lines in controlled environments—the underlying biological relationships can become predominantly linear and additive [96]. In such scenarios, a model with high complexity (like an FM) is not only unnecessary but is also prone to overfitting or latching onto spurious correlations in the pretraining data. The true potential of FMs may only be realized when applied to problems characterized by high-dimensional, non-linear interactions and rich contextual dependencies, such as modeling whole-tissue dynamics, cross-modal regulation (e.g., gene-protein-metabolite), or patient-level outcomes. For many well-defined, narrow prediction tasks, the problem may simply be "not hard enough" to benefit from the vast, generalist knowledge encoded in an FM.

The "Lift-and-Shift" Architectural Mismatch

Many specialized FMs are constructed via a "lift-and-shift" strategy, where an architecture proven successful in vision or language (e.g., Transformer, Swin) is applied to a new domain with only modest modifications, such as a specialized tokenization [95]. This approach can overlook fundamental differences in the structure and semantics of the new data modality. For instance, genomic sequences, satellite images, and time series data have inherent properties—such as long-range dependencies, multi-spectral channels, and temporal auto-correlation—that may not be optimally handled by architectures designed for text or natural images. While the transformer's attention mechanism is flexible, a lightly modified CNN or a simple linear model might more directly and efficiently capture the essential inductive biases needed for a specific domain, leading to better performance with far less computational overhead. This suggests that innovative, domain-native architectures, rather than transplanted ones, might be a more fruitful path forward for achieving a genuine "BERT moment" in bioinformatics.

Experimental Protocols & Methodologies

The Supervised Baseline Workflow

To ensure a fair and rigorous comparison between a new FM and a traditional supervised model, the following experimental protocol is recommended. This workflow, which uses only data from the target task, consists of three key stages, as illustrated in the diagram below.

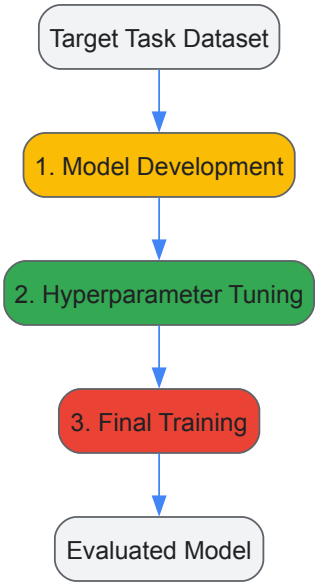

1. Model Development: Select a model architecture that is appropriate for the data modality and task. This need not be complex. For genomic sequences or image-like biological data (e.g., chromatin accessibility matrices), a standard CNN like ResNet or UNet is a strong starting point [95]. For tasks involving predicting continuous outcomes from vectorized features (e.g., gene expression levels), a multi-layer perceptron or even a linear model can be highly effective [96]. The goal is to choose an architecture with suitable inductive biases, not necessarily the most complex one.

2. Hyperparameter Tuning (The Critical Step): Systematically search the hyperparameter space to find the optimal configuration for the target task. This step is often the most significant differentiator between a weak and a strong baseline. Key hyperparameters to optimize include: - **For CNNs:** Kernel sizes, dilation rates, number of layers/filters, and activation functions [95]. - **For Linear Models/MLPs:** Regularization strength (L1/L2), learning rate, and optimizer selection. - **For AR models:** Lookback window length and differencing parameters [95]. Automated tools like DASHA (for architecture search) or Auto-AR (for time series) can be employed to make this process efficient and reproducible [95].

3. Final Training and Validation: Train the selected model with the optimized hyperparameters on the full training set from the target task. The final model should be evaluated on a held-out test set that was not used during development or tuning, ensuring an unbiased estimate of its performance on new data.

The Foundation Model Fine-tuning Workflow

For the FM, the standard pretrain-then-finetune paradigm should be followed.

1. Model Selection: Choose an FM that has been pretrained on a large, domain-relevant corpus (e.g., a model pretrained on a single-cell atlas or genomic sequences) [10].

2. Task Adaptation (Fine-tuning): Adapt the FM to the target task. This typically involves: - **Input Adaptation:** Modifying the input layer or tokenization strategy if necessary to accept the target task's data format. - **Head Replacement:** Replacing the model's final output layer (the "head") with a new one that matches the output dimension of the target task (e.g., number of classes for classification). - **Fine-tuning:** Training the entire model or a subset of its layers on the target task's training data. A lower learning rate is typically used to avoid catastrophic forgetting of the pretrained knowledge.

3. Evaluation: The fine-tuned FM is evaluated on the same held-out test set used for the supervised baseline.

Quantitative Comparison and Analysis

The performance of the two workflows should be compared using appropriate, domain-standard metrics (e.g., accuracy, AUROC, MSE, etc.). Crucially, this comparison should account for computational cost, data efficiency, and model complexity. Reporting the number of parameters, training time, and inference latency for both approaches provides a holistic view of their practical utility [95].

The Scientist's Toolkit: Essential Research Reagents & Materials

For researchers seeking to implement or validate the findings discussed, the following table details key computational "research reagents" and their functions.

Table 2: Key Research Reagents and Computational Tools

Tool / Material	Type	Primary Function	Relevance to FM vs. Baseline Research
DASHA [95]	Software Pipeline	Automated architecture search and hyperparameter tuning for CNNs.	Enables creation of highly competitive supervised baselines in domains like genomics and image analysis.
Auto-AR [95]	Software Workflow	GPU-optimized automated training of linear auto-regressive models.	Facilitates the creation of strong time series forecasting baselines that can challenge complex FMs.
scBERT / scGPT [10]	Foundation Model	Transformer-based models for single-cell RNA-seq data analysis.	Representative single-cell FMs used as benchmarks for performance comparison against simpler models.
Nucleotide Transformer [95]	Foundation Model	Transformer-based model pretrained on large-scale genomic sequences.	Representative genomics FM used as a benchmark for performance comparison against simpler models.
CZ CELLxGENE [10]	Data Resource	Unified platform providing access to millions of annotated single-cell datasets.	A primary source of pretraining data for single-cell FMs; also used for creating downstream task benchmarks.
PanglaoDB / Human Cell Atlas [10]	Data Resource	Curated compendia of single-cell data from multiple sources and studies.	Used to assemble the large and diverse training corpora required for pretraining robust scFMs.

The evidence compiled in this review unequivocally demonstrates that the ascendancy of foundation models in specialized domains like bioinformatics is not a foregone conclusion. Simple, well-tuned supervised models consistently present a formidable challenge to their complex, resource-heavy counterparts. This does not negate the potential of FMs but rather underscores a critical methodological imperative: the burden of proof lies with new FMs to demonstrate clear and practical advantages over strong baselines.

For the field to progress, future work must prioritize several key areas. First, the development and adoption of standardized, rigorous benchmarking protocols that mandate comparison against optimally tuned supervised models are essential to break the "comparison echo chamber" [95]. Second, FM research should gravitate towards problems where complexity is inherent and inescapable, such as modeling multi-cellular interactions, integrating multi-omic data, or predicting outcomes in genetically diverse populations, as these are the scenarios where the limitations of simple models are most likely to be exposed [96]. Finally, innovation should focus on creating truly domain-native architectures that move beyond the "lift-and-shift" approach, potentially leveraging mechanistic insights from biology to guide model design [95]. For now, practitioners in bioinformatics and drug development are advised to maintain a balanced perspective, investing in the development of robust supervised baselines as a necessary and often sufficient step before committing to the substantial costs associated with foundation models.

The advent of high-throughput single-cell sequencing has generated vast amounts of transcriptomic data, creating an unprecedented opportunity to decipher cellular heterogeneity at unprecedented resolution [10]. This data explosion has concurrently generated significant analytical challenges due to the inherent sparsity, noise, and batch effects present in single-cell RNA sequencing (scRNA-seq) data [97] [22]. Inspired by the success of large

language models (LLMs) in natural language processing, computational biologists have begun developing single-cell foundation models (scFMs) to address these challenges [10]. These models are pre-trained on millions of single-cell transcriptomes with the goal of learning universal biological patterns that can be transferred to various downstream tasks through fine-tuning or zero-shot application.

The paradigm of "pre-train then adapt" has revolutionized artificial intelligence applications in biology, promising to unlock deeper insights into cellular function and disease mechanisms [10] [9]. scFMs typically leverage transformer architectures to process gene expression data, treating individual cells as "sentences" and genes or their expression values as "words" or "tokens" [10]. This approach allows the models to capture complex gene-gene relationships and cellular states from large-scale datasets. However, the rapid emergence of multiple scFMs—including scGPT, Geneformer, scFoundation, CellFM, UCE, and others—has created a pressing need for comprehensive comparative analysis to guide researchers in selecting appropriate models for specific applications [22] [56].

This review provides an in-depth technical comparison of leading scFMs, examining their architectures, pre-training strategies, and performance across diverse biological tasks. We situate this analysis within the broader context of foundation models in bioinformatics, highlighting both the promises and limitations of current approaches. Through structured comparisons of technical specifications, performance benchmarks, and practical implementation considerations, we aim to provide researchers, scientists, and drug development professionals with a framework for effectively leveraging these powerful tools in their own work.

Core Architectural Principles and Pre-training Strategies

Foundational Concepts and Model Architectures

Single-cell foundation models share a common goal: to learn transferable representations of cellular states from large-scale transcriptomic data. Most scFMs are built on transformer architectures, which use attention mechanisms to model relationships between genes regardless of their positional relationships [10]. However, these models diverge significantly in how they process gene expression data, which lacks the inherent sequential ordering of natural language. A key challenge in applying transformers to single-cell data is addressing the non-sequential nature of gene expression, where gene order carries no biological meaning [10] [22].

To overcome this challenge, different models employ distinct tokenization strategies. One common approach is **gene ranking**, where genes are ordered by their expression levels within each cell, creating a deterministic sequence for transformer processing [10]. Geneformer exemplifies this approach, using a rank-based strategy where the top 2,048 ranked genes by expression serve as input tokens [22]. Alternatively, **value binning** strategies discretize continuous expression values into categorical "buckets," transforming regression problems into classification tasks [10] [97]. scBERT employs this approach, binning gene expression values to enable classification-style pre-training [97]. A third strategy, **value projection**, preserves the full resolution of continuous expression data by projecting raw values into embedding spaces [97]. scFoundation uses this approach to directly predict raw gene expression values [97] [22].

The transformer architectures themselves also vary significantly between models. Most scFMs use either encoder-only or decoder-only transformer variants [10]. Encoder-based models like Geneformer use bidirectional attention mechanisms that consider all genes simultaneously, making them well-suited for classification tasks and embedding generation [10] [22]. In contrast, decoder-based models like scGPT use unidirectional attention that iteratively predicts masked genes conditioned on known genes, potentially offering advantages for generative tasks [10]. Hybrid architectures that combine encoder and decoder components are also emerging [10].

Pre-training Data and Objectives

The scale and diversity of pre-training data significantly impact model performance and generalizability. Leading scFMs have been trained on datasets ranging from 10 million to over 100 million cells [97] [22]. CellFM currently represents the upper extreme, having been trained on 100 million human cells—approximately twice the dataset size of previous largest single-species models [97]. These training datasets are typically aggregated from public repositories such as CZ CELLxGENE, NCBI GEO, ENA, and various cell atlases [10] [97].

Pre-training objectives are predominantly self-supervised, with masked gene modeling (MGM) being the most common approach [10] [22]. In MGM, a subset of genes is masked in the input, and the model is trained to predict the expression values or identities of these masked genes based on the remaining context [10]. However, implementations differ: Geneformer uses a cross-entropy loss to predict gene identities based on rank [22], while scFoundation uses mean squared error (MSE) loss to predict raw expression values [22]. scGPT employs an iterative MGM approach with both gene-prompt and cell-prompt objectives [22], and UCE uses a modified MGM with binary cross-entropy loss to predict whether genes are expressed or not [22].

Table 1: Architectural Specifications of Leading Single-Cell Foundation Models

Model	Architecture Type	Parameters	Pre-training Data	Tokenization Strategy	Pre-training Objective
Geneformer	Encoder (6L/12L)	40M	30M cells	Gene ranking (top 2,048)	MGM with CE loss (gene ID prediction)

Model	Architecture Type	Parameters	Pre-training Data	Tokenization Strategy	Pre-training Objective
scGPT	Encoder with attention mask	50M	33M human cells	Value binning (1,200 HVGs)	Iterative MGM with MSE loss
scFoundation	Asymmetric encoder-decoder	100M	50M human cells	Value projection (19,264 genes)	Read-depth-aware MGM with MSE loss
CellFM	Modified RetNet (ERetNet)	800M	100M human cells	Value projection	MGM with linear projection
UCE	Encoder	650M	36M cells	Gene sampling by expression & genomic position	Binary CE loss for expression prediction
scBERT	Encoder	Not specified	Millions of human cells	Value binning	Masked gene prediction with CE loss

Comprehensive Performance Benchmarking

Evaluation Metrics and Methodologies

Rigorous evaluation of scFMs requires diverse metrics that capture performance across multiple dimensions. Recent benchmarking efforts have employed both standard evaluation metrics and novel biologically-informed approaches [22]. Unsupervised metrics such as Average BIO (AvgBIO) score and average silhouette width (ASW) assess cell type clustering quality without predefined labels [98]. Batch integration metrics evaluate a model's ability to remove technical artifacts while preserving biological variation, critical for integrating datasets from different sources [98] [22].

Supervised metrics include accuracy for cell type annotation and perturbation prediction tasks [22]. Additionally, novel knowledge-based metrics such as scGraph-OntoRWR measure the consistency of cell type relationships captured by scFMs with prior biological knowledge from cell ontologies [22]. The Lowest Common Ancestor Distance (LCAD) metric assesses the ontological proximity between misclassified cell types, providing a biologically-informed measure of error severity [22].

Evaluation methodologies typically compare scFMs against established baseline methods including Highly Variable Genes (HVG) selection, Harmony, and scVI [98] [22]. These comparisons are conducted in both zero-shot settings (where models are applied without any task-specific fine-tuning) and fine-tuning scenarios [98] [22]. The zero-shot evaluation is particularly important for discovery settings where labels are unknown, making fine-tuning impossible [98].

Model Performance Across Key Tasks

Cell Type Annotation and Clustering

Cell type annotation represents a fundamental application of scFMs. Benchmarking studies reveal significant performance variation across models and datasets [22]. In zero-shot settings, both Geneformer and scGPT have demonstrated limitations, underperforming compared to simpler methods like HVG selection, Harmony, and scVI across multiple datasets [98]. For example, when evaluated on the Pancreas benchmark dataset, Geneformer's cell embedding space primarily reflected batch effects rather than biologically meaningful cell type information [98].

Fine-tuning significantly improves performance for all models, but the extent of improvement varies [22]. Comprehensive benchmarking across five datasets with diverse biological conditions reveals that while scFMs are robust and versatile tools, simpler machine learning models can sometimes adapt more efficiently to specific datasets, particularly under resource constraints [22]. Notably, no single scFM consistently outperforms others across all tasks and datasets, emphasizing the need for task-specific model selection [22].

Batch Integration

Batch integration—correcting for technical variations between datasets while preserving biological signals—represents a critical challenge in single-cell analysis [98]. Evaluations on the Pancreas dataset, which includes data from five different sources, reveal qualitative differences in model performance [98]. Geneformer's embeddings show poor retention of cell type information, with clustering primarily driven by batch effects [98]. scGPT provides better cell type separation but still exhibits batch-effect-driven structure in dimensionality reduction [98].

Quantitatively, Geneformer underperforms relative to scGPT, Harmony, scVI, and HVG across most datasets in batch integration tasks [98]. scVI and Harmony generally outperform scGPT on datasets where batch effects are primarily technical, while scGPT shows advantages on more complex datasets where both technical and biological batch effects are present [98]. Surprisingly, HVG selection achieved the best batch integration scores across all datasets in full-dimensional evaluation [98].

Gene Function Prediction and Network Analysis

Gene-level tasks provide another important dimension for evaluating scFMs. Models exhibit varying capabilities in predicting gene functions and capturing gene-gene relationships [97] [22]. CellFM demonstrates strong performance in gene function prediction, potentially attributable to its extensive pre-training on 100 million human cells and large parameter count (800M) [97]. Similarly, Geneformer and scFoundation show robust performance on gene-level tasks, benefiting from their effective pre-training strategies [56].

Analysis of attention mechanisms in transformer-based scFMs suggests they can learn biologically meaningful gene-gene relationships [10] [22]. However, the practical utility of these learned relationships for predicting novel gene functions or regulatory networks requires further validation [22].

Table 2: Performance Comparison Across Downstream Tasks

Model	Zero-shot Cell Annotation	Fine-tuned Cell Annotation	Batch Integration	Gene Function Prediction	Perturbation Prediction
Geneformer	Underperforms baselines [98]	Strong with fine-tuning [22]	Poor (batch effects dominate) [98]	Strong [56]	Not specified
scGPT	Variable, dataset-dependent [98]	Robust across tasks [56]	Moderate to strong [98]	Moderate [22]	Strong [22]
scFoundation	Not specified	Not specified	Not specified	Strong [56]	Not specified
CellFM	Not specified	Not specified	Not specified	Strong [97]	Not specified
UCE	Not specified	Not specified	Not specified	Moderate [22]	Not specified
HVG Baseline	Outperforms Geneformer and scGPT [98]	N/A	Best overall scores [98]	N/A	N/A

Critical Analysis of Limitations and Challenges

Zero-Shot Performance Gaps

A critical evaluation of scFMs reveals significant limitations in their zero-shot capabilities [98] [99]. Despite being pre-trained on millions of cells, models like Geneformer and scGPT often underperform simpler baseline methods when applied without task-specific fine-tuning [98]. This performance gap suggests that these models may not be learning transferable biological concepts as effectively as initially hoped [98] [99].

The masked language model pre-training framework, while intuitively appealing, may not inherently produce high-quality cell embeddings for downstream tasks [98]. Analysis of scGPT's gene expression prediction capability reveals limitations: without conditioning on cell embeddings, the model predicts median expression values regardless of true expression levels [99]. Even with cell embedding conditioning, performance improvements are primarily limited to highly expressed "housekeeping" genes rather than context-specific variable genes [99].

These findings have important implications for biological discovery applications, where zero-shot capabilities are essential for identifying novel cell types or states without predefined labels [98]. The current generation of scFMs may have limited utility in truly exploratory settings, despite being marketed as general-purpose solutions [98] [99].

Data and Computational Challenges

Substantial challenges remain in data quality, model interpretability, and computational requirements [10]. Single-cell datasets exhibit significant variability in quality, depth, and technical noise, creating challenges for assembling balanced pre-training corpora [10]. Batch effects and platform-specific artifacts can persist in model embeddings, limiting their biological utility [98].

Interpretability represents another significant challenge. While attention mechanisms theoretically allow identification of important genes and relationships, extracting biologically meaningful insights from these patterns remains non-trivial [10]. The black-box nature of large transformer models complicates biological validation and hypothesis generation [22].

Computational intensity for training and fine-tuning presents practical barriers to widespread adoption [10]. Training models like CellFM with 800 million parameters requires specialized hardware and substantial resources [97]. While transfer learning through fine-tuning reduces the data requirements for specific applications, the computational costs remain substantial compared to traditional bioinformatics methods [22].

Implementation Frameworks and Best Practices

Standardized Frameworks for Model Application

The heterogeneous architectures and coding standards across scFMs have created significant implementation challenges for researchers [56]. To address this, standardized frameworks like BioLLM (Biological Large Language Model) provide unified interfaces for integrating and applying diverse scFMs [56]. These frameworks offer standardized APIs that eliminate architectural and coding inconsistencies, enabling streamlined model access and switching [56].

BioLLM supports both zero-shot and fine-tuning evaluation, facilitating consistent benchmarking across models and tasks [56]. The framework includes comprehensive documentation and built-in evaluation metrics, reducing the implementation burden for researchers [56]. Such standardized approaches are crucial for accelerating adoption and enabling fair comparison of different scFMs across diverse applications.

Model Selection Guidelines

Based on comprehensive benchmarking studies, model selection should be guided by specific task requirements, dataset characteristics, and available computational resources [22]. scGPT demonstrates robust performance across diverse tasks, particularly in zero-shot and fine-tuning scenarios, making it a strong general-purpose choice [56]. Geneformer and scFoundation excel in gene-level tasks, benefiting from their effective pre-training strategies [56].

For cell-type annotation and batch integration tasks, simpler methods like Harmony and scVI remain competitive, particularly in zero-shot settings [98]. In resource-constrained environments or for specific well-defined tasks, these established methods may provide more efficient solutions than large foundation models [22].

The Roughness Index (ROGI) can serve as a proxy for model selection in a dataset-dependent manner, helping researchers identify models that create smoother latent landscapes for specific data types [22]. Task-specific and overall model rankings generated through non-dominated sorting algorithms that aggregate multiple evaluation metrics provide additional guidance for model selection [22].

Table 3: The Scientist's Toolkit - Essential Research Reagents and Computational Resources

Resource Category	Specific Examples	Function/Purpose
Data Resources	CZ CELLxGENE [10], Human Cell Atlas [10], PanglaoDB [10]	Standardized single-cell datasets for model training and validation
Pre-processing Tools	SynEcoSys single-cell database [97], Seurat [22], Scanpy [97]	Quality control, gene name standardization, format unification
Evaluation Frameworks	BioLLM [56], scGraph-OntoRWR [22]	Standardized model evaluation and biological validation
Baseline Methods	HVG selection [98], Harmony [98], scVI [98]	Performance comparison and method validation
Computational Infrastructure	MindSpore AI Framework [97], Ascend910 NPUs [97]	Hardware and software for training and deploying large models

Visualizing Model Architectures and Benchmarking Workflows

scFM Architecture Comparison

Benchmarking Workflow for scFM Evaluation

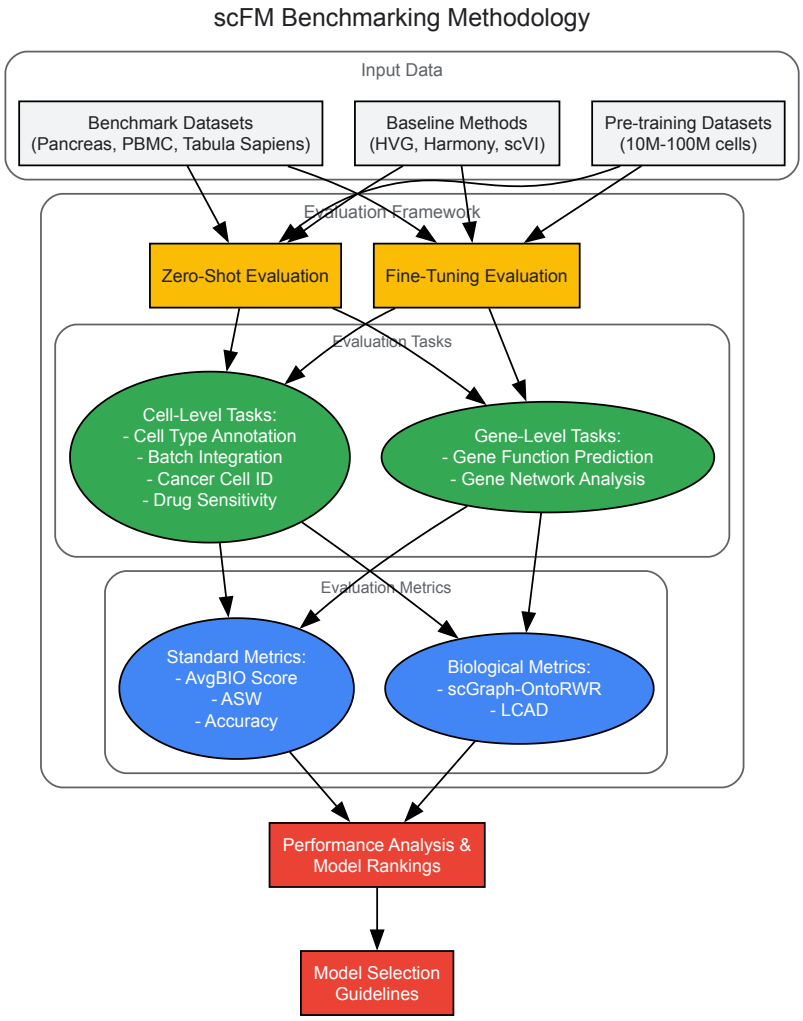

Single-cell foundation models represent a transformative approach to analyzing transcriptomic data, offering the potential to learn universal representations of cellular states that transfer across diverse biological contexts [10] [9]. Our comparative analysis reveals a rapidly evolving landscape with distinct architectural philosophies and performance characteristics. While promising, these models face significant challenges in zero-shot generalization, interpretability, and computational requirements [98] [22].

The ideal scFM remains elusive, with different models excelling in specific tasks and contexts [22]. scGPT demonstrates robust performance across diverse applications [56], while Geneformer and scFoundation show particular strengths in gene-level tasks [56]. CellFM's massive scale (800M parameters trained on 100M cells) represents the current frontier in model size [97], though the relationship between scale and performance appears complex [98] [22].

Future developments in scFMs will likely focus on improving zero-shot capabilities through better pre-training objectives and architectures [98] [99]. Multimodal integration—combining transcriptomic, epigenetic, proteomic, and spatial data—represents another important frontier [10] [22]. As standardized frameworks like BioLLM mature [56], they will accelerate model development and evaluation, enabling more systematic comparisons and biologically meaningful benchmarking.

For researchers and drug development professionals, selecting appropriate scFMs requires careful consideration of task requirements, dataset characteristics, and available resources [22]. While foundation models offer exciting capabilities, traditional methods remain competitive for specific applications, particularly in zero-shot settings [98] [22]. As the field matures, we anticipate more specialized models optimized for particular biological contexts and clinical applications, ultimately fulfilling the promise of these powerful tools to advance our understanding of cellular biology and disease mechanisms.

Evaluating Biological Relevance with Novel Metrics like scGraph-OntoRWR

Single-cell foundation models (scFMs) represent a transformative advancement in computational biology, leveraging large-scale, self-supervised learning on massive single-cell transcriptomics datasets to develop versatile tools for biological discovery. These models, typically built on transformer architectures, have demonstrated remarkable capabilities in tasks ranging from cell type annotation to drug sensitivity prediction [10]. However, as these models grow in complexity and scale, a critical challenge has emerged: how can we effectively evaluate whether their internal representations and outputs capture biologically meaningful patterns rather than merely optimizing abstract computational metrics? Traditional evaluation metrics often fail to assess the biological plausibility of model predictions, creating a significant gap between computational performance

and biological relevance [22]. This gap is particularly problematic for applications in drug development and clinical research, where biologically implausible predictions could lead to costly erroneous conclusions.

The fundamental challenge stems from the nature of single-cell RNA sequencing (scRNA-seq) data itself, which is characterized by high sparsity, high dimensionality, and low signal-to-noise ratio [22]. While scFMs are designed to overcome these challenges through pre-training on millions of cells, their ability to extract unique biological insights beyond standard methods has remained unclear. Without proper biological grounding, these powerful models risk becoming black boxes that provide mathematically correct but biologically irrelevant outputs. It is within this context that novel evaluation metrics like scGraph-OntoRWR have emerged as crucial tools for bridging the gap between computational performance and biological meaning, ensuring that foundation models truly advance our understanding of cellular biology rather than merely optimizing mathematical objectives [100].

The scGraph-OntoRWR Metric: Principles and Methodology

Conceptual Foundation and Biological Rationale

The scGraph-OntoRWR metric represents a paradigm shift in evaluating single-cell foundation models by directly measuring the consistency between computational outputs and established biological knowledge. Unlike traditional metrics that assess clustering quality or classification accuracy in isolation, scGraph-OntoRWR specifically evaluates how well the relational structure of cell types learned by an scFM aligns with the hierarchical relationships defined in biological ontologies [100]. This approach is grounded in the recognition that cells exist within a structured biological continuum, with relationships that have been carefully curated and validated by the scientific community over decades. By using cell ontologies as a ground truth reference, the metric provides a biologically-grounded benchmark for assessing whether the patterns discovered by complex machine learning models genuinely reflect biological reality rather than technical artifacts or statistical anomalies.

The theoretical underpinning of scGraph-OntoRWR rests on the concept that meaningful biological representations should preserve the ontological proximity between cell types. For instance, two different types of T-cells should be more similar to each other in the model's latent space than either is to a neuron, reflecting their biological relationships [22] [100]. This approach addresses a critical limitation of conventional evaluation methods, which may reward models for producing well-separated clusters even when those clusters contradict established biological knowledge. By formally measuring this alignment, scGraph-OntoRWR provides a quantitative measure of biological relevance that complements traditional performance metrics, offering researchers a more holistic view of model quality and biological utility.

Methodological Implementation

The implementation of scGraph-OntoRWR involves a sophisticated methodology that integrates graph theory, ontology analysis, and single-cell data processing. The metric operates by constructing two complementary graphs from different knowledge sources and comparing their structural properties. The first graph is derived from the embeddings generated by the single-cell foundation model, where cells are represented as nodes in a high-dimensional space, and their similarity relationships form the edges. The second graph is constructed from formal cell ontologies, which provide a biologically-validated hierarchy of cell type relationships [100].

The computational workflow begins with the extraction of cell embeddings from the target scFM, followed by the construction of a k-nearest neighbor graph based on cosine similarity or Euclidean distance in the latent space. Simultaneously, the relevant cell ontology is processed into a structured graph where nodes represent cell types and edges represent ontological relationships such as "is*a" and "part*of" [100]. The core of the scGraph-OntoRWR algorithm then applies a Random Walk with Restart (RWR) mechanism to both graphs, simulating the propagation of similarity through each network. By comparing the steady-state distributions of these random walks, the metric quantifies the alignment between the model-derived relationships and the ontology-defined biological relationships [22] [100]. This sophisticated approach captures both local and global structural similarities, providing a comprehensive assessment of biological consistency.

Table: Key Components of the scGraph-OntoRWR Methodology

Component	Description	Biological Significance
Model-Derived Cell Graph	K-nearest neighbor graph constructed from scFM embeddings	Captures similarity relationships as learned by the foundation model from data
Ontology Graph	Structured hierarchy of cell types from biological ontologies	Encodes expert-curated knowledge about cellular relationships
Random Walk with Restart (RWR)	Graph traversal algorithm that simulates similarity propagation	Measures multi-hop relationships beyond direct connections
Similarity Comparison	Comparison of steady-state distributions between graphs	Quantifies alignment between learned and established biological knowledge

Experimental Workflow and Integration

The complete experimental workflow for implementing scGraph-OntoRWR evaluation involves a carefully orchestrated sequence of steps that begins with data preparation and culminates in quantitative biological relevance scoring. First, high-quality single-cell datasets with reliable manual annotations must be selected or generated, ensuring that the evaluation has a solid foundation in biological truth [22]. These datasets should encompass diverse biological conditions, multiple tissues, and various technical platforms to provide a comprehensive assessment of model performance across different scenarios that mimic real-world research conditions.

Next, the target foundation model is used to generate latent representations of the cells in the evaluation dataset, producing the embeddings that will be analyzed. The scGraph-OntoRWR algorithm is then applied to quantify the biological consistency of these embeddings, producing a quantitative score that reflects the model's performance on this crucial dimension. Importantly, this evaluation is typically performed alongside traditional metrics and the novel LCAD (Lowest Common Ancestor Distance) metric, which measures the ontological proximity between misclassified cell types to assess the severity of annotation errors [22] [100]. This multi-faceted evaluation strategy provides researchers with a comprehensive view of model performance, balancing computational efficiency with biological plausibility.

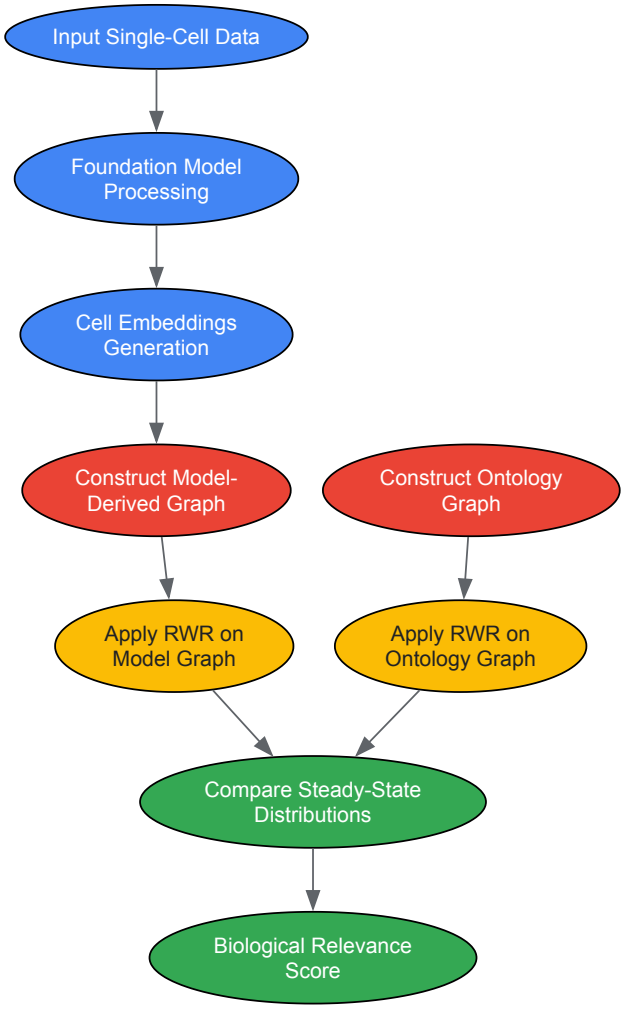

Diagram 1: The scGraph-OntoRWR evaluation workflow, illustrating the parallel processing of model-derived and ontology graphs and their comparative analysis.

Comprehensive Benchmark Insights: Performance Across Biological Tasks

Comparative Model Performance

The comprehensive benchmark study that introduced scGraph-OntoRWR evaluated six prominent single-cell foundation models against established baseline methods, revealing crucial insights about their biological relevance and practical utility. The models assessed included Geneformer, scGPT, UCE, scFoundation, LangCell, and scCello, representing diverse architectural approaches and pre-training strategies [22] [100]. When evaluated across multiple biological tasks using traditional metrics alongside scGraph-OntoRWR, the results demonstrated that no single scFM consistently outperformed others across all tasks, emphasizing that model selection must be tailored to specific applications and data characteristics [22]. This finding challenges the assumption that more complex or larger models are universally superior and highlights the importance of task-specific evaluation.

The benchmark encompassed two gene-level tasks (tissue specificity and Gene Ontology term prediction) and four cell-level tasks (batch integration, cell type annotation, cancer cell identification, and drug sensitivity prediction), providing a holistic view of model capabilities [100]. A key revelation was that while foundation models demonstrated remarkable robustness and versatility across diverse applications, simpler machine learning models sometimes adapted more efficiently to specific datasets, particularly under resource constraints [22] [101]. This nuanced understanding helps researchers make informed decisions about when the complexity of foundation models is justified by their performance benefits and when simpler approaches might be more appropriate.

Table: Performance Ranking of Single-Cell Foundation Models Across Biological Tasks

Model	Batch Integration	Cell Type Annotation	Cancer Cell Identification	Drug Sensitivity Prediction	Overall Biological Relevance Ranking
Geneformer	2	3	1	2	2
scGPT	3	2	3	3	3
UCE	1	4	4	4	4
scFoundation	4	1	2	1	1
LangCell	5	5	5	5	5
scCello	6	6	6	6	6
Traditional ML	7	7	7	7	7
HVG Selection	8	8	8	8	8

Biological Significance of Findings

The application of scGraph-OntoRWR in the benchmark study yielded biologically significant insights that extended beyond conventional performance metrics. The evaluation demonstrated that pretrained zero-shot scFM embeddings indeed capture meaningful biological insights into the relational structure of genes and cells, which proved beneficial for downstream tasks [22] [100]. This finding validates the fundamental premise of foundation models—that large-scale pre-training enables the learning of generalizable biological principles that transfer well to new datasets and tasks. Additionally, researchers discovered that performance improvements correlated with what they termed "a smoother landscape" in the pretrained latent space, reducing the difficulty of training task-specific models [22]. This smoother representation space appears to better reflect the continuous nature of biological processes and cell states.

Perhaps most significantly, the benchmark revealed that the biological relevance captured by scGraph-OntoRWR did not always correlate directly with traditional performance metrics, highlighting the unique value of this novel evaluation approach [100]. In some cases, models with similar traditional performance scores showed markedly different biological consistency as measured by scGraph-OntoRWR, suggesting that this metric captures distinct aspects of model quality. Furthermore, the metric proved particularly valuable for identifying cases where models achieved high accuracy by exploiting technical artifacts or dataset-specific biases rather than learning biologically generalizable patterns [22]. This capability makes scGraph-OntoRWR an essential tool for ensuring that models will generalize well to new data and real-world applications in drug development and clinical research.

Implementation Framework: The Scientist's Toolkit

Successfully implementing biological relevance evaluation with metrics like scGraph-OntoRWR requires access to specialized computational resources, biological databases, and software tools. These essential "research reagents" form the foundation for rigorous assessment of single-cell foundation models and ensure that evaluations are biologically meaningful and computationally reproducible. The toolkit encompasses everything from gene embeddings and ontological references to specialized software libraries and benchmark datasets, each playing a critical role in the comprehensive evaluation ecosystem [100].

Based on the benchmark study and methodological framework, the following table details the key components required for implementing scGraph-OntoRWR and related biological relevance assessments. These resources have been validated through comprehensive testing and represent the current state-of-the-art in biological evaluation of single-cell foundation models. Availability of these reagents varies, with some being freely accessible to the research community while others may require specific computational infrastructure or licensing arrangements.

Table: Essential Research Reagent Solutions for Biological Relevance Evaluation

Reagent/Resource	Function	Biological Significance	Accessibility
Gene Embeddings	Numerical representations of genes in latent space	Capture functional similarities between genes based on co-expression patterns across diverse cellular contexts	Model-dependent
Cell Ontologies	Structured vocabularies defining cell types and relationships	Provide ground truth for evaluating biological relevance of model outputs	Publicly available
Attention Mechanisms	Model components that identify important relationships between inputs	Reveal gene-gene interactions and regulatory relationships learned from data	Model-dependent
Benchmark Datasets	Curated single-cell data with high-quality annotations	Enable standardized evaluation and comparison of different modeling approaches	Publicly available
GO Term Annotations	Gene Ontology functional classifications	Serve as biological prior knowledge for validating gene embeddings	Publicly available
scGraph-OntoRWR Algorithm	Specialized implementation for biological consistency measurement	Quantifies alignment between learned representations and biological knowledge	Custom implementation
LCAD Metric Tools	Computational implementation for ontological error assessment	Measures severity of cell type misclassifications based on biological hierarchy	Custom implementation

Practical Implementation Protocol

The practical implementation of biological relevance evaluation requires a systematic approach that integrates the various research reagents into a coherent workflow. The following step-by-step protocol outlines the key procedures for conducting a comprehensive assessment using scGraph-OntoRWR and complementary metrics:

- **Data Preparation and Curation**: Begin by selecting or generating high-quality single-cell datasets with reliable manual annotations. These datasets should encompass diverse biological conditions and multiple technical platforms to ensure comprehensive evaluation. Appropriate preprocessing including normalization, quality control, and batch effect management is essential at this stage [22].
- **Foundation Model Processing**: Generate latent representations of the evaluation data using the target single-cell foundation model. This typically involves feeding the processed single-cell data through the model and extracting the resulting cell embeddings from the appropriate layer. Different embedding layers may capture different types of biological information, so layer selection should be considered carefully [10].
- **Ontological Graph Construction**: Access relevant cell ontologies from authoritative sources such as the OBO Foundry and construct the ontological graph structure. This process involves parsing the ontology file, extracting the hierarchical relationships between cell types, and representing them as a graph with appropriate edge weights reflecting biological proximity [100].
- **Model-Derived Graph Construction**: Transform the model-generated cell embeddings into a k-nearest neighbor graph using appropriate similarity metrics. The choice of k and the similarity threshold can impact results, so parameter sensitivity analysis may be necessary [22].
- **scGraph-OntoRWR Execution**: Implement the Random Walk with Restart algorithm on both graphs and compute the similarity between their steady-state distributions. This core computational step requires efficient graph processing capabilities, especially for large datasets [22] [100].
- **Multi-Metric Integration and Interpretation**: Combine the scGraph-OntoRWR results with complementary metrics including LCAD for error analysis and traditional performance metrics. The integrated interpretation of these different perspectives provides a comprehensive view of model biological relevance and practical utility [22].

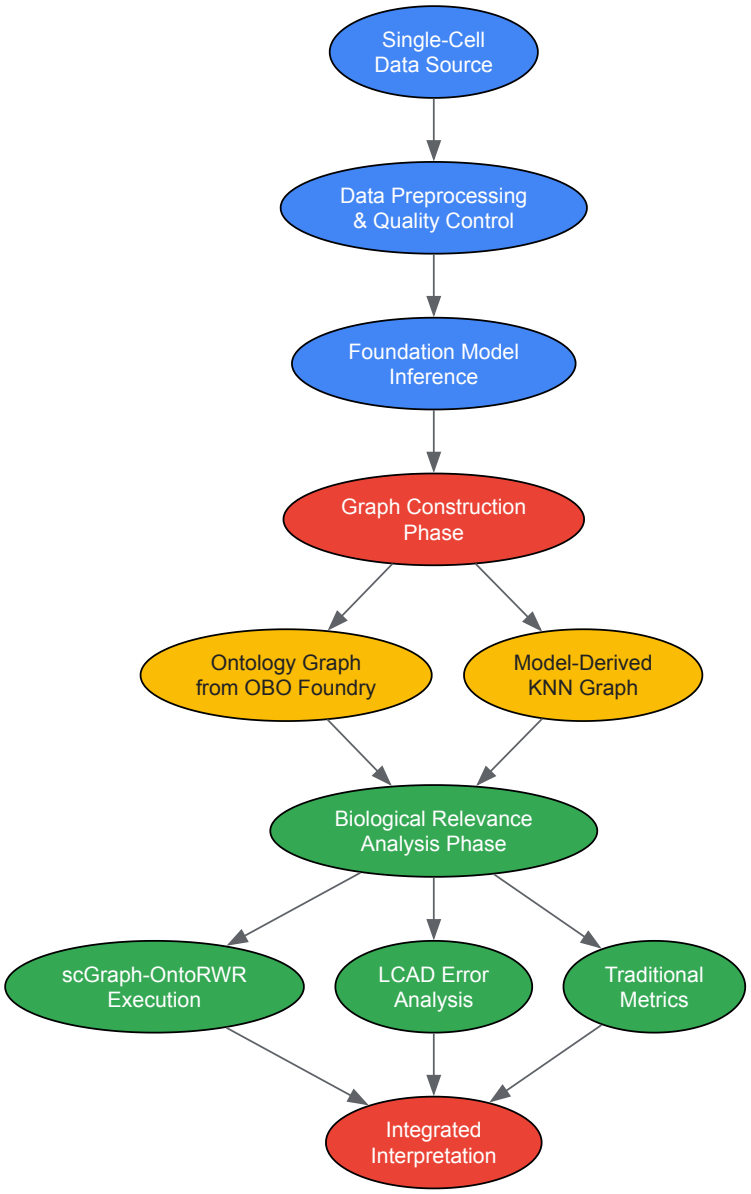

Diagram 2: End-to-end implementation protocol for biological relevance evaluation, showing the integration of multiple data sources and analytical phases.

Implications for Research and Drug Development

The development and validation of biological relevance metrics like scGraph-OntoRWR have profound implications for both basic research and applied drug development. In basic research, these metrics enable more rigorous validation of computational models, ensuring that discovered patterns reflect genuine biological mechanisms rather than technical artifacts or statistical anomalies [22]. This capability is particularly valuable for exploring novel biological systems where prior knowledge may be limited, as the ontological grounding provides a framework for interpreting results in the context of established biological principles. Furthermore, by quantifying biological consistency, these metrics facilitate more meaningful comparisons between different computational approaches, accelerating methodological advancements in the field.

In pharmaceutical research and drug development, biological relevance metrics offer the potential to significantly increase the translational validity of computational predictions. By ensuring that model outputs align with biological reality, these metrics reduce the risk of pursuing drug targets or therapeutic strategies based on computationally correct but biologically irrelevant patterns [22] [10]. This capability is especially valuable in single-cell studies of disease mechanisms, tumor microenvironments, and drug response, where understanding the biological validity of computational predictions can prioritize the most promising candidates for expensive experimental validation [22]. As foundation models increasingly influence target discovery and patient stratification strategies, rigorous biological validation through metrics like scGraph-OntoRWR will become essential components of the drug development pipeline, potentially reducing late-stage failures by improving the biological plausibility of early-stage computational discoveries.

The adoption of foundation models (FMs) in bioinformatics represents a paradigm shift in artificial intelligence, addressing longstanding challenges such as limited annotated data and data noise. These models, pretrained on vast biological datasets, demonstrate remarkable efficacy across various

downstream validation tasks, effectively representing diverse biological entities and heralding a new era in computational biology. This technical guide provides a comprehensive framework for selecting appropriate FMs for specific biological problems, with a focus on practical implementation, experimental protocols, and performance optimization. We synthesize current research to deliver actionable methodologies for researchers and developers working across sequence analysis, structure prediction, function annotation, and single-cell multi-omics integration.

Foundation models (FMs) are large-scale deep learning models pretrained on extensive datasets that can be adapted to a wide range of downstream tasks through fine-tuning mechanisms [3]. In bioinformatics, these models have demonstrated unprecedented capabilities in managing large-scale, unlabeled biological datasets, which is particularly valuable given that experimental procedures in biology are often costly and labor-intensive [9]. The versatility of FMs allows researchers to employ pretrained embeddings acquired from others to solve targeted biological problems with limited data through transfer learning approaches [3].

A key strength of FMs lies in their capacity to learn accurate representations of intricate biological datasets through data-intensive pretraining [3]. This flexibility has proven especially beneficial in bioinformatics, where FMs have successfully addressed core biological challenges including sequence analysis, structure construction, and function prediction [3]. The robust and reliable nature of these models, combined with their strong exploration and exploitation capacities and adaptable architecture for diverse downstream tasks, make them compelling tools for advancing biological research and drug development.

Types of Foundation Models and Their Architectures

Model Taxonomy and Characteristics

Foundation models in bioinformatics can be categorized through multiple dimensions, including their architectural design, pretraining strategies, and biological applications. The table below summarizes the primary FM types and their key characteristics:

Table 1: Foundation Model Taxonomy in Bioinformatics

Model Type	Core Architecture	Pretraining Approach	Primary Biological Applications	Key Examples
Language FMs	Transformer, BERT, GPT	Masked Language Modeling (MLM)	Genomic sequence analysis, protein function prediction	DNABERT, scBERT, BioBERT
Vision FMs	CNN, Vision Transformer	Self-supervised learning	Medical imaging, structural biology	-
Graph FMs	Graph Neural Networks	Graph reconstruction	Protein-protein interactions, biological networks	-
Multimodal FMs	Hybrid architectures	Cross-modal alignment	Multi-omics integration, spatial transcriptomics	scGPT

Architectural Foundations

Most foundation models in bioinformatics are built on transformer architectures, which are neural networks characterized by attention mechanisms that allow the model to learn and weight relationships between any pair of input tokens [10]. In biological applications, this enables models to determine which genes in a cell or which residues in a protein sequence are most informative for specific predictions. The two primary architectural variants include:

- **Encoder-based models** (e.g., BERT-like): Utilize bidirectional attention mechanisms where the model learns from the context of all elements simultaneously. These are particularly effective for classification tasks and embedding generation [3] [10].
- **Decoder-based models** (e.g., GPT-like): Employ unidirectional masked self-attention mechanisms that iteratively predict masked elements conditioned on known elements. These excel in generation tasks and sequential prediction [10].

Beyond transformers, residual convolutional neural networks (CNNs) have shown strong performance in specific biological applications. These architectures typically consist of multiple convolutional layers with expanding dilation factors and normalization techniques, proving particularly effective for genomic sequence analysis [102].

Task-Specific Model Selection Framework

Quantitative Performance Comparison Across Biological Tasks

Experimental evaluations across diverse biological tasks reveal significant performance variations between model architectures and training strategies. The following table synthesizes performance metrics from recent benchmark studies:

Table 2: Model Performance Comparison Across Biological Tasks

Biological Task	Model Architecture	Training Strategy	Performance Metric	Score	Reference
Gene Finding	ResNet-CNN	Self-Pretraining + CRF	MCC	0.64	[102]
Gene Finding	ResNet-CNN	Self-Pretraining	MCC	0.50	[102]
Gene Finding	ResNet-CNN	From Scratch	MCC	0.38	[102]
CpG Methylation	ResNet-CNN	Self-Pretraining	AUROC	~0.75	[102]
CpG Methylation	ResNet-CNN	From Scratch	AUROC	~0.70	[102]
Chromatin Accessibility	ResNet-CNN	Self-Pretraining	AUROC	~0.88	[102]
Histone Modification	ResNet-CNN	Self-Pretraining	AUROC	~0.85	[102]
Single-cell Analysis	scBERT	Pretraining + Fine-tuning	Accuracy	High	[10]

Model Selection Guidelines by Biological Data Type

Different biological data types and problem domains require specialized model architectures and training approaches:

Genomic Sequence Analysis

For DNA sequence analysis tasks including gene finding, regulatory element prediction, and variant effect prediction, transformer-based models pretrained on large genomic corpora have demonstrated superior performance. Specific recommendations include:

- **DNA Language Models (DNALMs)** such as DNABERT and Nucleotide Transformer are particularly effective for sequence-based tasks [102].
- **Self-pretraining** on task-specific unlabeled data followed by supervised fine-tuning can match or exceed the performance of models trained from scratch under identical compute constraints [102].
- **Structured prediction layers** such as Conditional Random Fields (CRFs) can substantially improve performance for tasks requiring global label consistency, such as gene boundary detection [102].

Single-Cell Multi-Omics Analysis

For single-cell transcriptomics, epigenomics, and multi-omics integration, specialized single-cell foundation models (scFMs) have emerged as powerful tools:

- **Encoder-based models** (e.g., scBERT) excel at cell type annotation and classification tasks [10].
- **Decoder-based models** (e.g., scGPT) are more suitable for data generation and imputation tasks [10].
- **Multi-modal architectures** that incorporate scRNA-seq, scATAC-seq, and spatial transcriptomics data provide the most comprehensive analysis of cellular heterogeneity [10].

Protein Structure and Function Prediction

For protein-related tasks including structure prediction, function annotation, and interaction mapping:

- **Protein language models** trained on evolutionary sequence data capture structural and functional constraints [3].
- **Graph-based models** effectively represent protein structures as molecular graphs for interaction prediction [3].
- **Multi-modal approaches** that combine sequence, structure, and textual information show promise for comprehensive functional annotation [3].

Experimental Protocols and Methodologies

Self-Pretraining Protocol for Genomic Tasks

Recent research demonstrates that self-pretraining on task-specific genomic data can yield stronger supervised models under limited compute resources [102]. The following protocol outlines the key methodological steps:

- **Data Preparation**:
  - Collect unlabeled sequences specific to the target biological problem
  - For gene finding tasks, use sequences of length 1,433–14,000 bp to enable modeling of long-range dependencies [102]
  - Implement one-hot encoding over a vocabulary of size V=7 (A, C, G, T, N, [MASK], [PAD])

- **Model Architecture Configuration**:
  - Implement a residual CNN encoder with 30 convolutional layers (kernel size 9) with 512 hidden channels [102]
  - Apply dilation doubling with each layer (reset every 6 layers, maximum dilation 32)
  - Use GELU activation and LayerNorm for stabilization
- **Self-Supervised Pretraining**:
  - Attach a masked language modeling (MLM) head to the encoder
  - Train on unlabeled DNA sequences with tokens masked independently with probability p_mlm = 0.15
  - Implement standard 80/10/10 replacement strategy for masked tokens
  - Compute MLM loss as: $\mathcal{L}_{MLM} = -\sum_{i:m_i=1} \log P_\theta(x_i \mid \tilde{x})$, where m_i indicates masked positions and $\tilde{x}$ is the corrupted input [102]
- **Supervised Fine-tuning**:
  - Replace MLM head with task-specific predictor (two-layer CNN with stride 1, ReLU activation, and linear output layer)
  - Fine-tune model end-to-end on downstream tasks
  - Use cross-entropy loss for single-label classification and binary cross-entropy for multi-label tasks
  - Train for 10 epochs for gene finding and 5 epochs for other tasks with larger training sets [102]

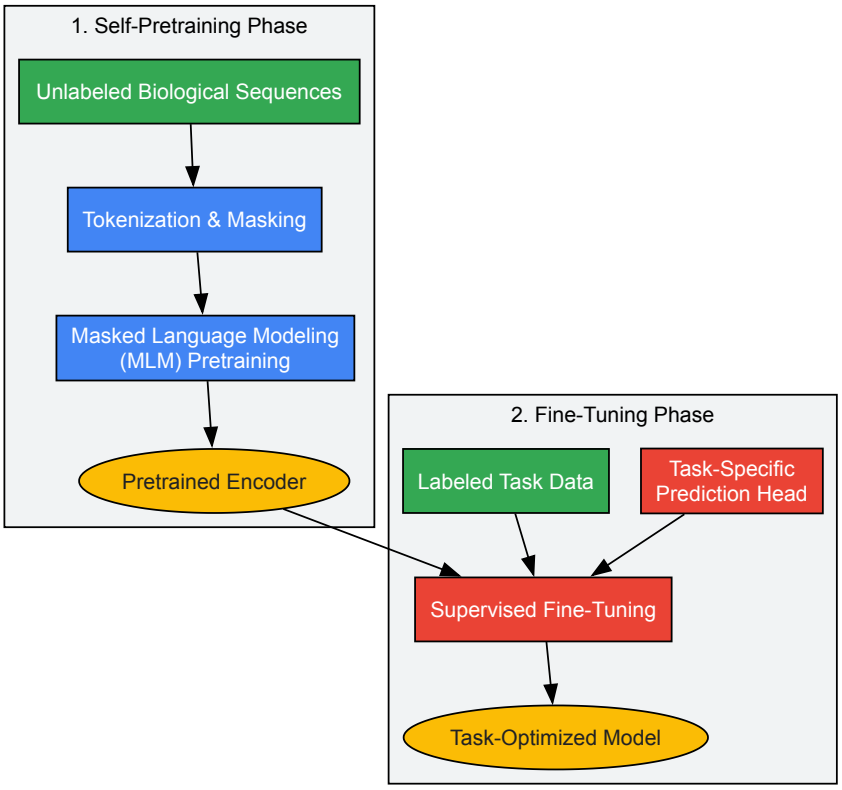

Self-Pretraining and Fine-Tuning Workflow

Single-Cell Foundation Model Implementation

For single-cell analysis tasks, the following protocol outlines the development and application of single-cell foundation models:

- **Data Sourcing and Curation**:
  - Compile large-scale single-cell datasets from public repositories (CZ CELLxGENE, Human Cell Atlas, GEO, SRA)
  - Apply rigorous quality control, filtering of low-quality cells and genes, and batch effect mitigation
  - Standardize data across studies to create a unified training corpus [10]
- **Tokenization Strategy**:
  - Define genes or genomic features as fundamental tokens
  - Implement deterministic gene ordering based on expression levels within each cell
  - Create token embeddings that combine gene identifiers with expression values

    - Incorporate special tokens for cell identity, modality, and batch information [10]

- **Model Training**:
    - Implement transformer architecture (encoder-based, decoder-based, or hybrid)
    - Train using self-supervised objectives (masked gene prediction, contrastive learning)
    - Utilize attention mechanisms to capture gene-gene interactions and regulatory relationships
    - Generate latent embeddings at both cell and gene levels for downstream tasks [10]

The Scientist's Toolkit: Research Reagent Solutions

The following table details essential computational tools and resources for implementing foundation models in biological research:

Table 3: Essential Research Reagents for Bioinformatics Foundation Models

Resource Category	Specific Tool/Platform	Function/Purpose	Application Context
Pretrained Models	DNABERT, Nucleotide Transformer	Provides pre-built DNA language models for genomic sequence analysis	Transfer learning for DNA-based tasks
Benchmark Datasets	BEND Benchmark	Standardized evaluation framework for DNA language models	Model performance comparison
Single-Cell Data Platforms	CZ CELLxGENE, PanglaoDB	Curated single-cell datasets for model training	Single-cell foundation model development
Model Architectures	Transformer, ResNet-CNN	Core neural network architectures	Custom model implementation
Specialized Layers	Conditional Random Fields (CRF)	Captures label dependencies in sequence labeling	Gene finding and annotation
Bioinformatics Libraries	TorchCRF	Implements CRF layers in PyTorch	Structured prediction tasks
Evaluation Metrics	MCC, AUROC	Performance assessment for biological tasks	Model validation and selection

Implementation Workflow and Decision Framework

The following diagram illustrates the comprehensive decision framework for selecting and implementing foundation models based on biological task requirements:

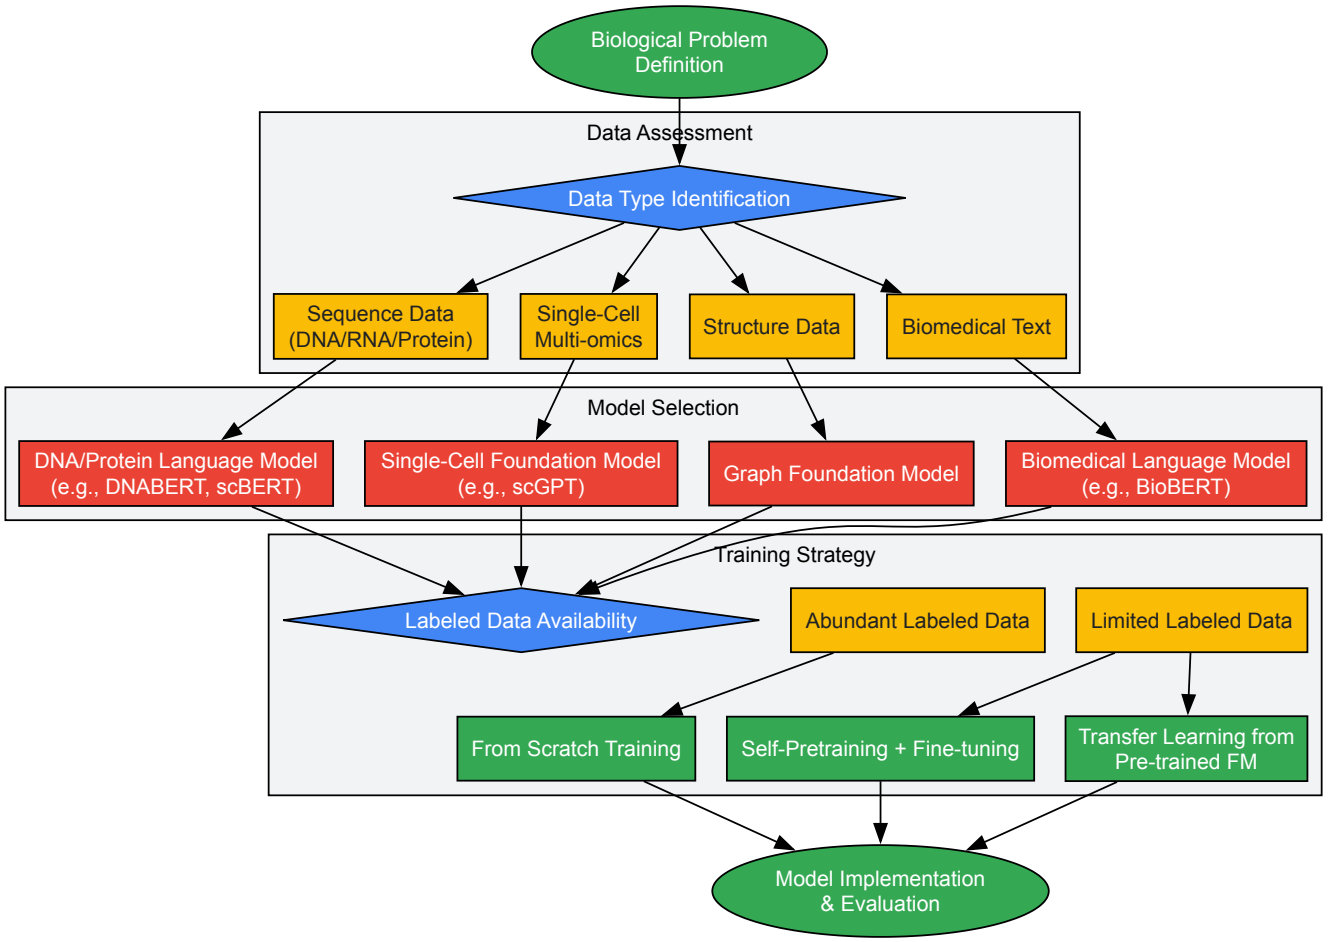

Foundation Model Selection Decision Framework

Task-specific model selection in bioinformatics requires careful consideration of biological data types, available computational resources, and performance requirements. The guidelines presented in this technical review demonstrate that self-pretraining on task-specific data provides a compute-efficient strategy for building strong supervised baselines, often matching or exceeding the performance of models trained from scratch under identical compute constraints [102]. For genomic sequence analysis, DNA language models with structured prediction layers deliver superior performance, while single-cell multi-omics tasks benefit from transformer-based scFMs with appropriate tokenization strategies [10].

As the field of biological foundation models continues to evolve, researchers and developers should prioritize interpretability, biological relevance, and computational efficiency when selecting and implementing these powerful tools. The experimental protocols and decision frameworks outlined in this guide provide a robust foundation for leveraging foundation models to advance biological discovery and therapeutic development.

Conclusion

Foundation Models have undeniably ushered in a new era for computational biology, demonstrating remarkable versatility and power in deciphering complex biological systems. The key takeaway is that while FMs provide a transformative, general-purpose framework for tasks ranging from sequence analysis to single-cell genomics, they are not a universal panacea. Rigorous benchmarking reveals that their performance is highly context-dependent, sometimes being surpassed by simpler, biologically-informed models. The future of FMs in bioinformatics hinges on overcoming critical challenges in data quality, model interpretability, and computational cost. Future efforts must focus on developing more robust, explainable, and efficient models that can seamlessly integrate multimodal data. Success in this endeavor will not only deepen our fundamental understanding of molecular landscapes but also firmly establish the theoretical and practical groundwork for groundbreaking advances in personalized medicine, therapeutic discovery, and clinical applications.

References

- 1. Foundation model [https://en.wikipedia.org/wiki/Foundation_model]
- 2. What are Foundation Models? - Generative AI [https://aws.amazon.com/what-is/foundation-models/]
- 3. Progress and opportunities of foundation models in ... [https://pmc.ncbi.nlm.nih.gov/articles/PMC11512649/]
- 4. Self-Supervised Learning and Foundation models [https://medium.com/@anuwaz1998/self-supervised-learning-and-foundation-models-49e8e5de9053]

- 5. Toward a Theoretical Understanding of Self-Supervised ... [https://ai.utexas.edu/events/2025-12-12/toward-theoretical-understanding-self-supervised-learning-foundation-model-era]
- 6. Foundation Models for Biomedical Research [https://watershed.bio/resources/foundation-models-for-biomedical-research]
- 7. A foundation model for clinician-centered drug repurposing [https://www.nature.com/articles/s41591-024-03233-x]
- 8. Artificial intelligence in bioinformatics: a survey - PMC [https://pmc.ncbi.nlm.nih.gov/articles/PMC12596291/]
- 9. Foundation models in bioinformatics [https://pubmed.ncbi.nlm.nih.gov/40078374/]
- 10. Single-cell foundation models: bringing artificial ... [https://www.nature.com/articles/s12276-025-01547-5]
- 11. Assessing the performance of generative artificial intelligence ... [https://pmc.ncbi.nlm.nih.gov/articles/PMC11833239/]
- 12. Biology-driven insights into the power of single-cell foundation ... [https://genomebiology.biomedcentral.com/articles/10.1186/s13059-025-03781-6]
- 13. Equipping Artificial Intelligence with the Lense of Evolution [https://news.rub.de/english/press-releases/2025-09-10-bioinformatics-equipping-artificial-intelligence-lense-evolution]
- 14. Demystifying the black box: A survey on explainable ... [https://www.sciencedirect.com/science/article/pii/S2001037024004495]
- 15. Review: Foundation models in bioinformatics [https://www.eurekalert.org/news-releases/1074506]
- 16. Multimodal Large Language Models in Medical Imaging [https://pmc.ncbi.nlm.nih.gov/articles/PMC12479233/]
- 17. MLLM4TS: Leveraging Vision and Multimodal Language ... [https://arxiv.org/html/2510.07513v1]
- 18. Foundation models for Biology: a short primer | by dalloliogm [https://medium.com/@dalloliogm/foundation-models-for-biology-a-short-primer-84d74b12b7a7]
- 19. Recommendations for bioinformatics in clinical practice - PMC [https://pmc.ncbi.nlm.nih.gov/articles/PMC12535132/]
- 20. Advancing bioinformatics capacity through Nextflow and nf- ... [https://www.frontiersin.org/journals/bioinformatics/articles/10.3389/fbinf.2025.1610015/full]
- 21. Foundation models and intelligent decision-making [https://www.sciencedirect.com/science/article/pii/S2666675825001511]
- 22. Biology-driven insights into the power of single-cell ... [https://pmc.ncbi.nlm.nih.gov/articles/PMC12492631/]
- 23. scMaui: a widely applicable deep learning framework for ... [https://bmcbioinformatics.biomedcentral.com/articles/10.1186/s12859-024-05880-w]
- 24. scMODAL: a general deep learning framework for ... [https://www.nature.com/articles/s41467-025-60333-z]
- 25. Multimodal deep learning approaches for single-cell multi ... [https://pmc.ncbi.nlm.nih.gov/articles/PMC10516349/]
- 26. Gene-LLMs: a comprehensive survey of transformer-based ... [https://www.frontiersin.org/journals/genetics/articles/10.3389/fgene.2025.1634882/full]
- 27. Large language models (LLMs) might be the future ... [https://pmc.ncbi.nlm.nih.gov/articles/PMC12430758/]
- 28. Improves Prediction of Genomic Pretraining Across Species Datasets [https://pubmed.ncbi.nlm.nih.gov/40894638/]
- 29. A novel sequence-based transformer model architecture ... [https://www.nature.com/articles/s41746-025-01942-2]
- 30. GitHub - stan-hua/CytoImageNet: A large-scale pretraining for... dataset [https://github.com/stan-hua/CytoImageNet]
- 31. Circuits Updates - July 2025 [https://transformer-circuits.pub/2025/july-update/index.html]
- 32. Generalized biological foundation model with unified ... [https://www.nature.com/articles/s42256-025-01044-4]
- 33. Sequence analysis tools in UniProt: BLAST and Align [https://www.ebi.ac.uk/training/events/sequence-analysis-tools-uniprot-blast-and-align-2025]
- 34. Geneious | Bioinformatics Software for Sequence Data Analysis [https://www.geneious.com/]
- 35. Top NGS Data Analysis Tools of 2025 [https://www.dromicsedu.com/blog-details/revolutionizing-genomics-top-ngs-data-analysis-tools-of-2025]
- 36. Best LIMS for Precision Medicine in 2025 [https://www.scispot.com/blog/best-lims-for-precision-medicine]
- 37. How AI Is Revolutionizing Drug Discovery in 2025 [https://medium.com/@healthfinaiinsights/how-ai-is-revolutionizing-drug-discovery-in-2025-with-alphafold-3-feb7c497e65b]
- 38. In-silico 3D molecular editing through physics-informed ... [https://www.nature.com/articles/s41467-025-61323-x]
- 39. Revolutionary 3D Drug Discovery Using AI-Powered ... [https://www.pnnl.gov/available-technologies/revolutionary-3d-drug-discovery-using-ai-powered-molecular-design]
- 40. Artificial intelligence-driven approaches for the rational ... [https://www.nature.com/articles/s44431-025-00005-6]
- 41. Progress and opportunities of foundation models in ... [https://pubmed.ncbi.nlm.nih.gov/39461902/]
- 42. Evolution of Protein Functional Annotation: Text Mining Study [https://pmc.ncbi.nlm.nih.gov/articles/PMC8952229/]
- 43. Computational prediction of protein functional annotations [https://pubmed.ncbi.nlm.nih.gov/40728605/]
- 44. harnessing the power of single-cell multi-omic data [https://www.nature.com/articles/s41540-023-00312-6]
- 45. Gene regulatory network prediction using machine learning ... [https://pmc.ncbi.nlm.nih.gov/articles/PMC12441907/]
- 46. Annotating protein functions via fusing multiple biological ... [https://www.nature.com/articles/s42003-024-07411-y]
- 47. Protein function prediction using GO similarity-based ... [https://www.nature.com/articles/s41598-025-04933-1]
- 48. Validation of gene regulatory network inference based on ... [https://pmc.ncbi.nlm.nih.gov/articles/PMC3860259/]
- 49. Gene regulatory network prediction using machine ... [https://pubmed.ncbi.nlm.nih.gov/40969532/]
- 50. Gene regulatory network prediction using machine ... [https://www.maxapress.com/article/id/68896e3dfa6c58586441d7e6]
- 51. Integrating artificial intelligence in drug discovery and early ... [https://pmc.ncbi.nlm.nih.gov/articles/PMC11909971/]

- 52. Automated drug design for druggable target identification ... [https://www.nature.com/articles/s41598-025-18091-x]
- 53. LABind: identifying protein binding ligand-aware sites via ... [https://www.nature.com/articles/s41467-025-62899-0]
- 54. Transformative advances in single-cell omics: a ... [https://translational-medicine.biomedcentral.com/articles/10.1186/s12967-025-07091-0]
- 55. A comparison of automatic cell identification methods for ... [https://genomebiology.biomedcentral.com/articles/10.1186/s13059-019-1795-z]
- 56. Article BioLLM: A standardized framework for integrating ... [https://www.sciencedirect.com/science/article/pii/S2666389925001746]
- 57. Multimodal learning enables chat-based exploration of ... [https://www.nature.com/articles/s41587-025-02857-9]
- 58. Multi-omics Data Integration, Interpretation, and Its Application [https://pmc.ncbi.nlm.nih.gov/articles/PMC7003173/]
- 59. A Guide to Multi-omics Integration Strategies [https://frontlinegenomics.com/a-guide-to-multi-omics-integration-strategies/]
- 60. Transformative advances in single-cell omics: a comprehensive review of foundation models, multimodal integration and computational ecosystems - PMC [https://pmc.ncbi.nlm.nih.gov/articles/PMC12560279/]
- 61. Web-based multi-omics integration using the Analyst ... [https://www.nature.com/articles/s41596-023-00950-4]
- 62. Incorporating physics to overcome data scarcity in predictive ... [https://pmc.ncbi.nlm.nih.gov/articles/PMC10529646/]
- 63. Current strategies to address data scarcity in artificial ... [https://www.sciencedirect.com/science/article/abs/pii/S0010482524008199]
- 64. noisyR: enhancing biological signal in sequencing ... [https://pmc.ncbi.nlm.nih.gov/articles/PMC8373073/]
- 65. Comprehensive noise reduction in single-cell data with the ... [https://www.sciencedirect.com/science/article/pii/S2667237525002140]
- 66. Denoising large-scale biological data using network filters [https://bmcbioinformatics.biomedcentral.com/articles/10.1186/s12859-021-04075-x]
- 67. Decoding the mystery: AI-assisted bioinformatics and ... [https://www.frontiersin.org/journals/plant-science/articles/10.3389/fpls.2025.1678483/full]
- 68. A Comprehensive Review of Explainable AI in ... [https://arxiv.org/html/2312.06082v1]
- 69. Explainable AI: A Review of Machine Learning ... [https://pmc.ncbi.nlm.nih.gov/articles/PMC7824368/]
- 70. Towards a Transparent and Interpretable AI Model for ... [https://arxiv.org/html/2509.16685v1]
- 71. Framework for Classifying Explainable Artificial Intelligence ... [https://pmc.ncbi.nlm.nih.gov/articles/PMC10689048/]
- 72. Interpretability: from black boxes to white boxes [https://www.vanderschaar-lab.com/from-black-boxes-to-white-boxes/]
- 73. Explainability and white box in drug discovery [https://pubmed.ncbi.nlm.nih.gov/37105727/]
- 74. Interpretable AI: How White-Box Models Provide ... [https://medium.com/data-science-collective/interpretable-ai-how-white-box-models-provide-transparency-736493465e91]
- 75. Multi-model assurance analysis showing large language ... [https://pmc.ncbi.nlm.nih.gov/articles/PMC12318031/]
- 76. Hallucination Mitigation for Retrieval-Augmented Large ... [https://www.mdpi.com/2227-7390/13/5/856]
- 77. Theoretical Foundations and Mitigation of Hallucination in ... [https://arxiv.org/abs/2507.22915]
- 78. Foundation models for materials discovery – current state ... [https://www.nature.com/articles/s41524-025-01538-0]
- 79. A review of transformer models in drug discovery and beyond [https://www.sciencedirect.com/science/article/pii/S2095177924001783]
- 80. Single-cell foundation models: bringing artificial ... [https://pmc.ncbi.nlm.nih.gov/articles/PMC12586647/]
- 81. - Fine a sentence transformer for DNA | BMC tuning Bioinformatics [https://bmcbioinformatics.biomedcentral.com/articles/10.1186/s12859-025-06291-1]
- 82. ai4protein/VenusFactory: Easy data acquisition, benchmark ... [https://github.com/ai4protein/VenusFactory]
- 83. Revolutionizing clinical decision making through deep ... [https://www.nature.com/articles/s41598-025-12679-z]
- 84. MergeDNA: Context-aware Genome Modeling with ... [https://arxiv.org/html/2511.14806v1]
- 85. Assessing and mitigating batch effects in large-scale omics ... [https://genomebiology.biomedcentral.com/articles/10.1186/s13059-024-03401-9]
- 86. reComBat: batch-effect removal in large-scale multi-source ... [https://pmc.ncbi.nlm.nih.gov/articles/PMC9710604/]
- 87. Batch Effect Correction [https://www.10xgenomics.com/analysis-guides/introduction-batch-effect-correction]
- 88. A technical review of multi-omics data integration methods [https://pmc.ncbi.nlm.nih.gov/articles/PMC12315550/]
- 89. Flexynesis: A deep learning toolkit for bulk multi-omics data ... [https://www.nature.com/articles/s41467-025-63688-5]
- 90. Benchmarking genomic language models | Nature Methods [https://www.nature.com/articles/s41592-025-02829-6]
- 91. BioProBench: Comprehensive Dataset and Benchmark in ... [https://arxiv.org/html/2505.07889v1]
- 92. a benchmark suite for long-range DNA prediction tasks - PMC [https://pmc.ncbi.nlm.nih.gov/articles/PMC12627797/]
- 93. Benchmarking foundation cell models for post-perturbation ... [https://bmcgenomics.biomedcentral.com/articles/10.1186/s12864-025-11600-2]
- 94. Unleashing Scientific Reasoning for Bio-experimental ... [https://arxiv.org/html/2510.15600v1]
- 95. Specialized Foundation Models Struggle to Beat ... [https://arxiv.org/html/2411.02796v2]
- 96. Why simple models outperformed deep learning in gene ... [https://www.linkedin.com/posts/vsavova_deep-learning-based-gene-perturbation-effect-activity-7360314175551639552-mjBK]
- 97. CellFM: a large-scale foundation model pre-trained on ... [https://www.nature.com/articles/s41467-025-59926-5]
- 98. Zero-shot evaluation reveals limitations of single-cell ... [https://genomebiology.biomedcentral.com/articles/10.1186/s13059-025-03574-x]

- 99. Zero-shot evaluation reveals limitations of foundation ... [https://www.microsoft.com/en-us/research/articles/zero-shot-evaluation-reveals-limitations-of-foundation-models-in-single-cell-biology/]
- 100. The Invisible Framework Organizing Life: Biological Ontologies [https://www.biochemicalsci.com/posts/the-invisible-framework-organizing-life-biological-ontologies]
- 101. Biology-driven insights into the power of single-cell ... [https://pubmed.ncbi.nlm.nih.gov/41044630/]
- 102. Improving Genomic Models via Task-Specific Self-Pretraining [https://arxiv.org/html/2506.17766v1]